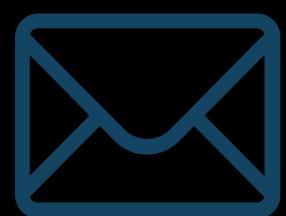


robw4

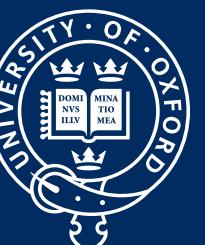
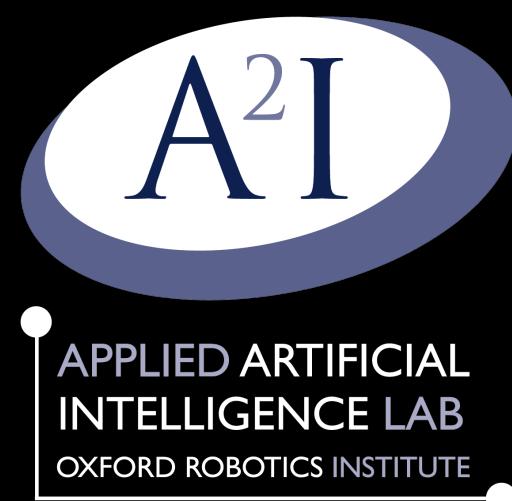


robw@robots.ox.ac.uk

# There And Back Again: Learning to Simulate Radar Data for Real World Applications

Rob Weston, Oiwi Parker Jones, Ingmar Posner

1524

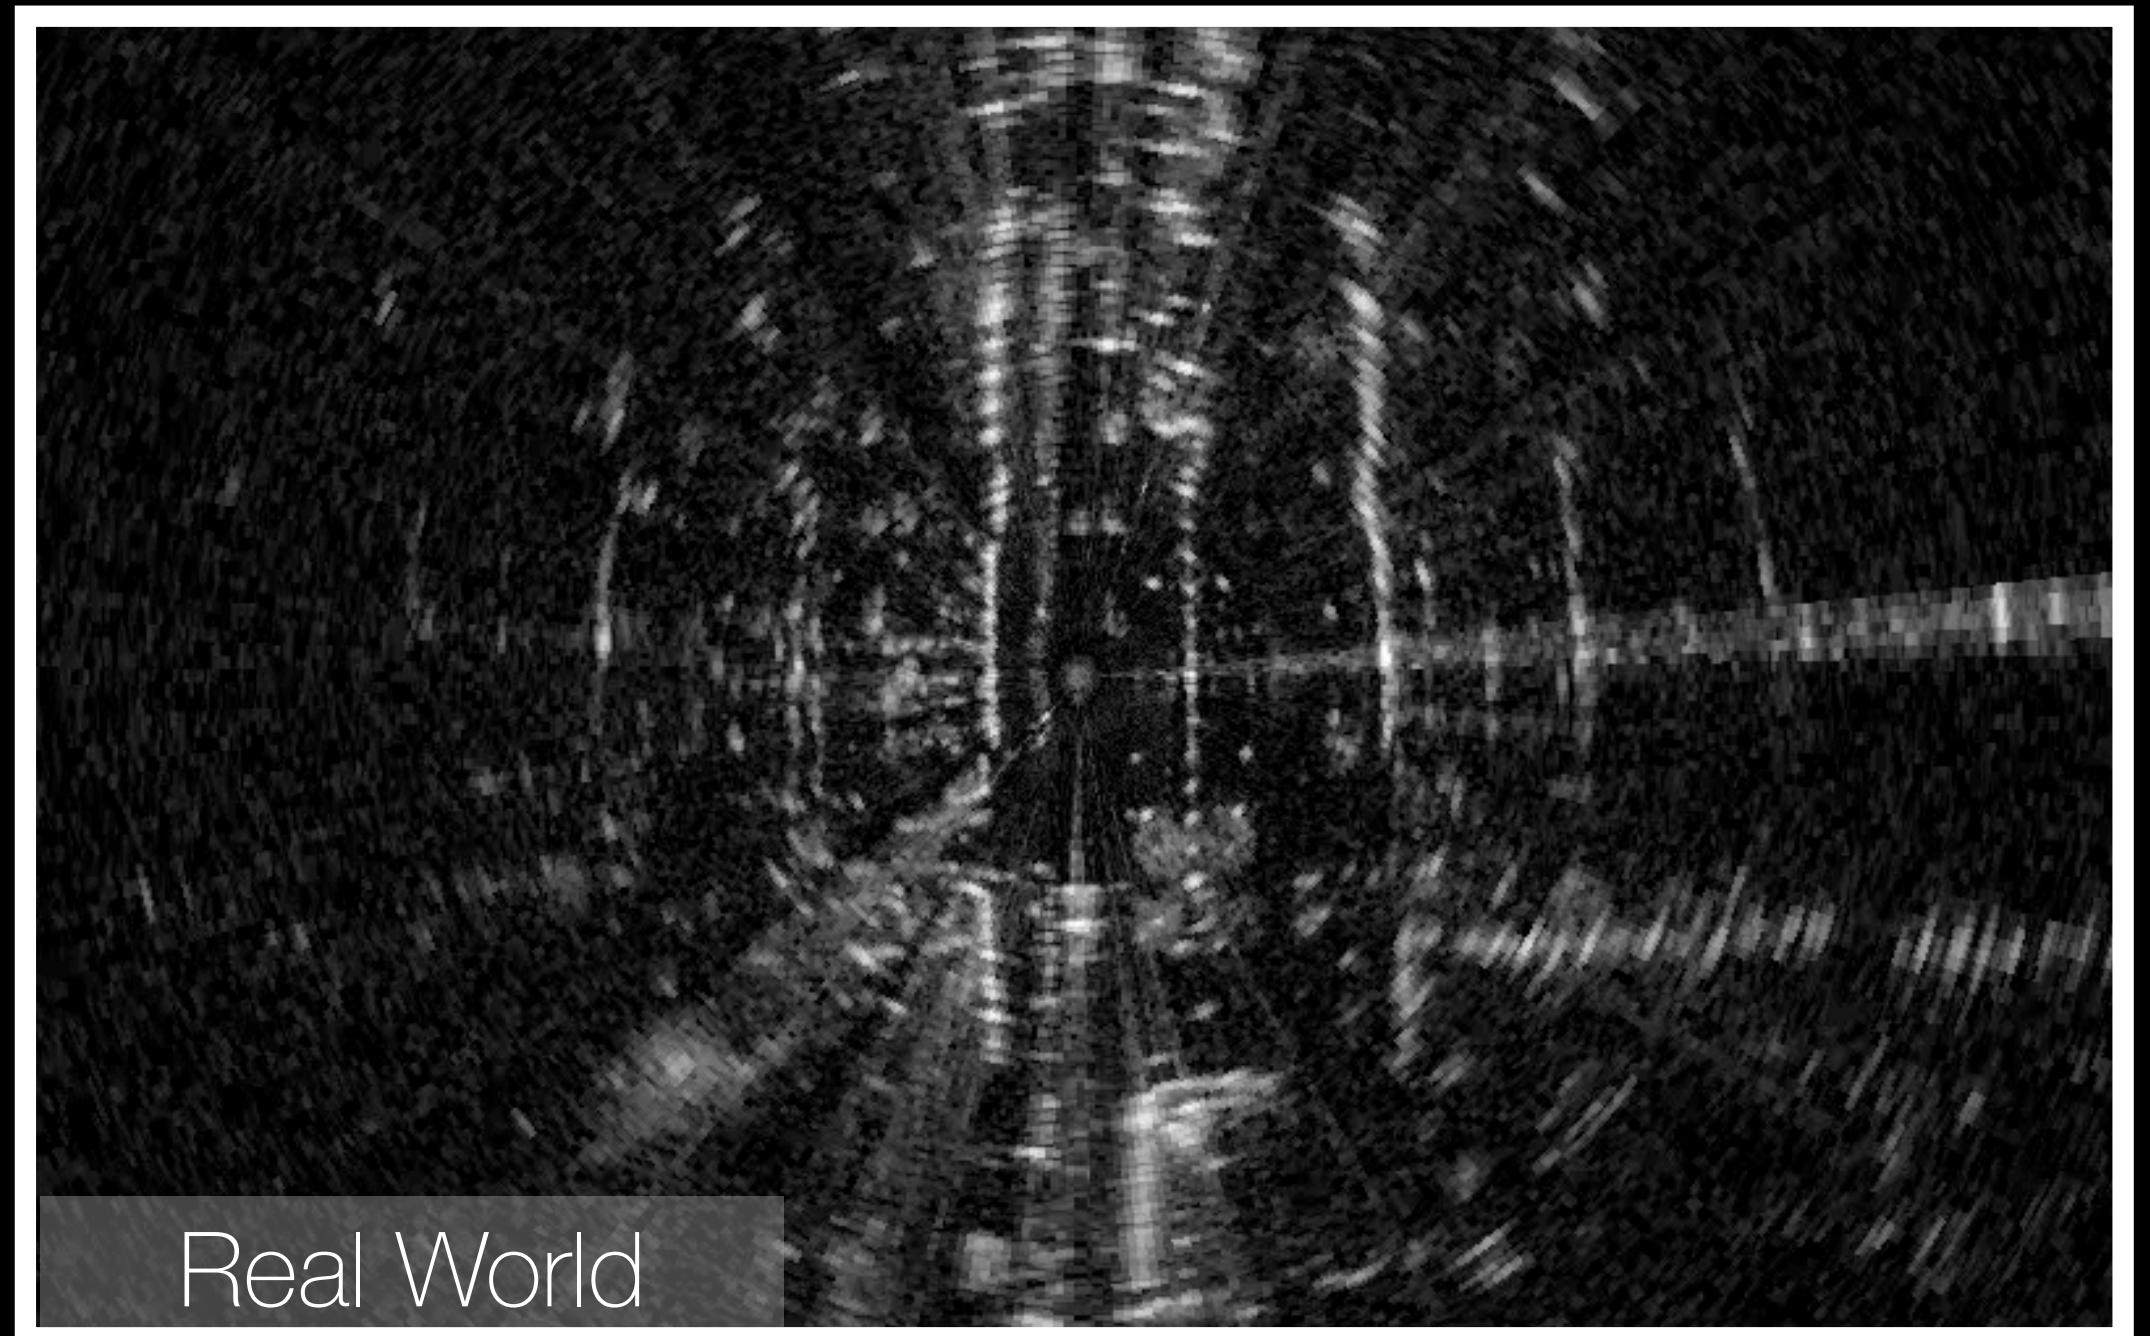


UNIVERSITY OF  
OXFORD

# Introduction

## Context

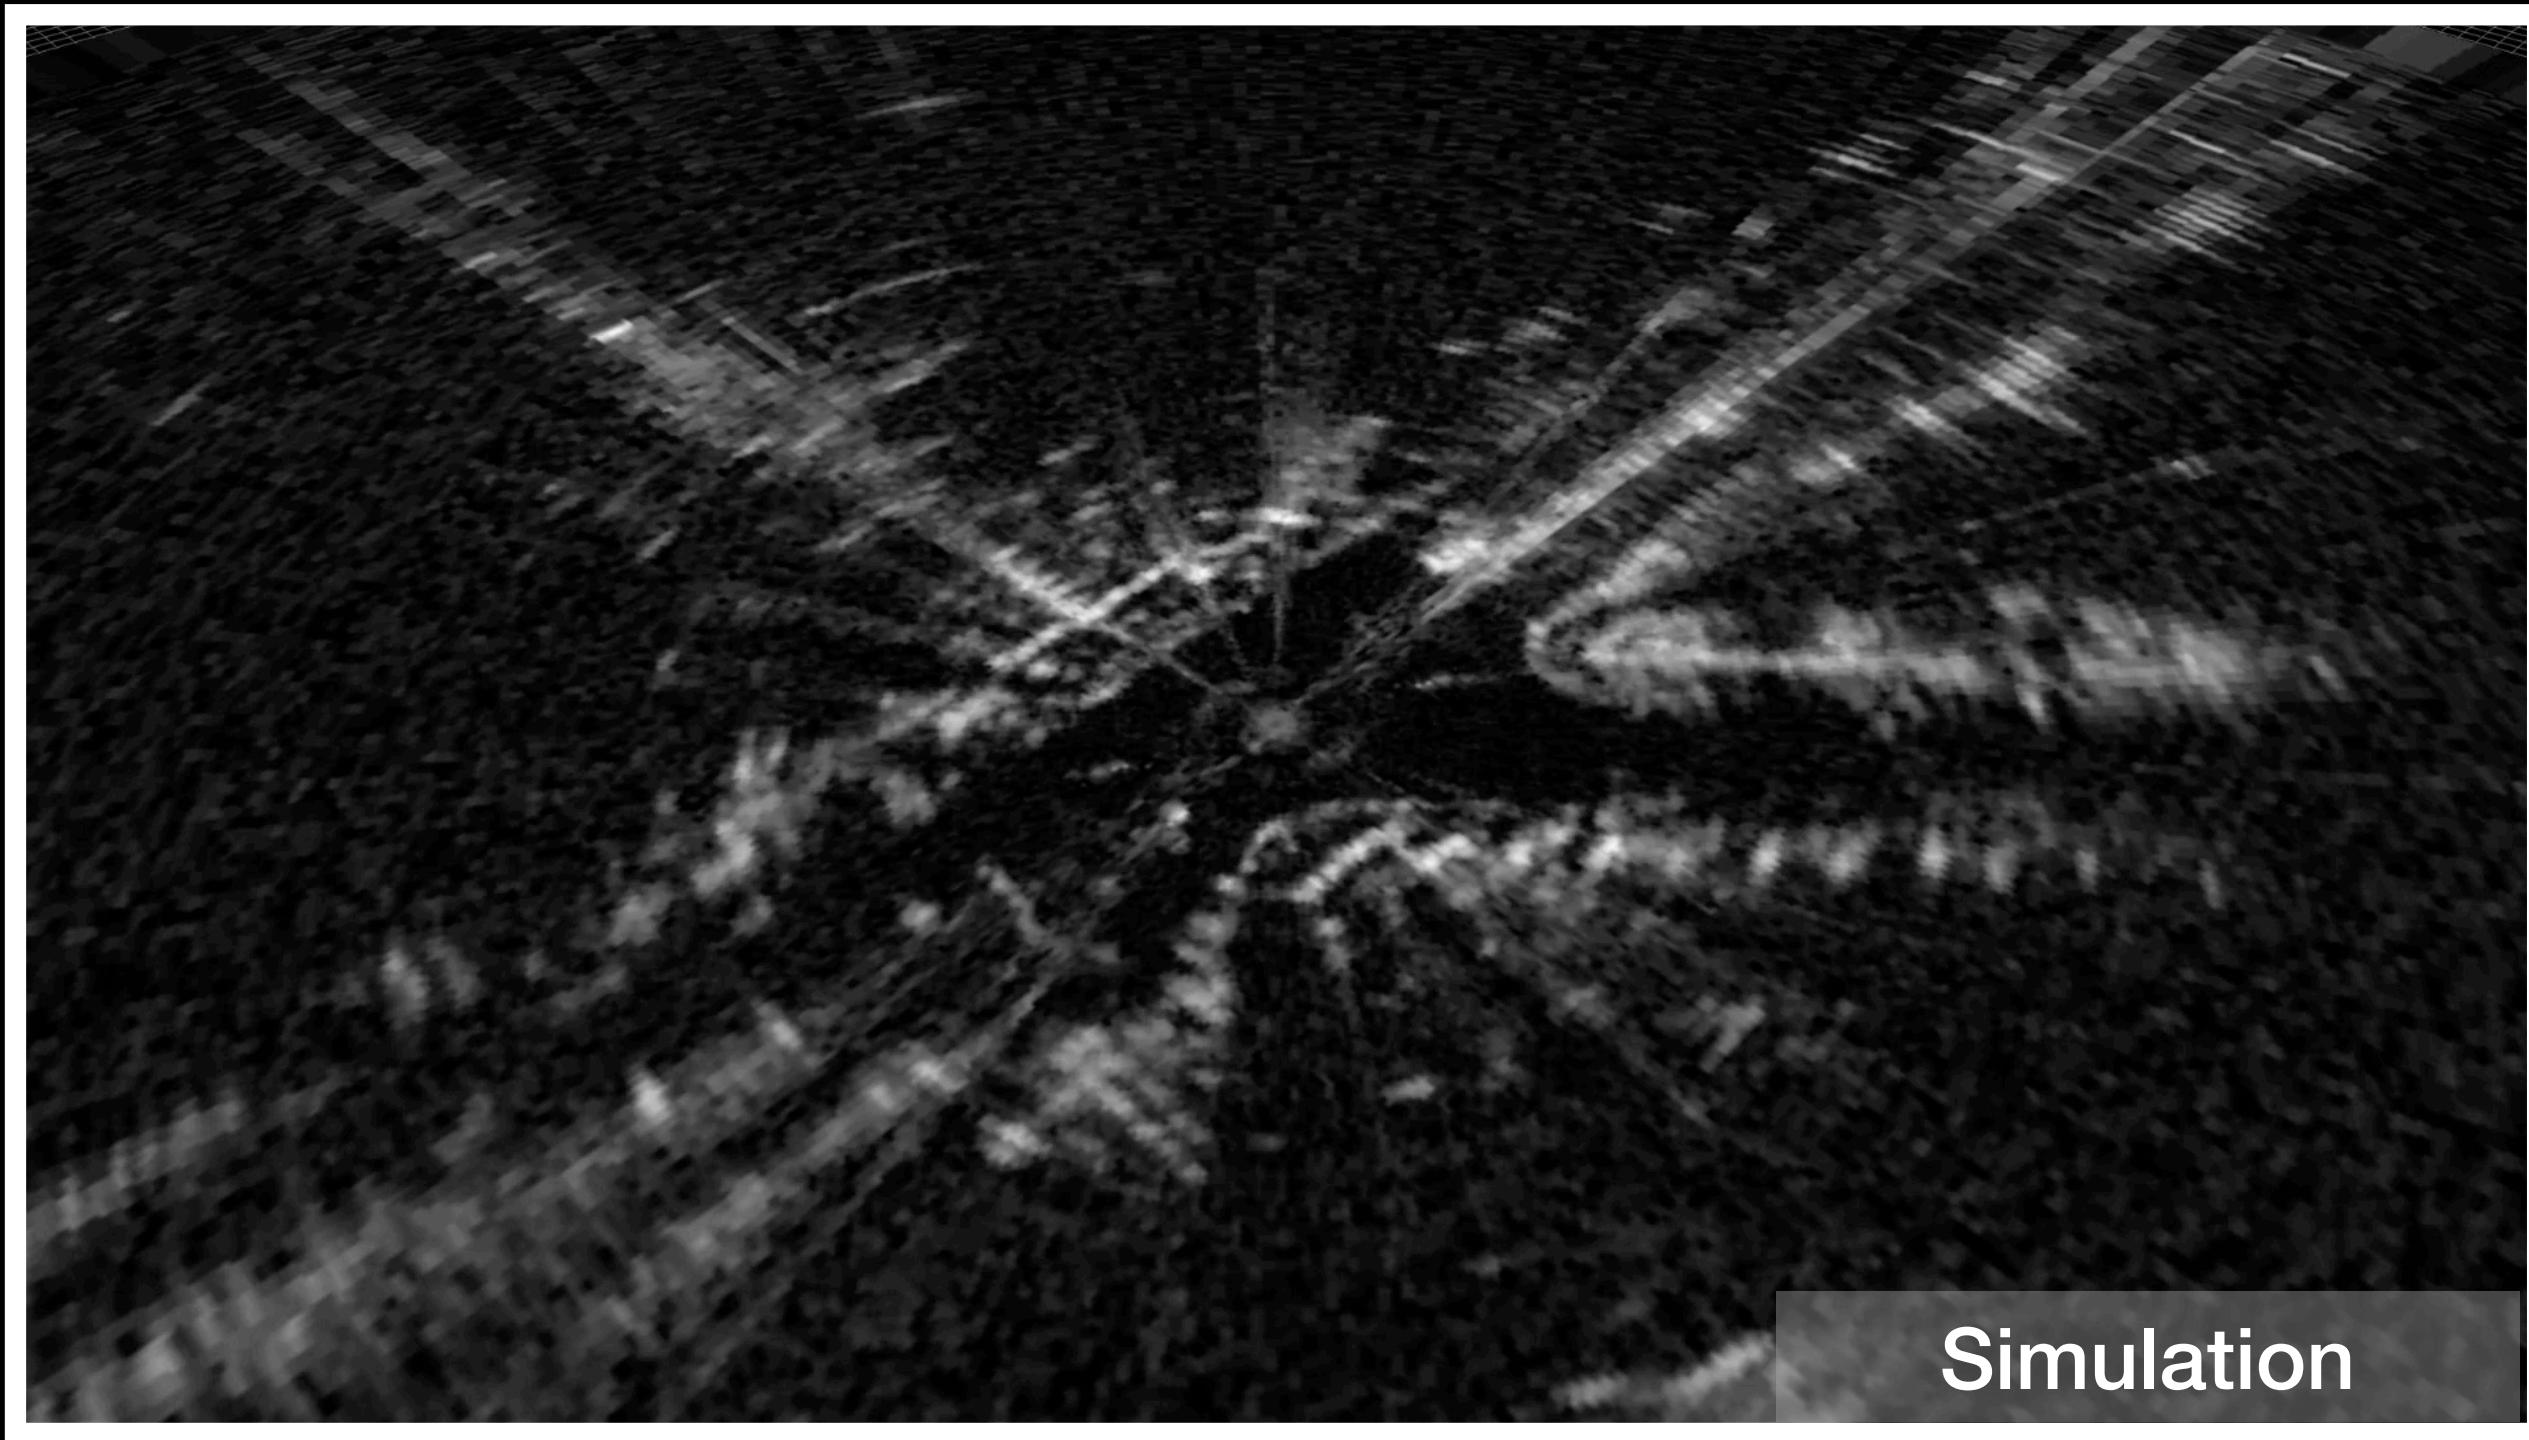
- Radar is a **promising** alternative to vision and lidar
- But notoriously challenging to work with as a result of challenging **noise artefacts**
- Data-driven approaches have made significant strides in using radar for a range of tasks in recent years
- But currently **limited by** size, quality and labelling of **datasets!**



Simulation has an **important** role to play...

# Introduction

## Our Aim



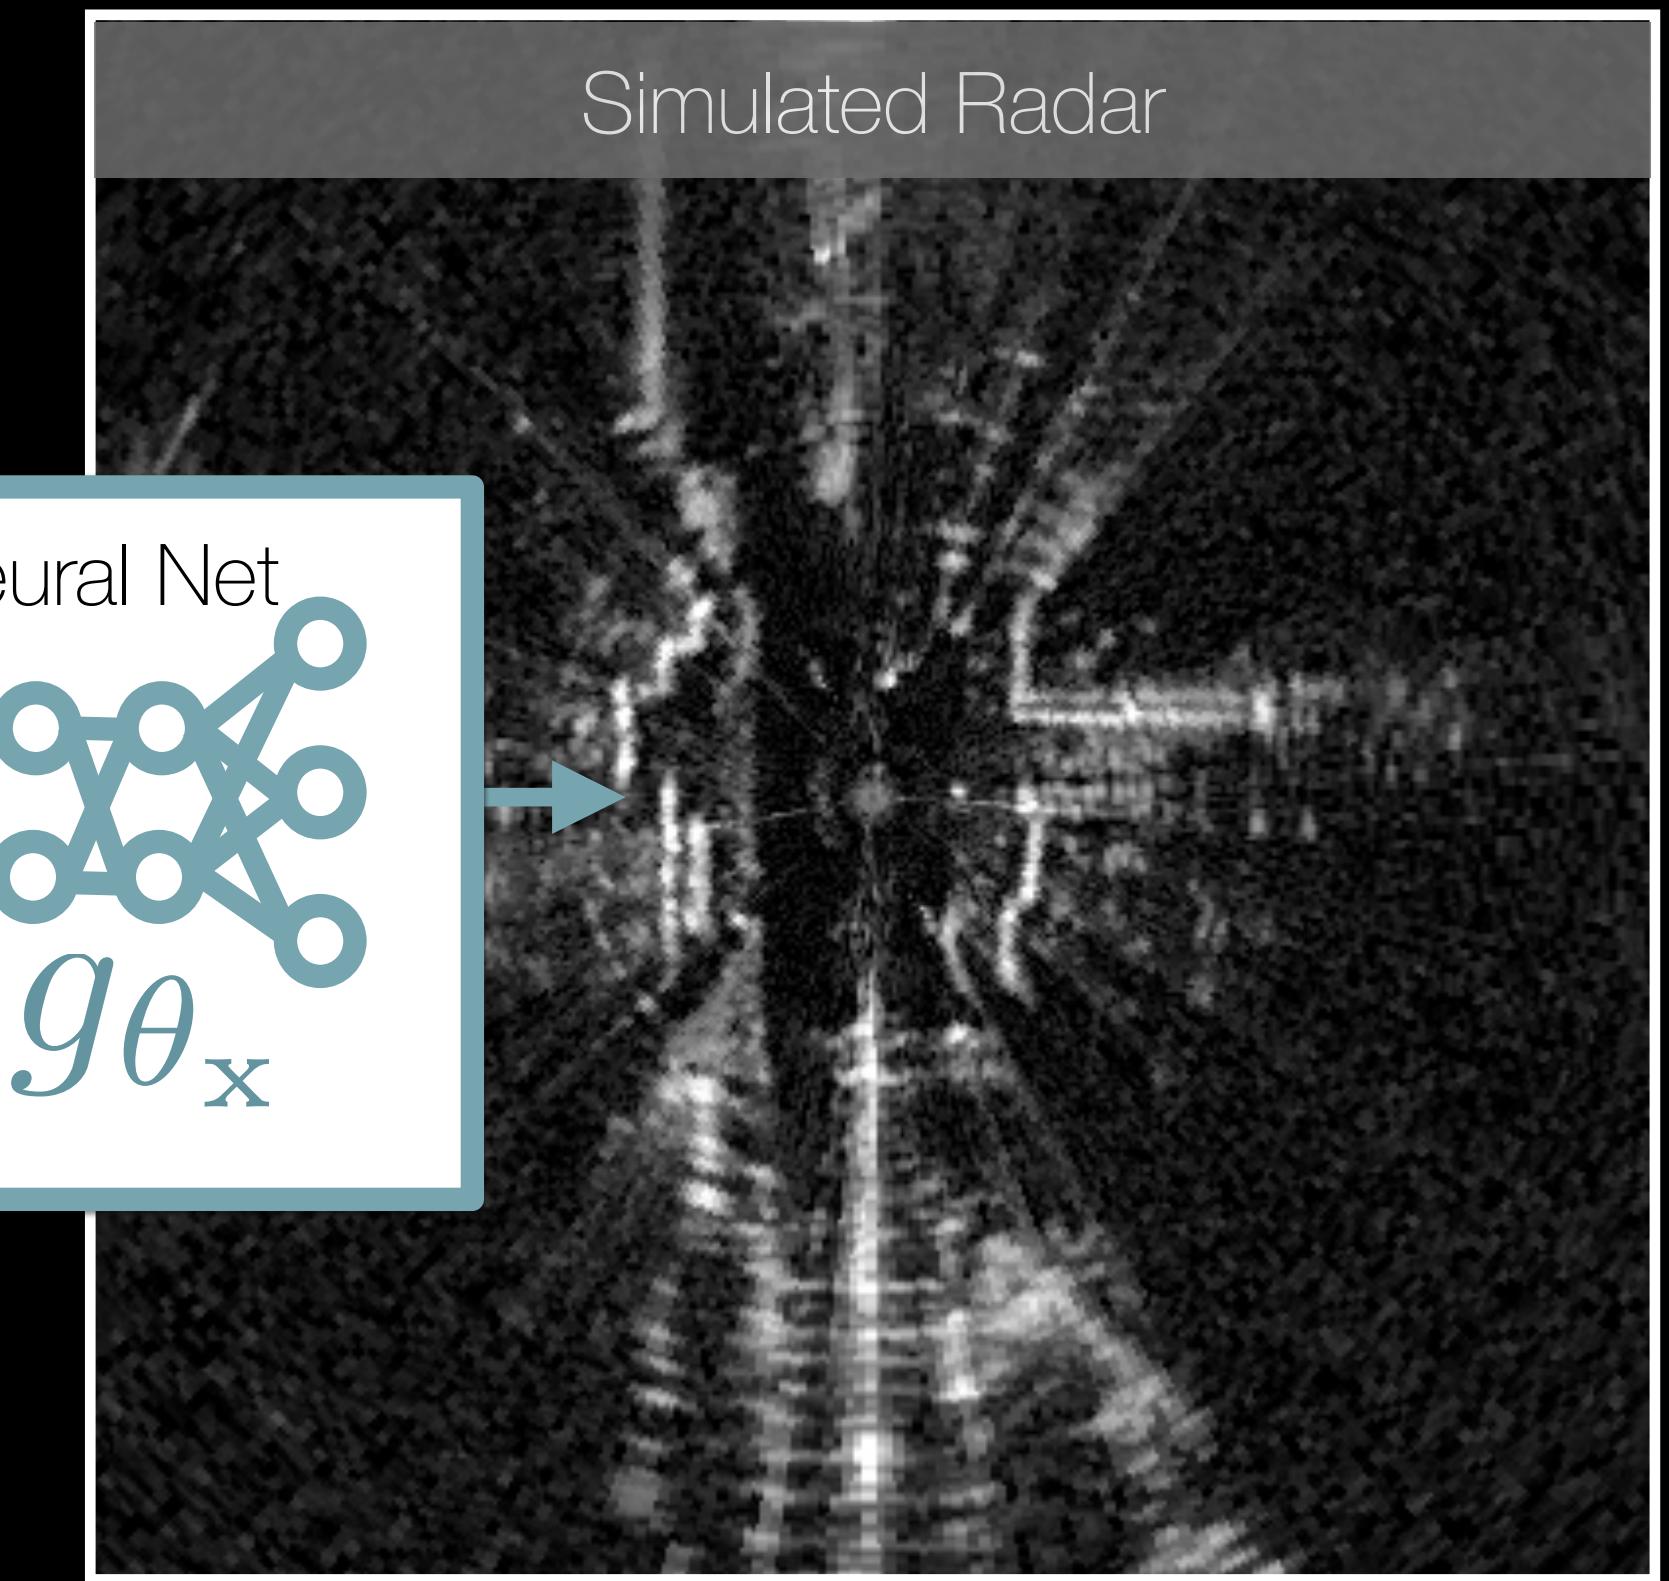
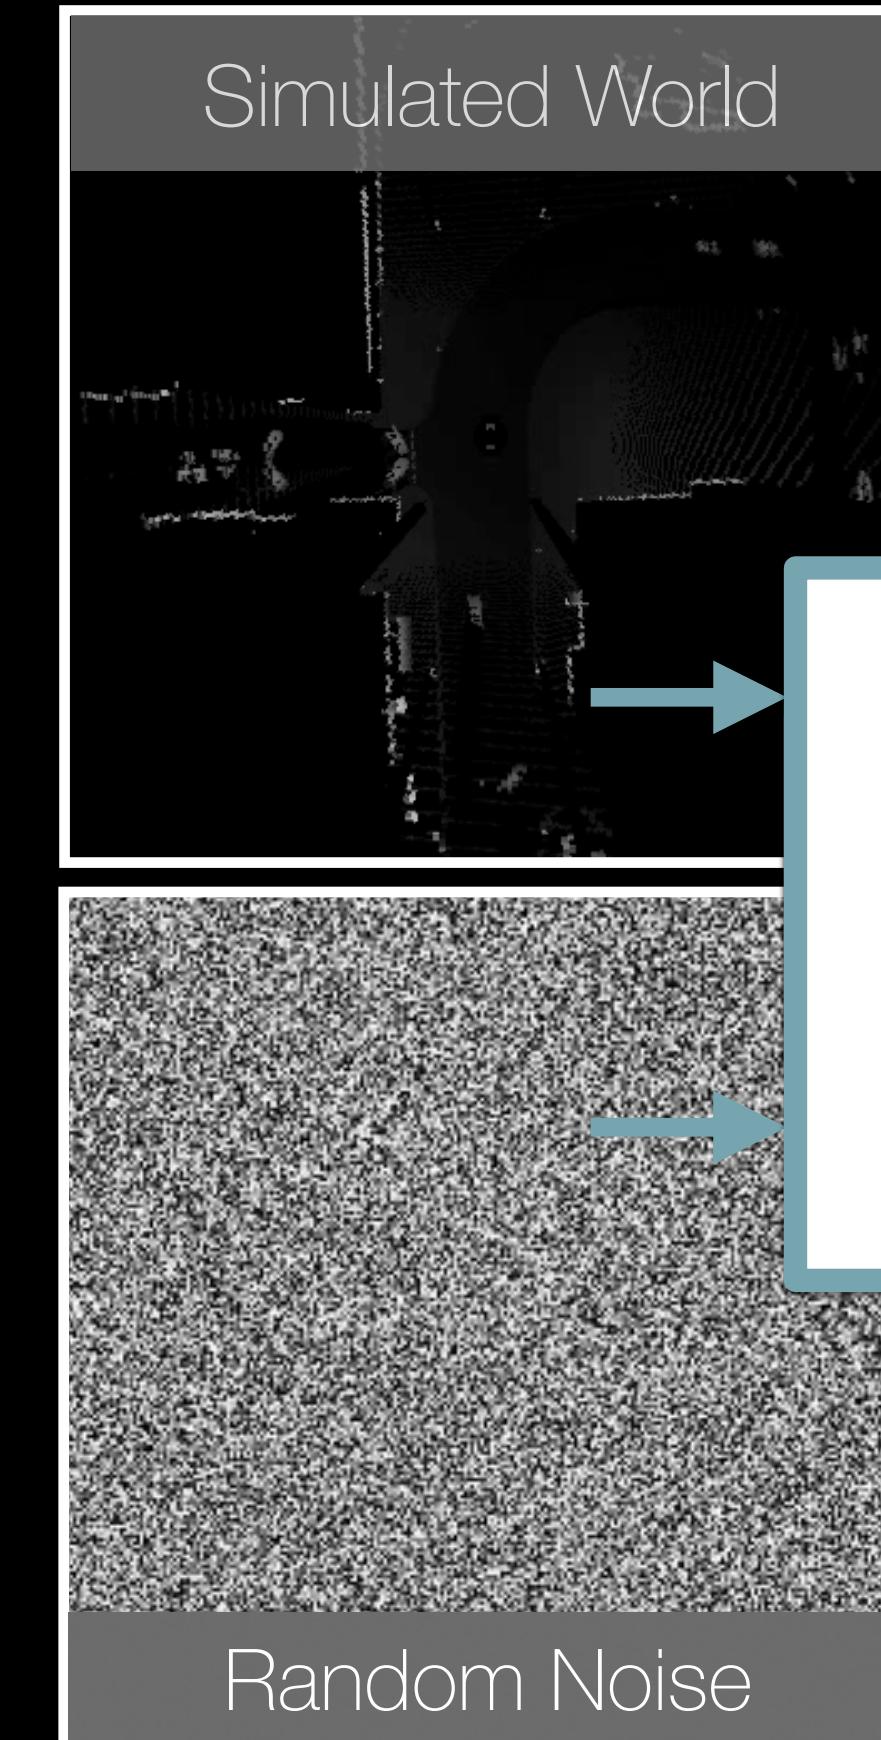
“Learn a radar sensor model, interfacing with an **existing** simulation environment, to train **new** models in simulation”

# Our Approach

## Stochastic Sensor Modelling with **Deep Implicit Model**



$\approx$



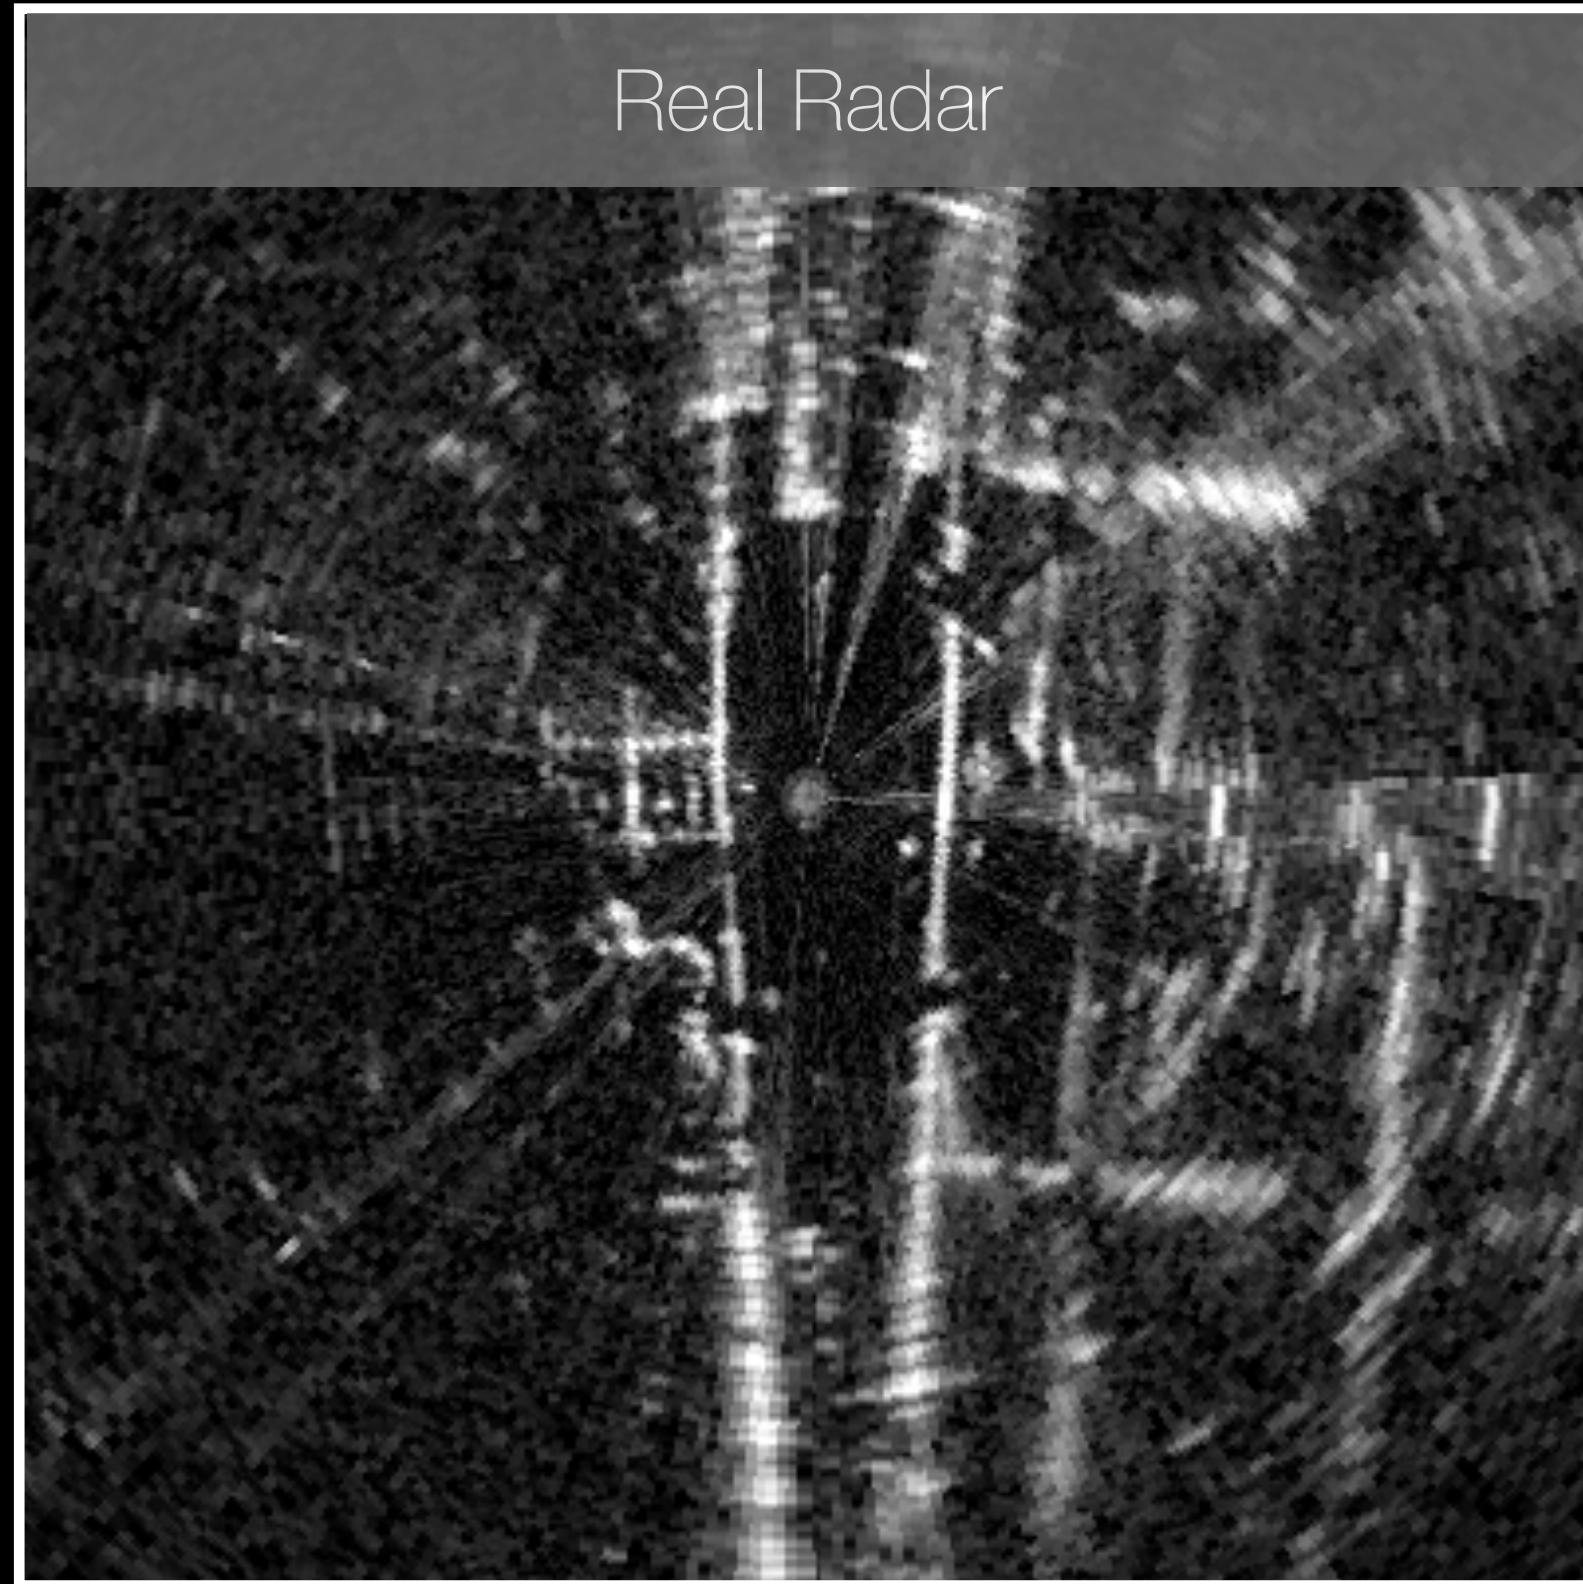
Real Radar is inherently **stochastic**

Simulate **stochastic** radar observations using deep implicit model

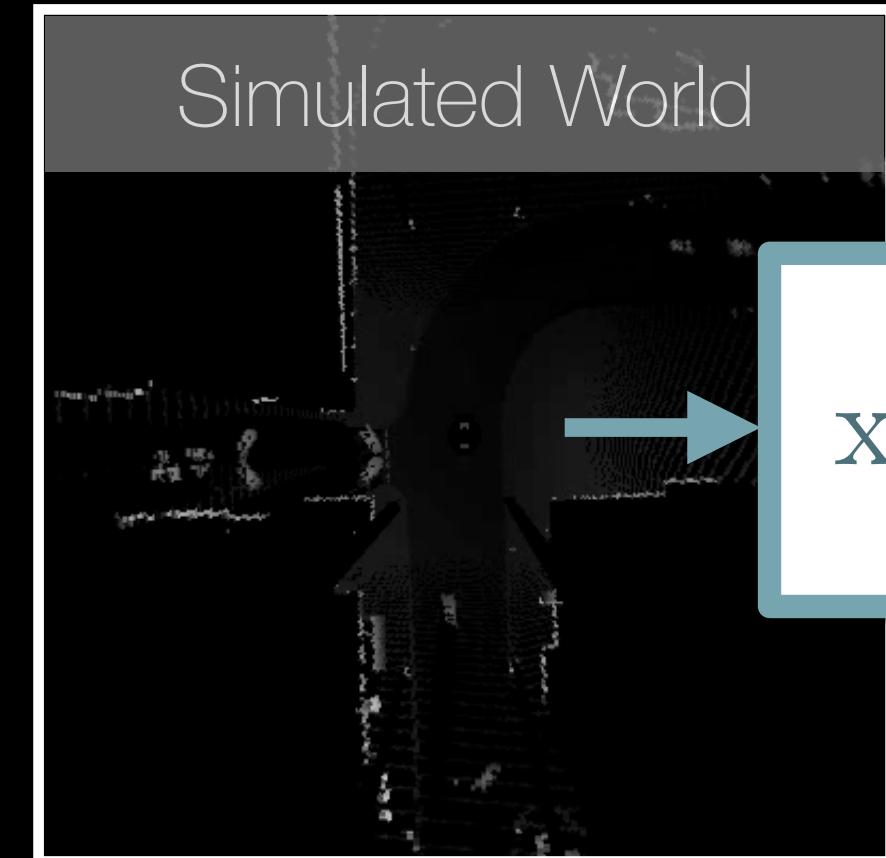
# Our Approach

## Stochastic Sensor Modelling with **Deep Implicit Model**

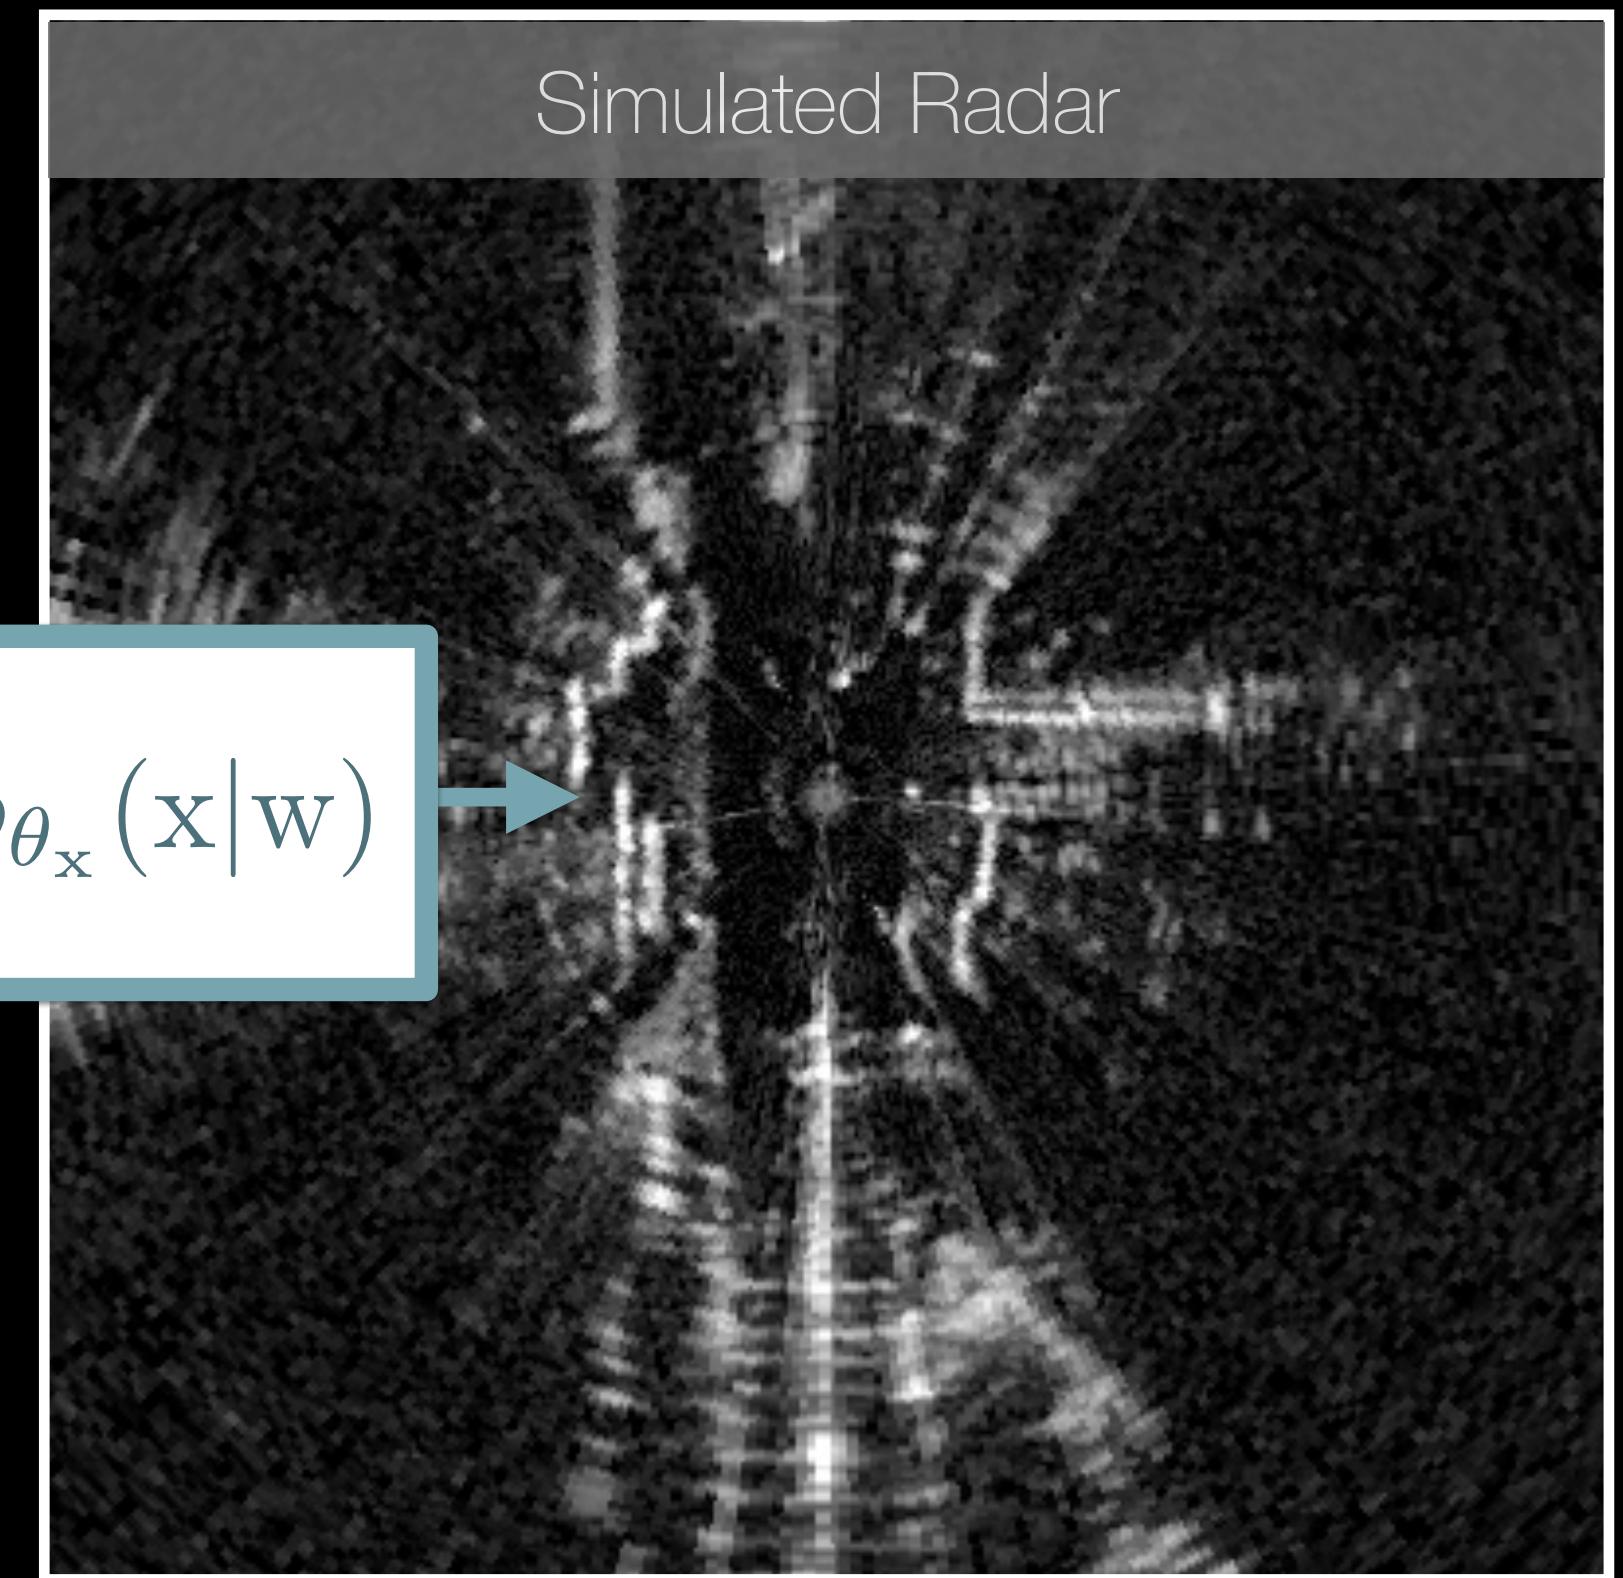
Implicitly capture **distribution** over radar observations



$\approx$



$$x \sim p_{\theta_x}(x|w)$$

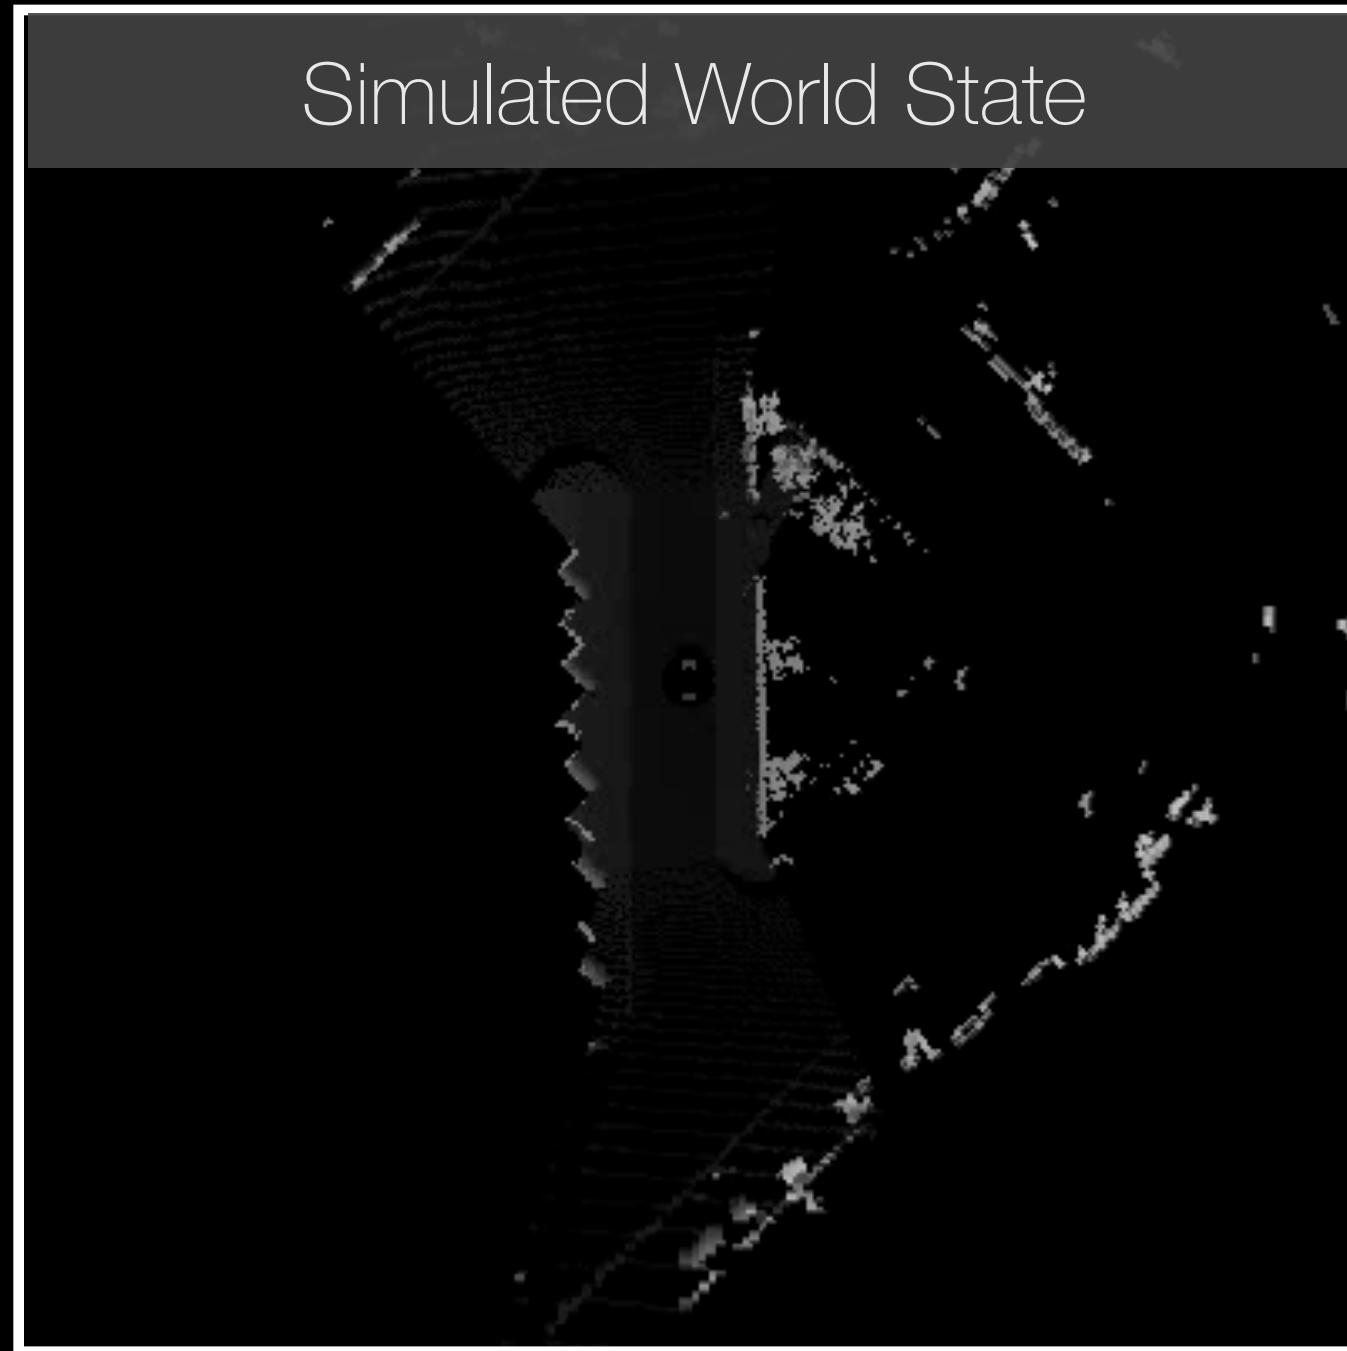


Real Radar is inherently **stochastic**

# Our Approach

Training data generated **automatically** in **simulation** or in the **real world**

$$\mathbf{S} = \{\mathbf{w}_l\}_{l=1}^L$$



World state **W** characterised as an  
**elevation map** generated using **CARLA**

# Our Approach

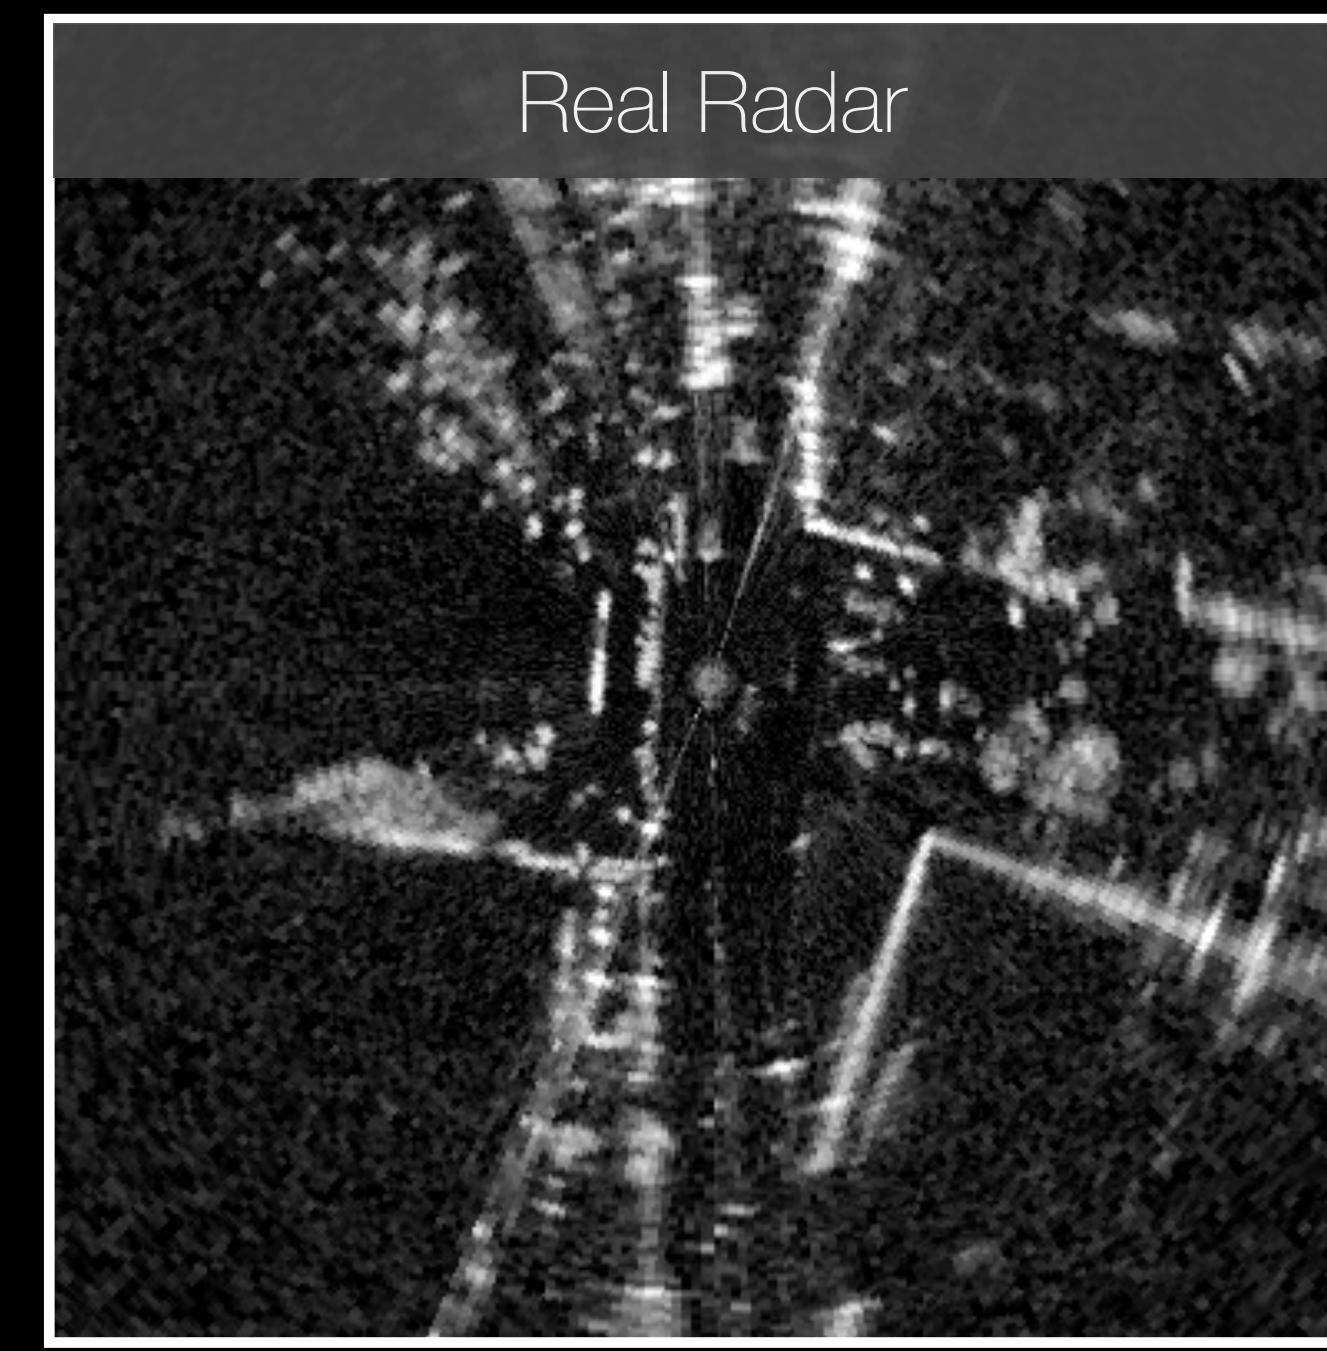
Training data generated **automatically** in **simulation** or in the **real world**

$$\mathbf{S} = \{\mathbf{w}_l\}_{l=1}^L$$

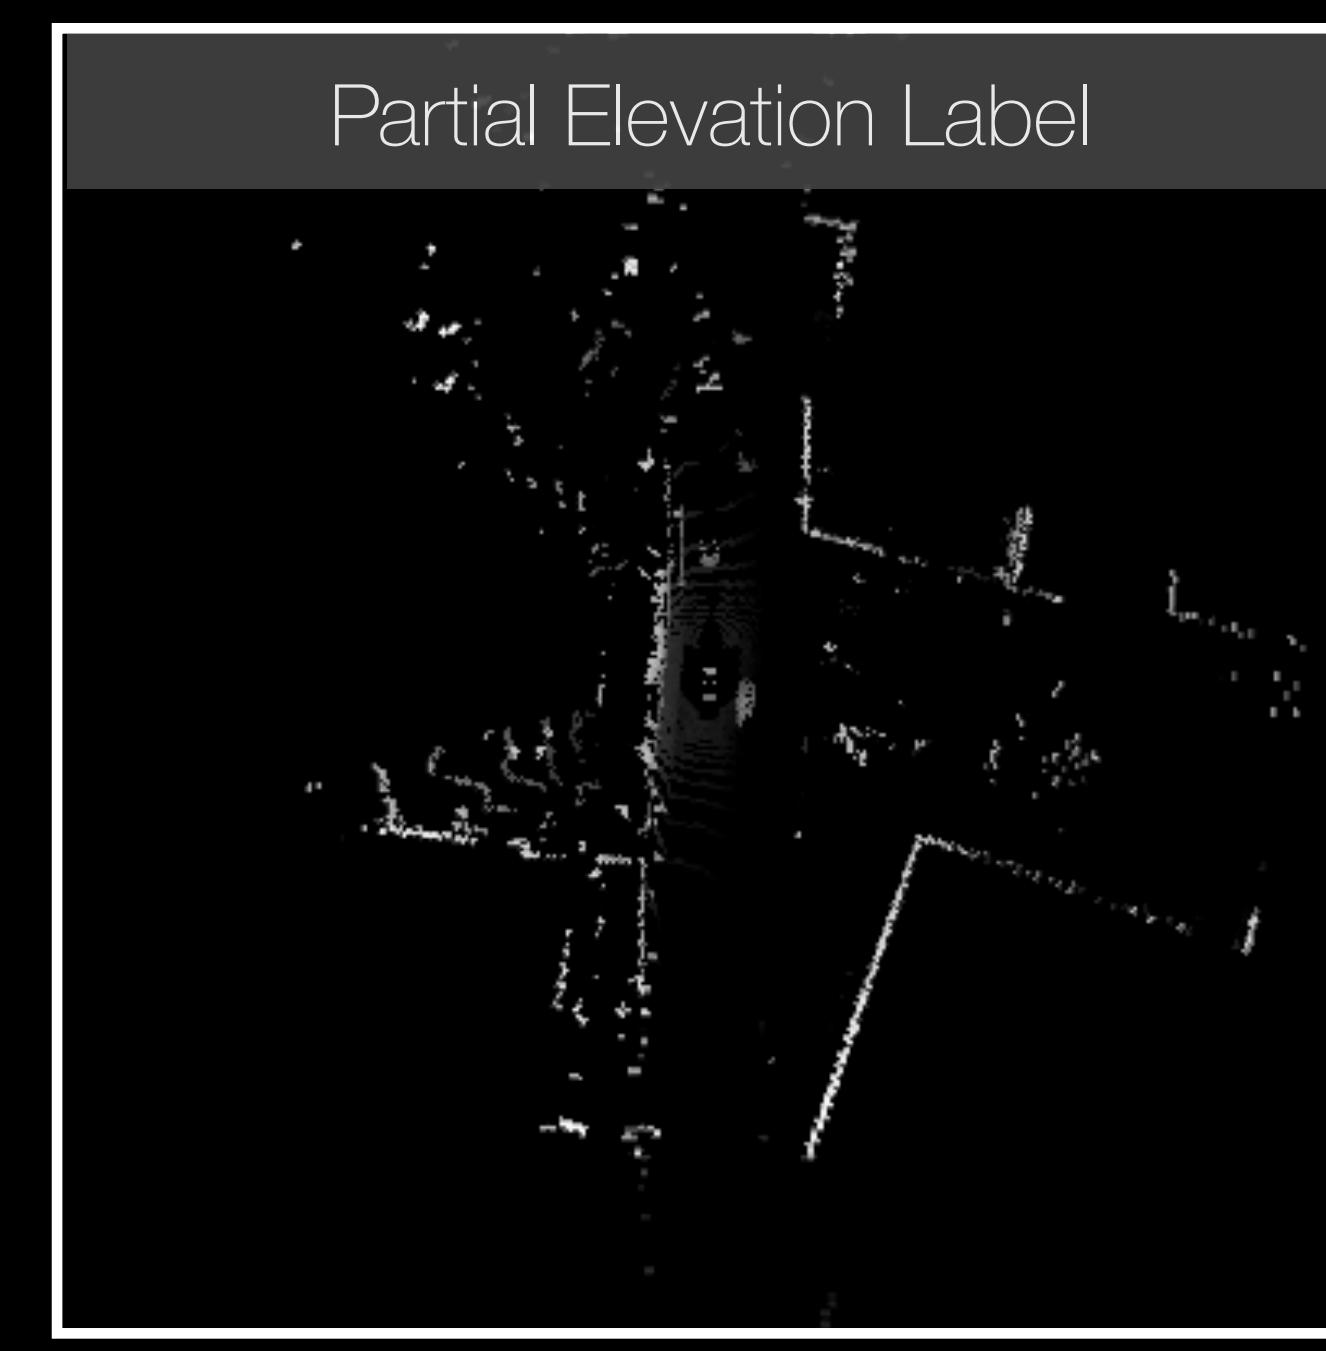
$$\mathbf{R} = \{(\mathbf{x}^*, \mathbf{m}^*)_n\}_{n=1}^N$$



Simulated World State



Real Radar



Partial Elevation Label

World state **W** characterised as an **elevation map** generated using **CARLA**

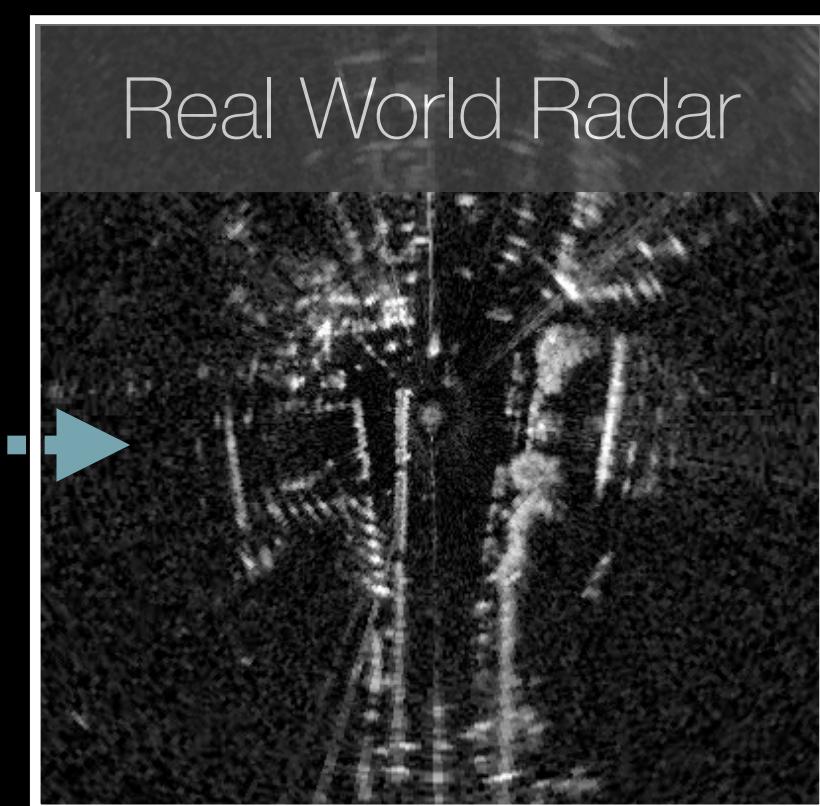
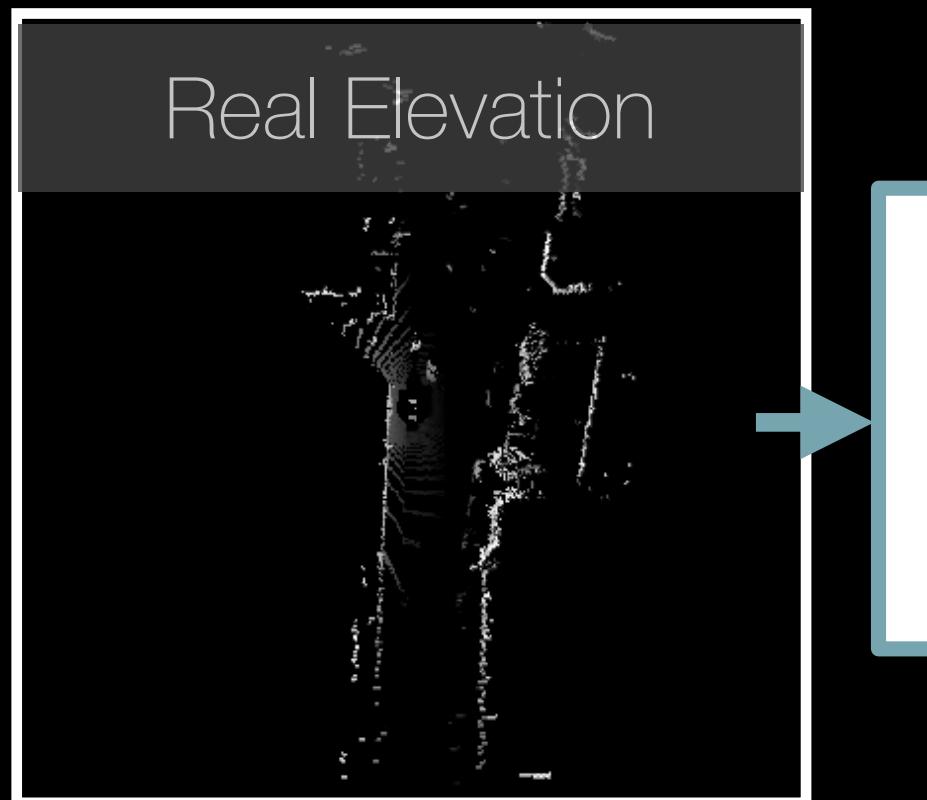
Real radar observations **X\*** from **Radar Oxford 10k** dataset

**Partial** elevation labels **m\*** **automatically** generated from lidar

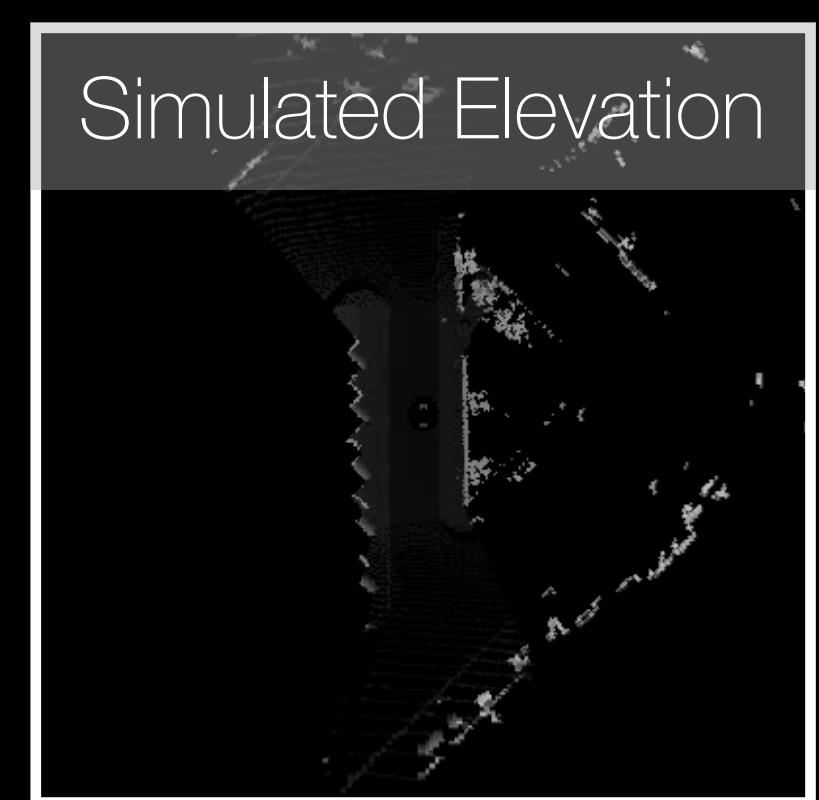
# One Possible Approach?

Train model to simulate radar from **real** world elevation

(1) Train on **real** data  $\mathbf{R} = \{(\mathbf{x}^*, \mathbf{m}^*)_n\}_{n=1}^N$



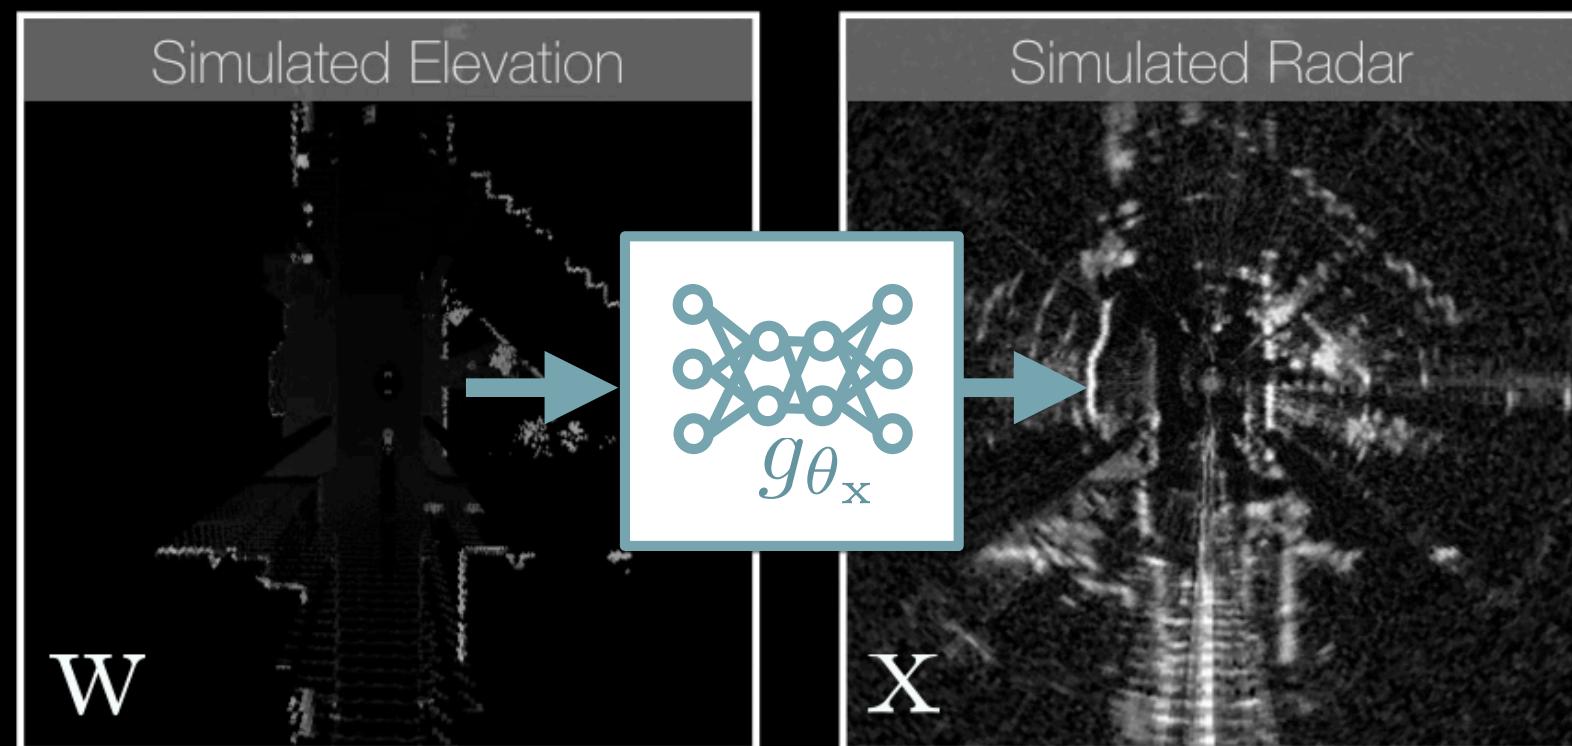
(2) Test on **simulated** data  $\mathbf{S} = \{\mathbf{w}_l\}_{l=1}^L$



Models trained **only** on real elevation  
poorly generalise to simulated elevation  
states

# Our Approach

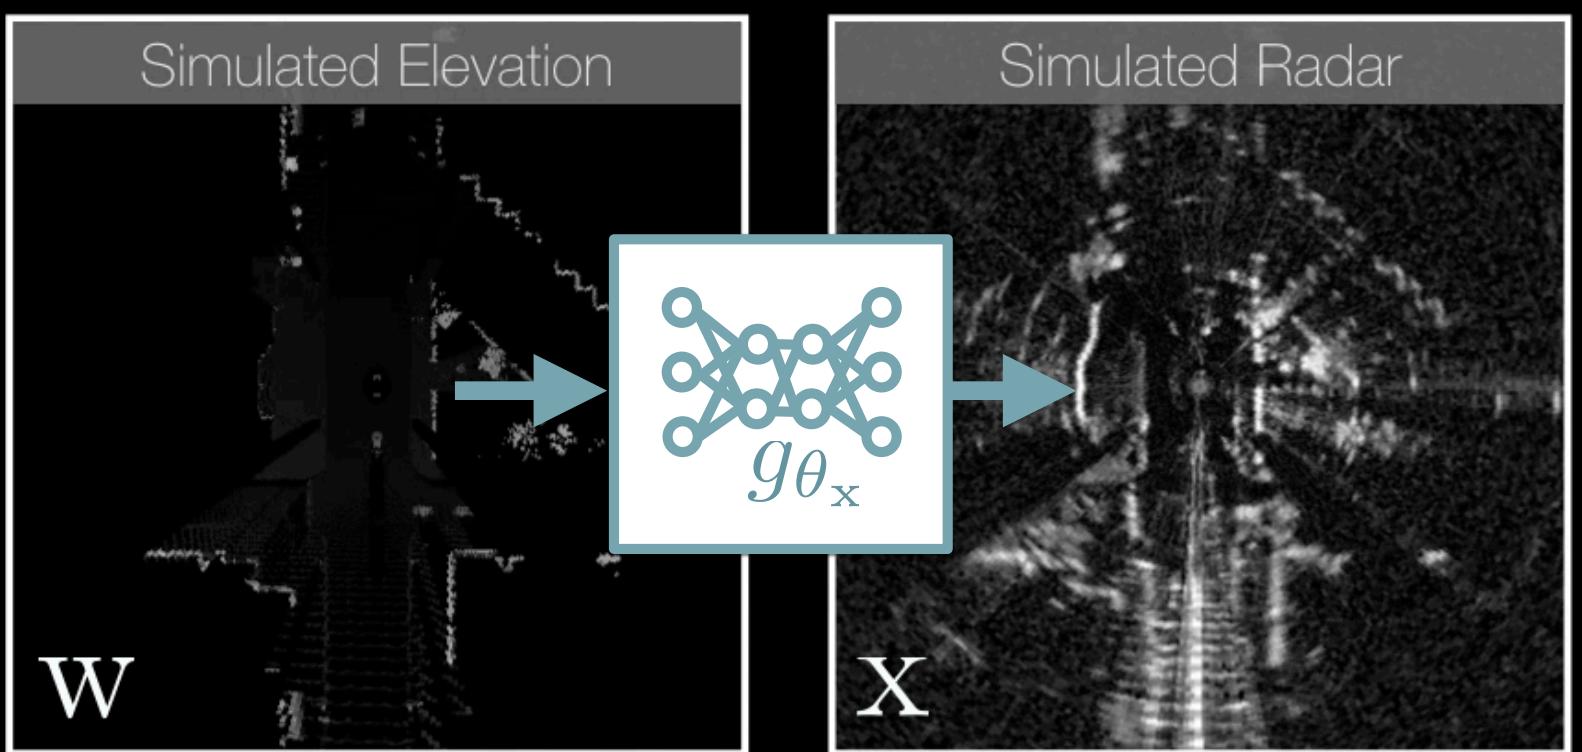
Predict radar from **simulated states** using **adversarial** loss



# Our Approach

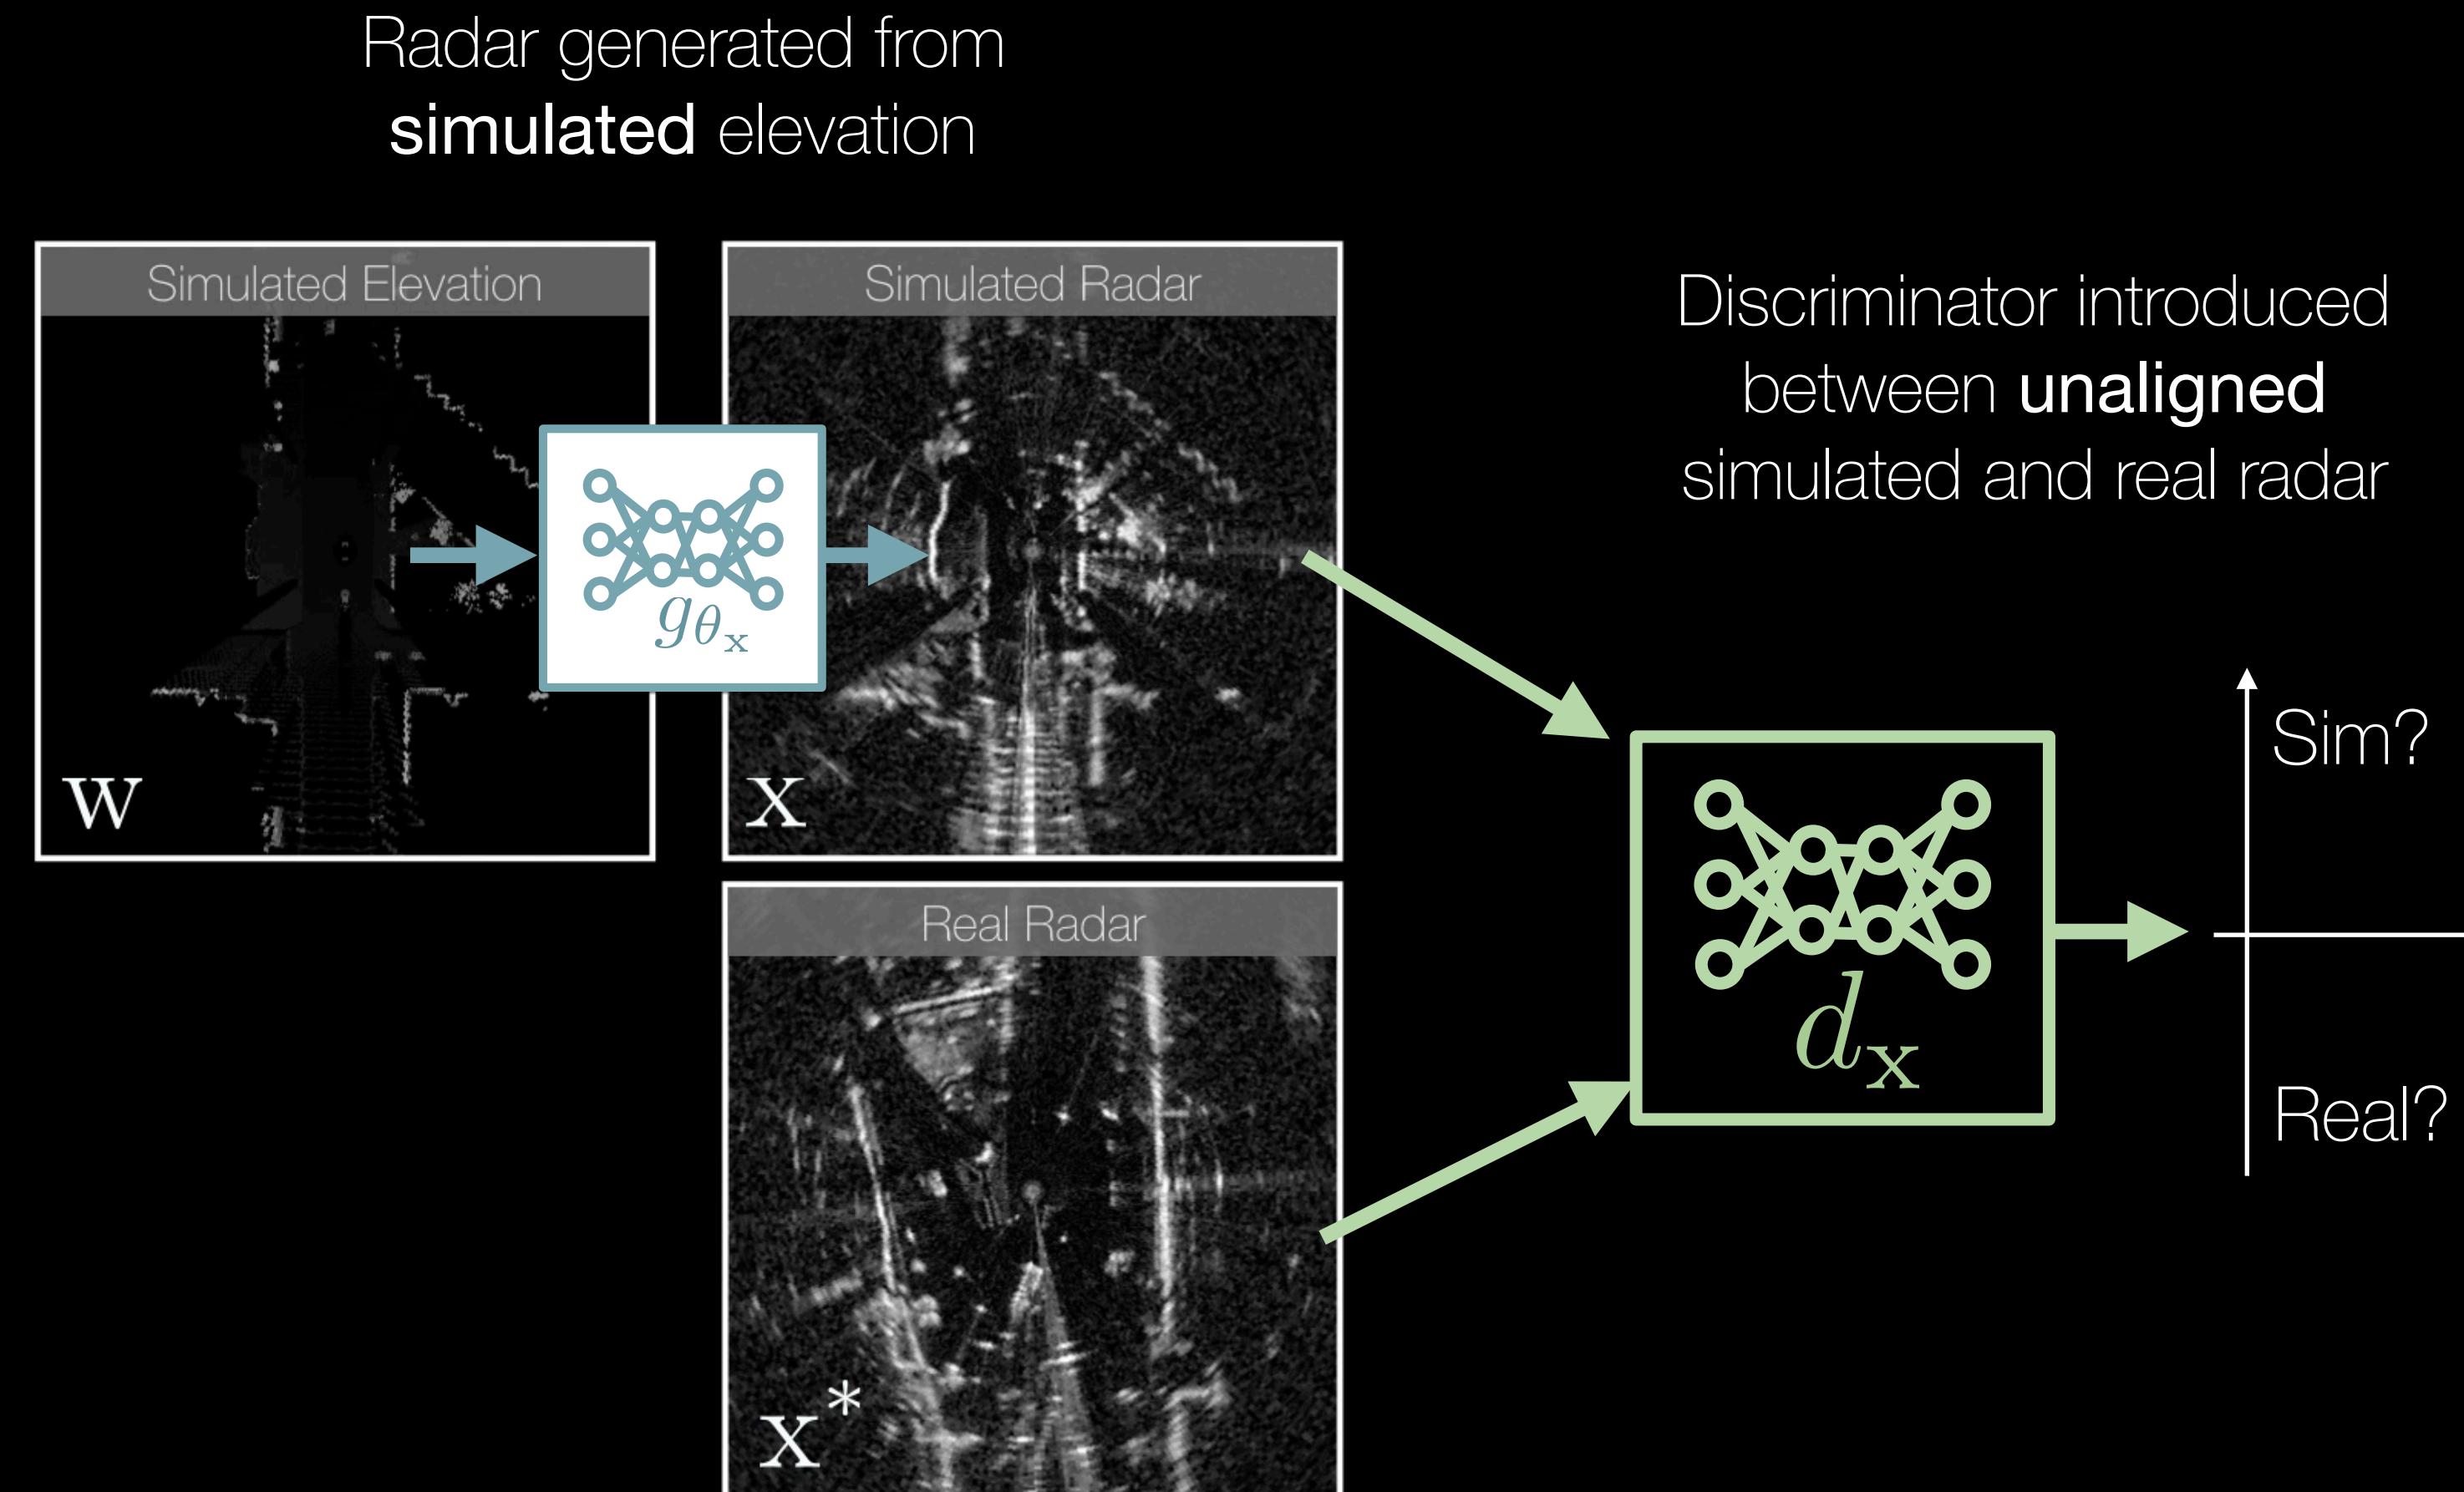
Predict radar from **simulated states** using **adversarial** loss

Radar generated from  
**simulated** elevation



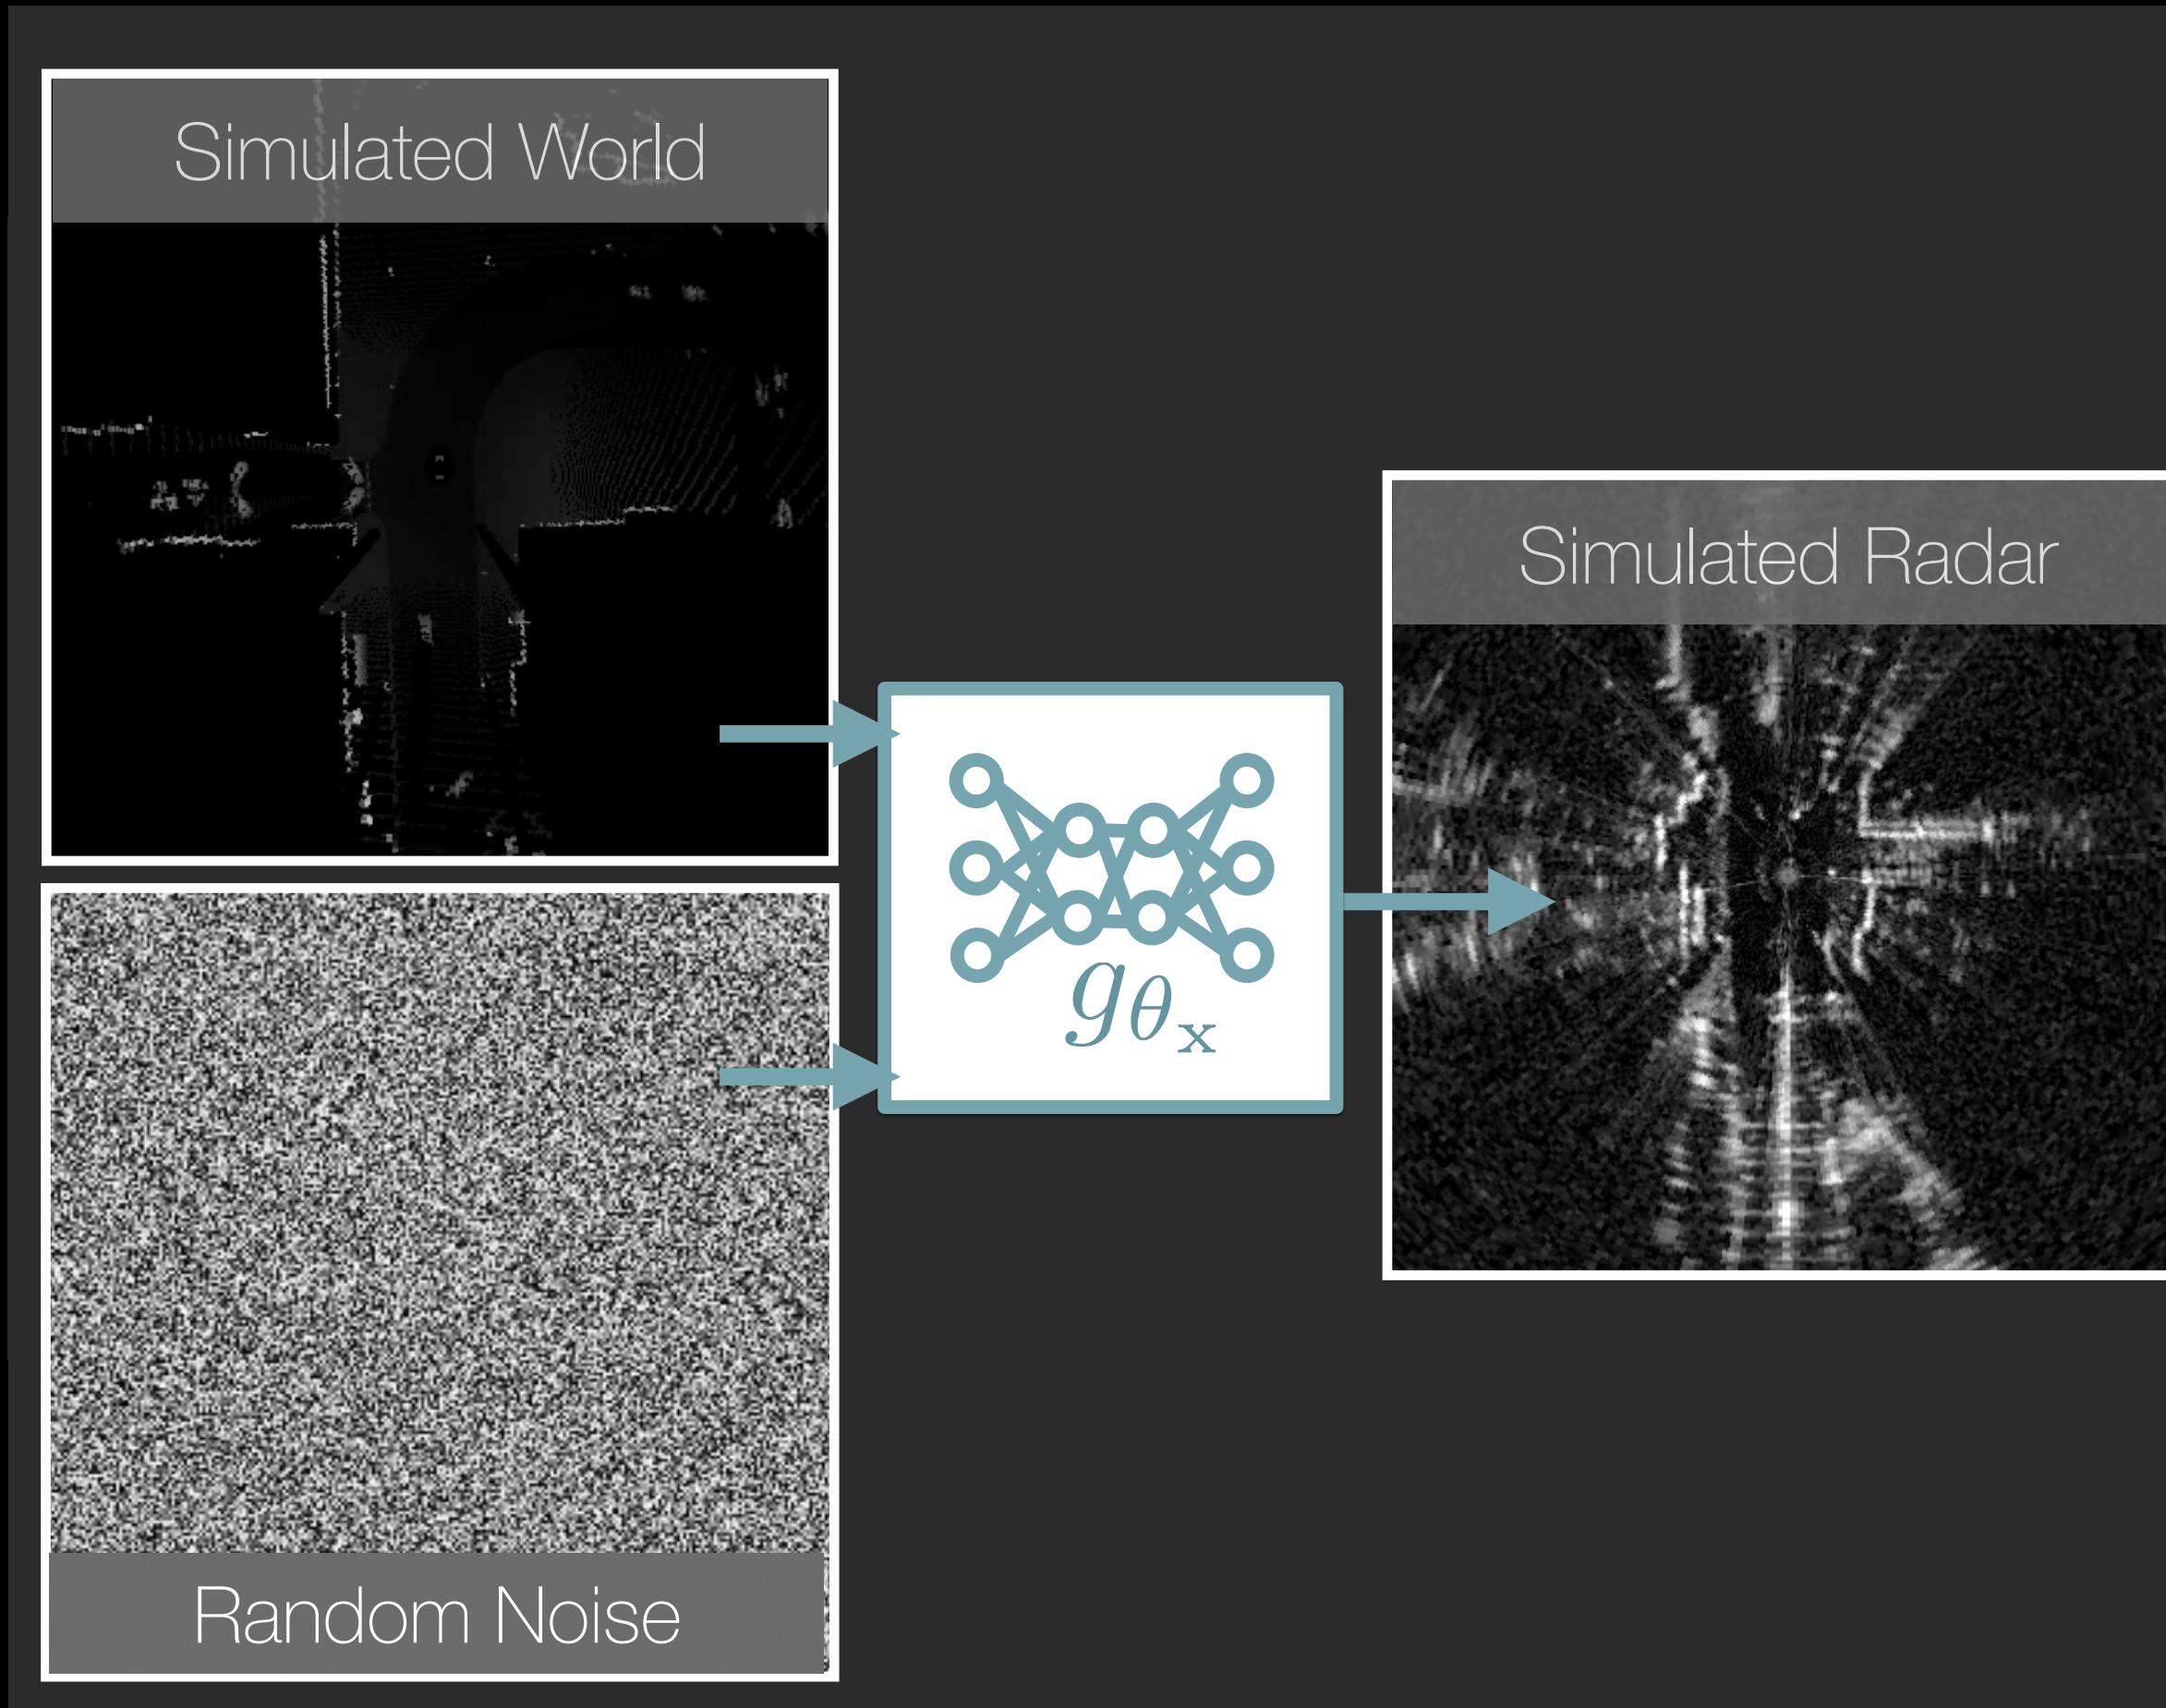
# Our Approach

Predict radar from **simulated states** using **adversarial** loss

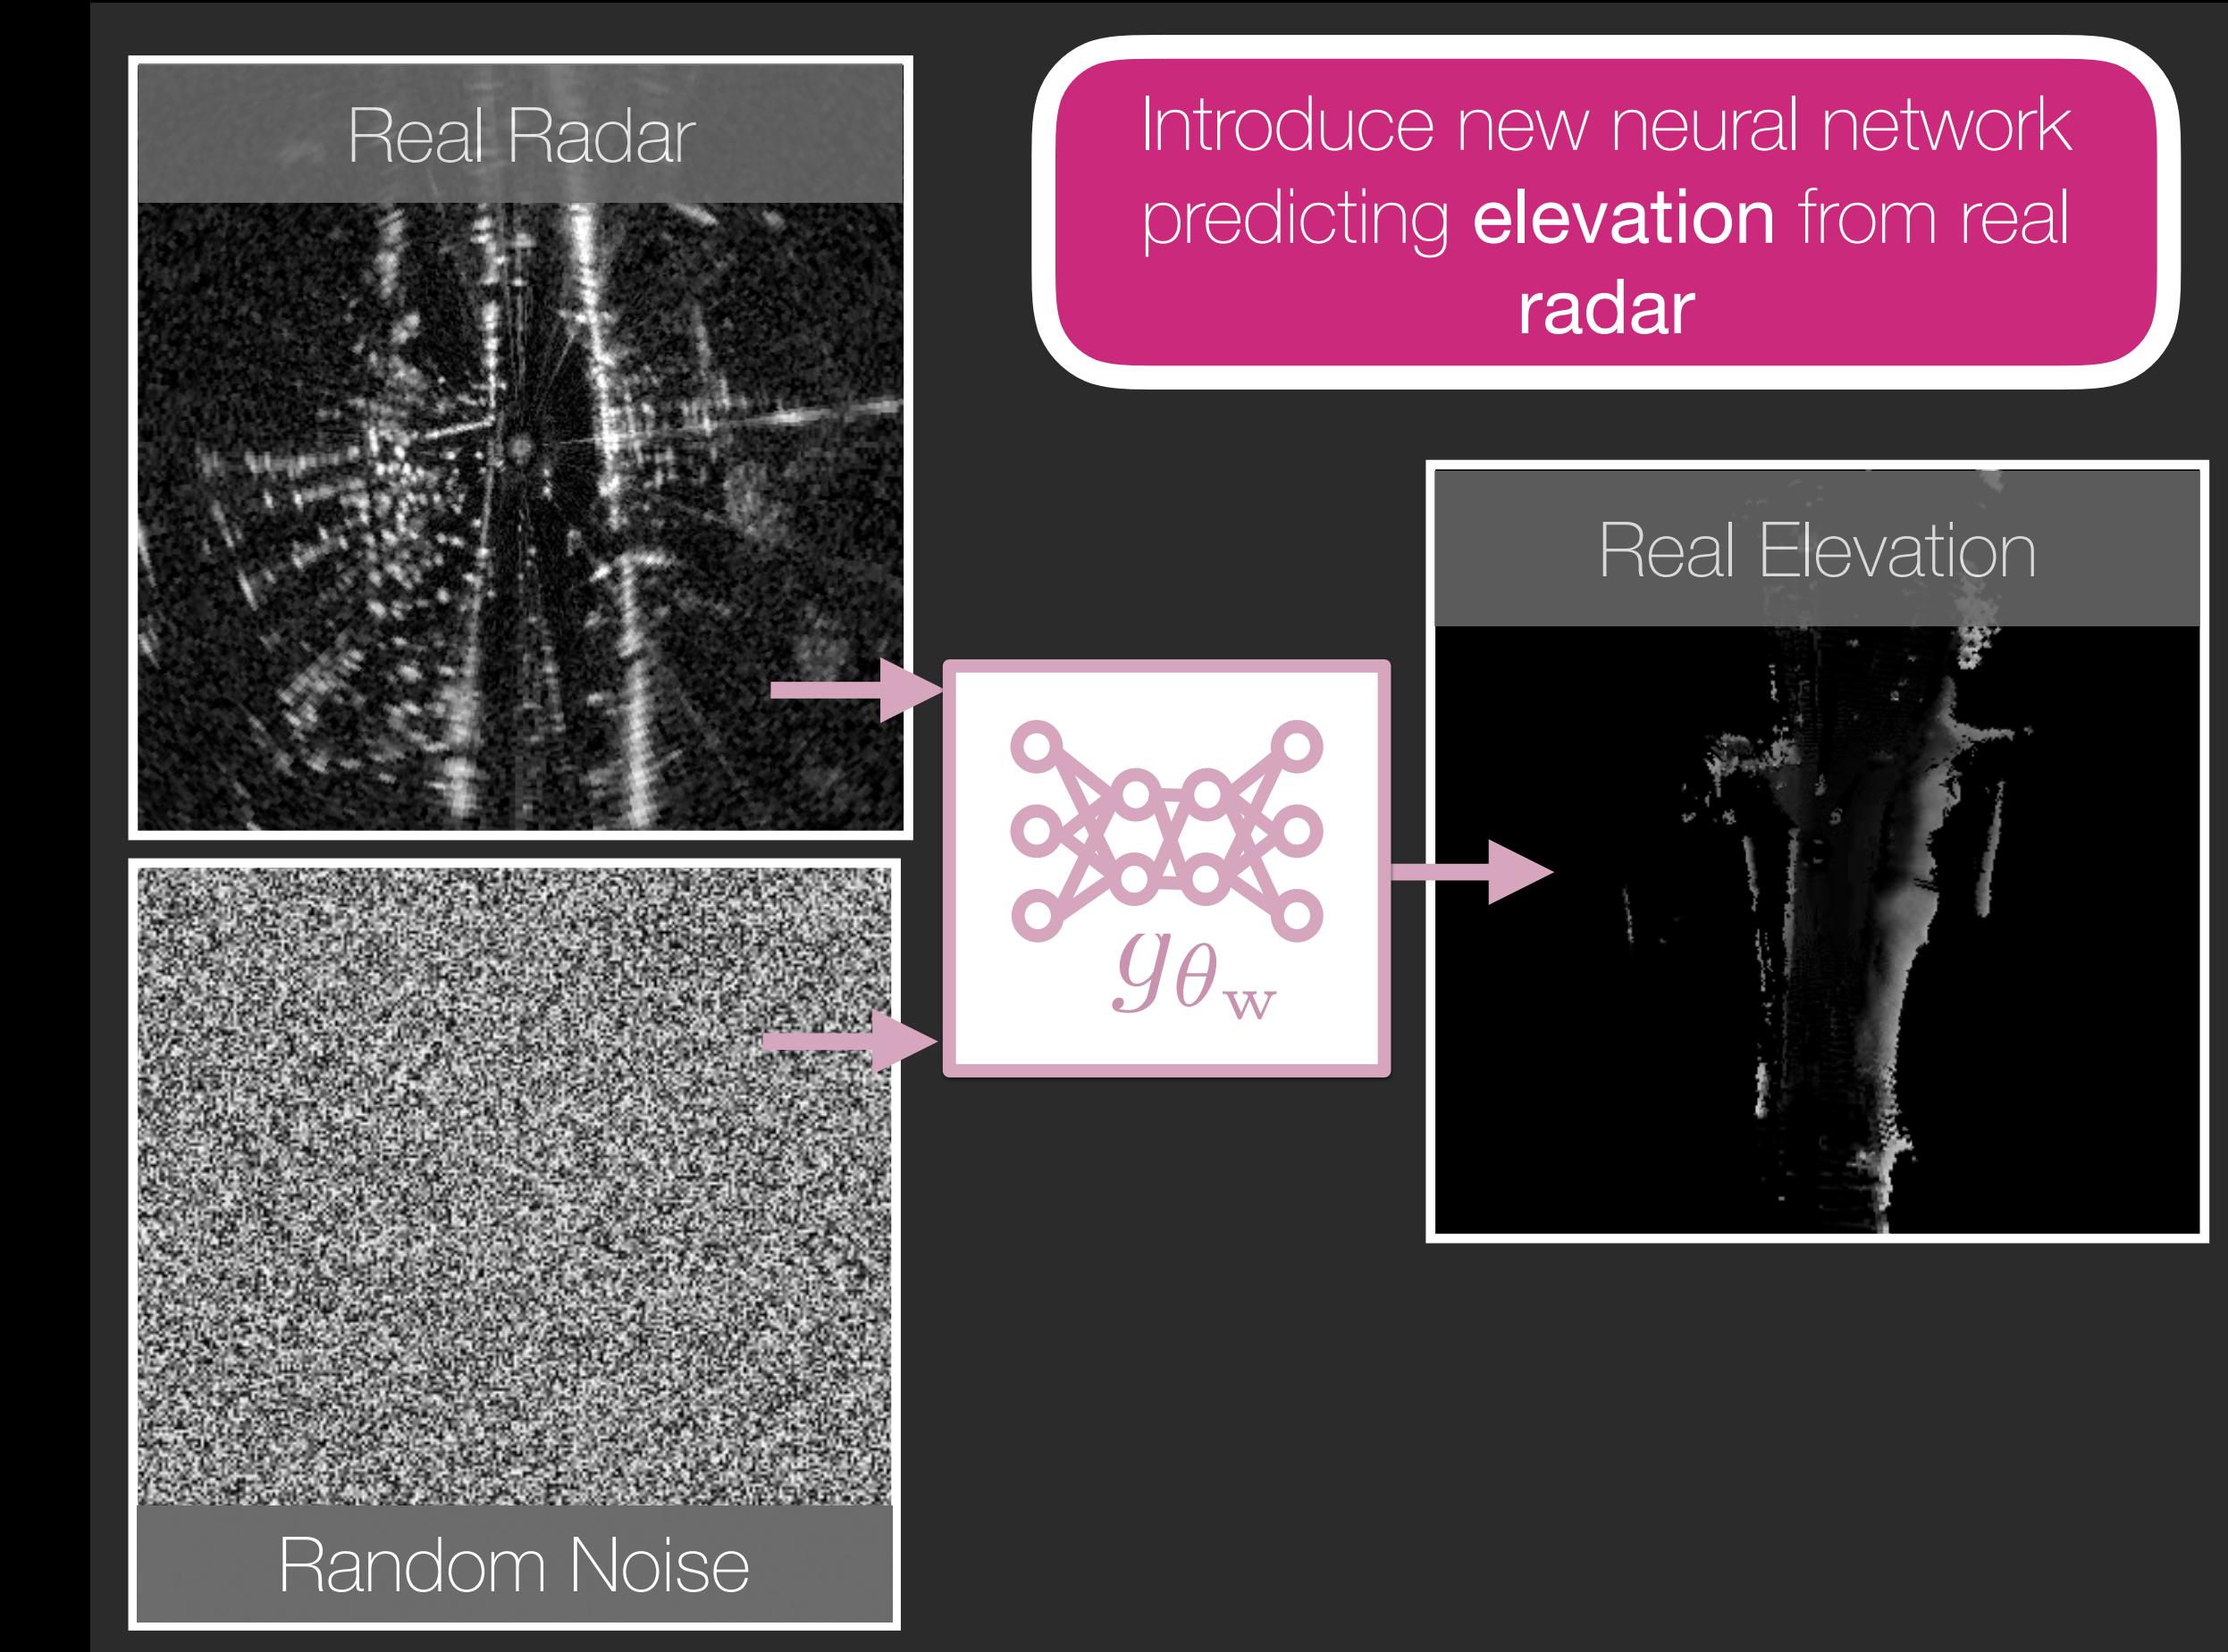


# Our Approach

Model both the **forward** and **backward** processes



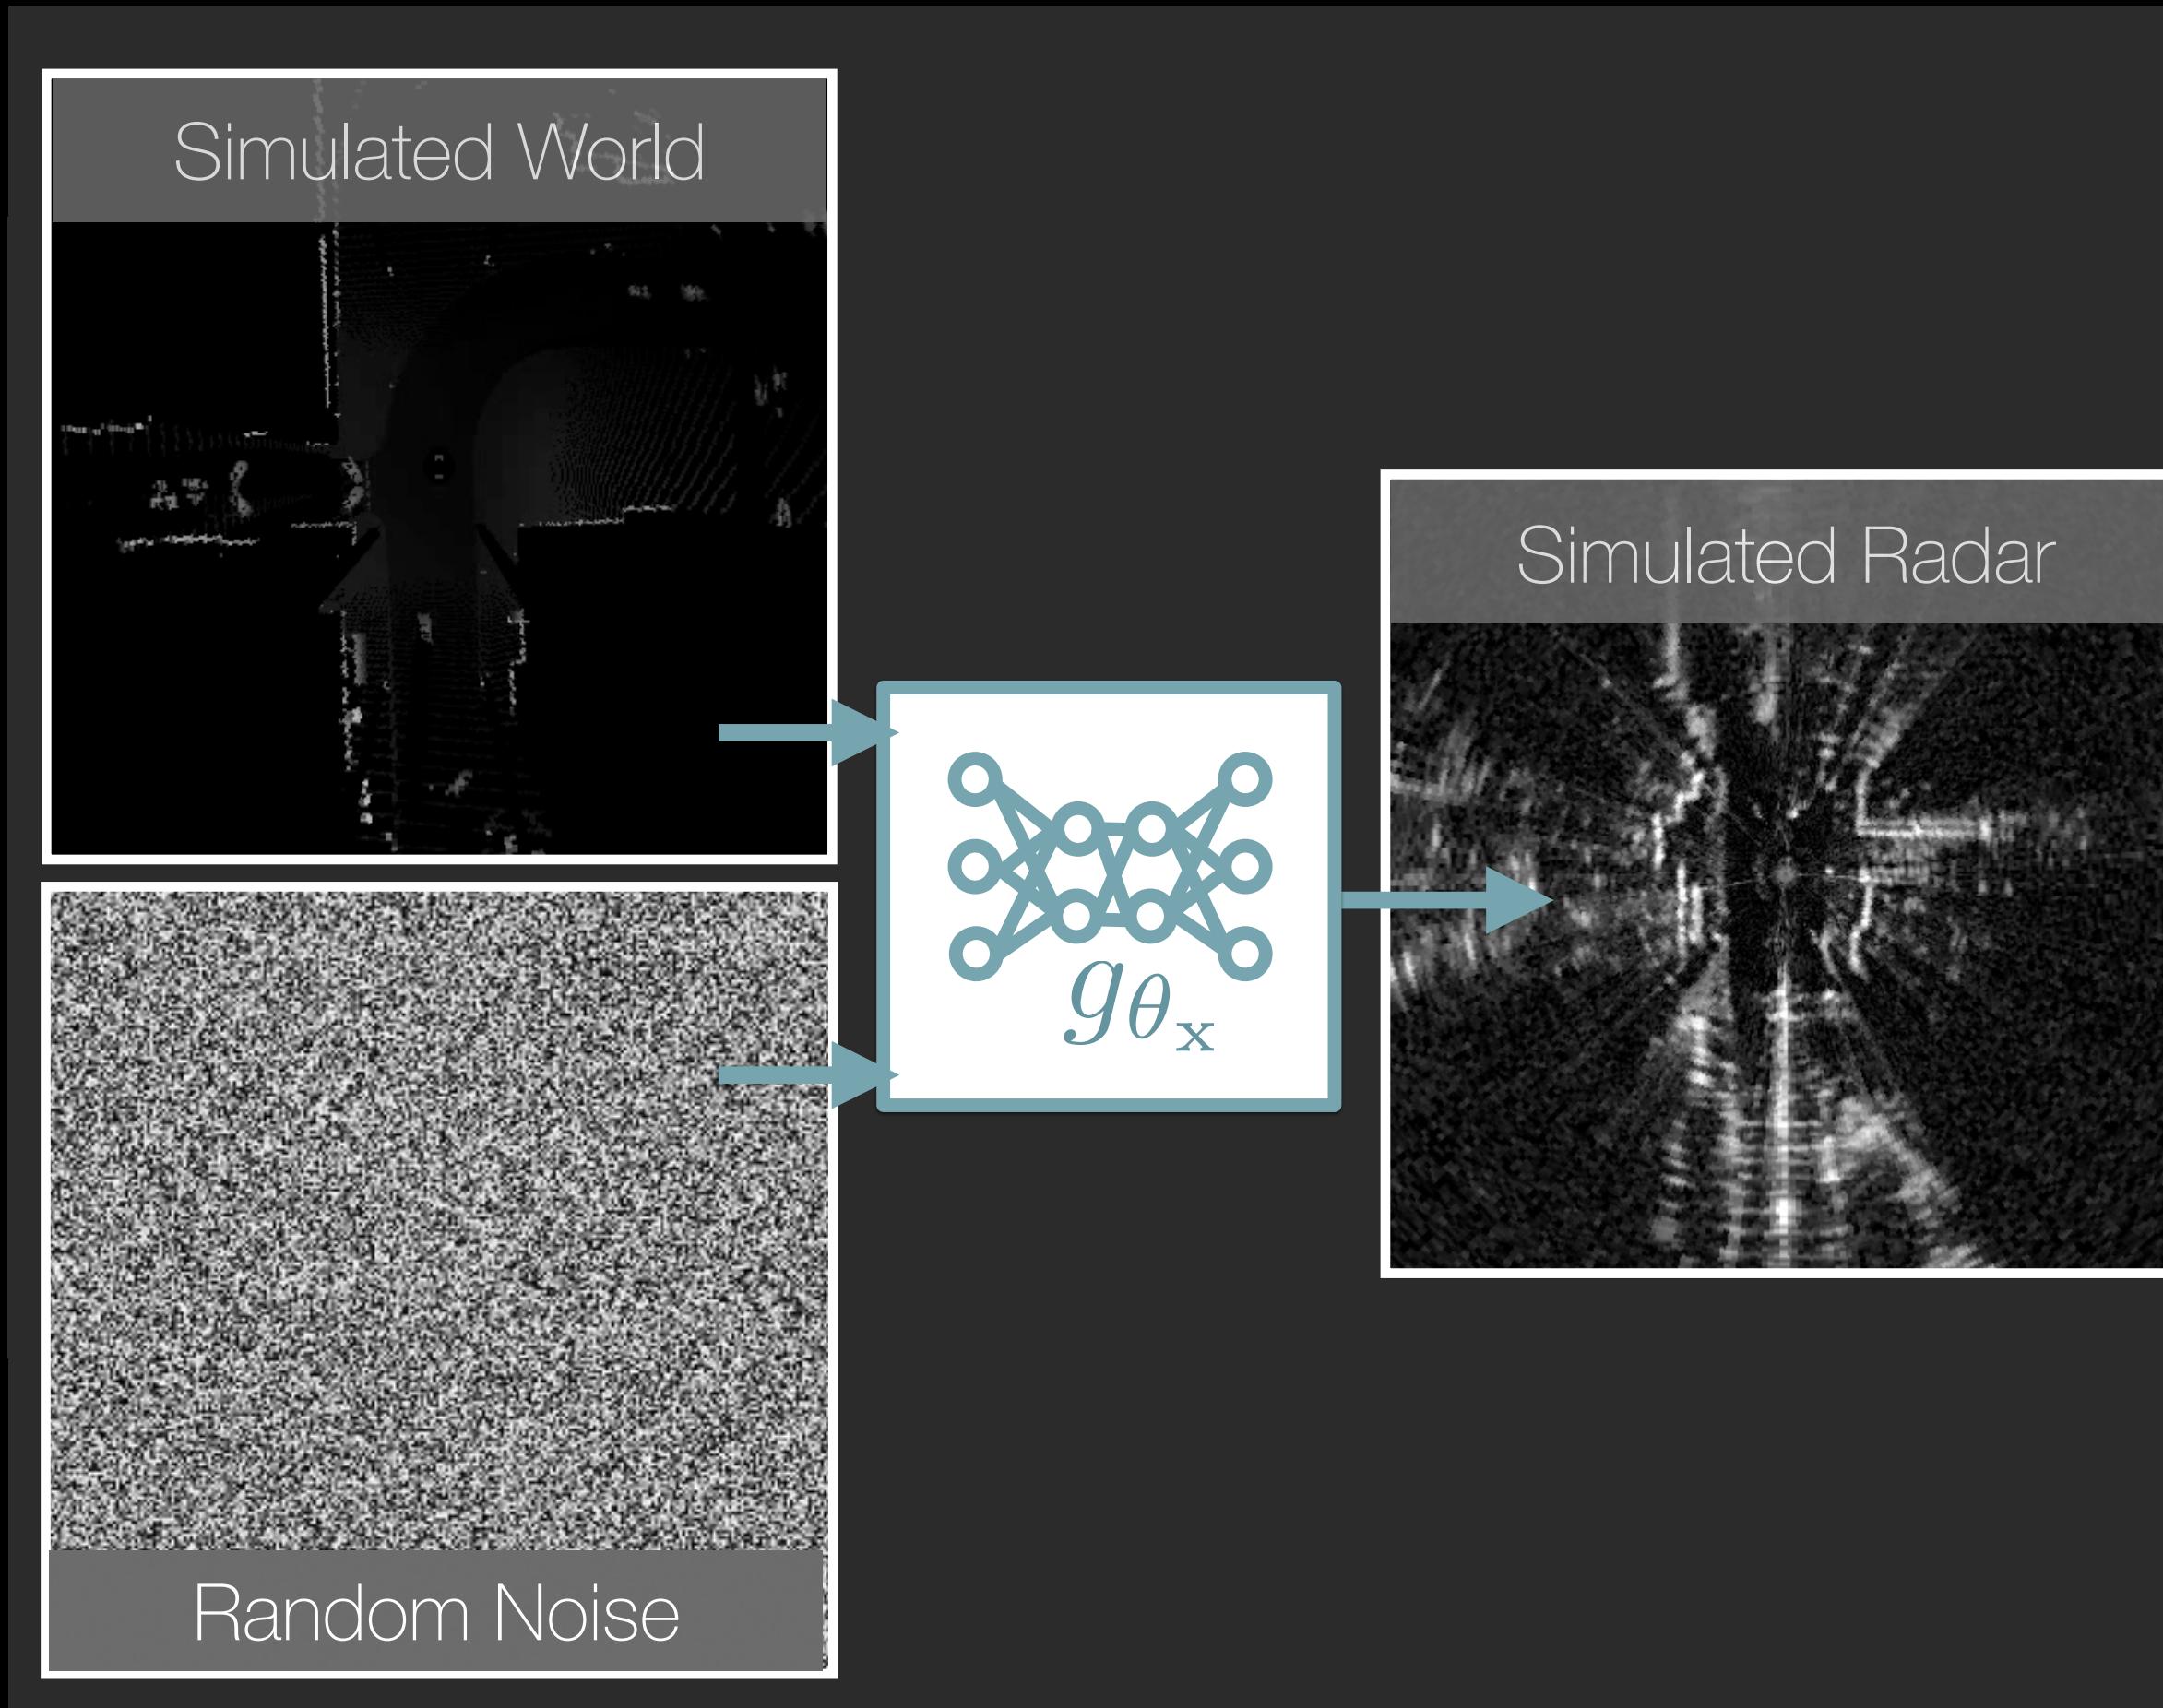
Forward



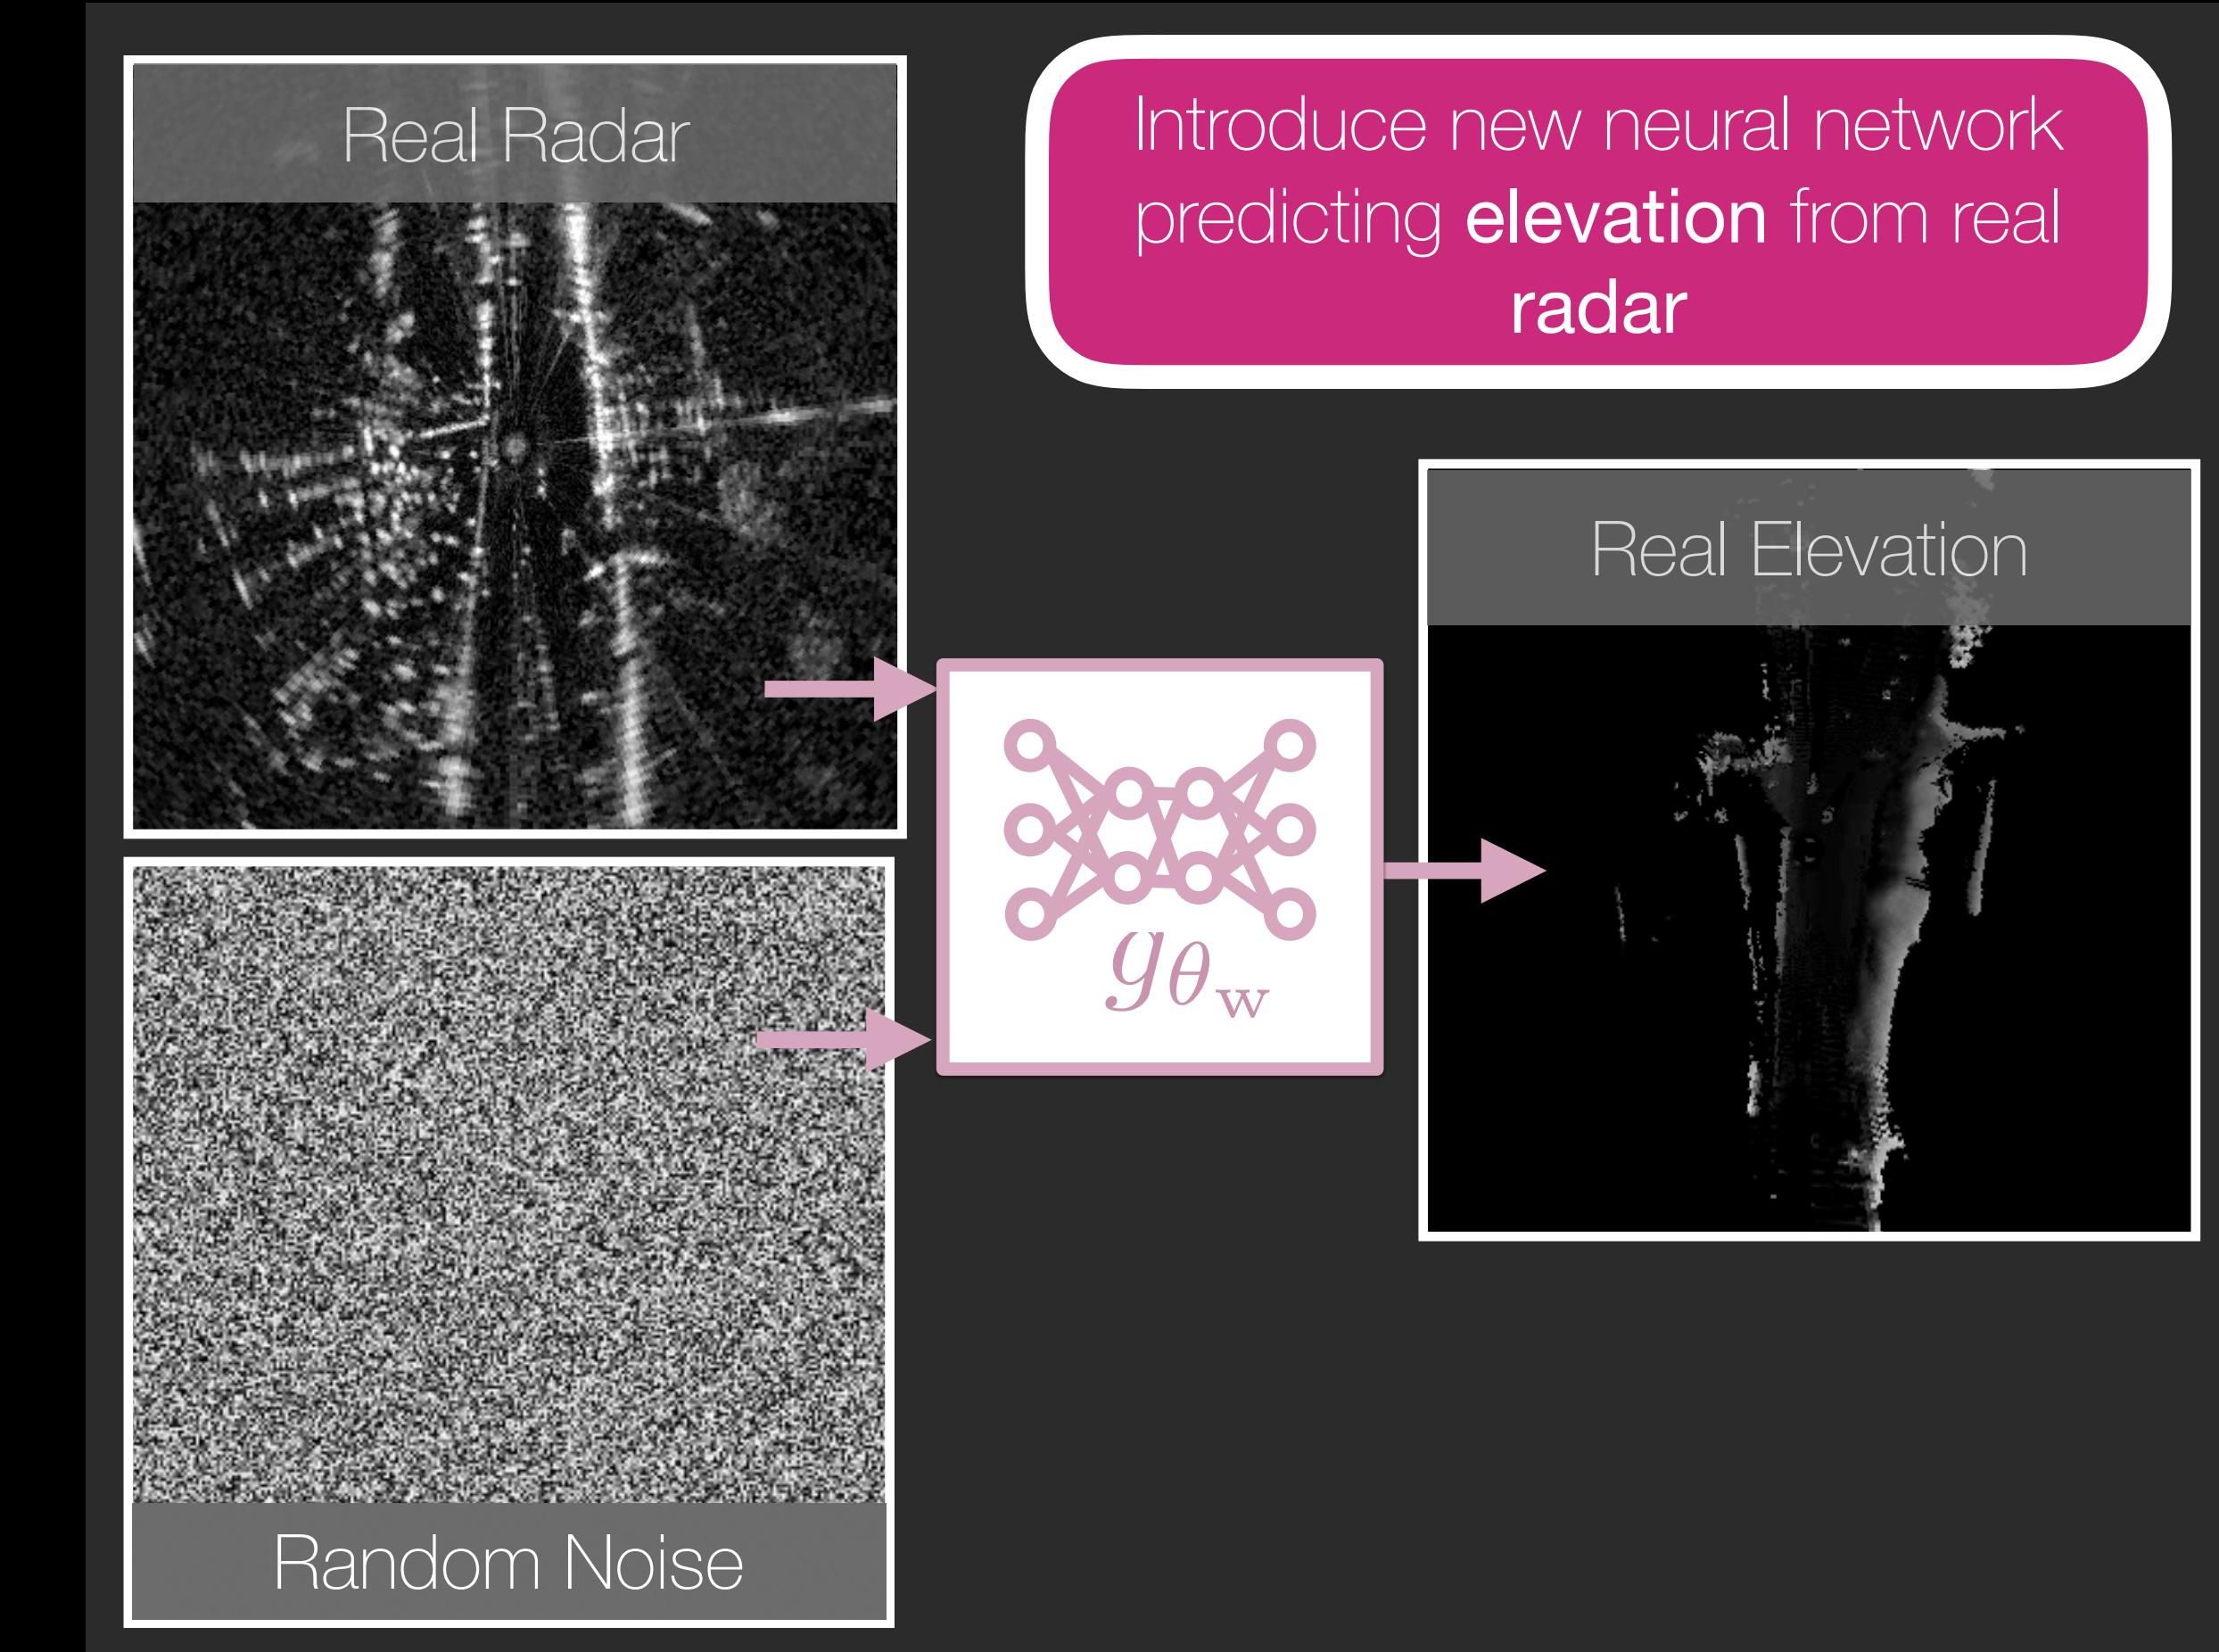
Backward

# Our Approach

Model both the **forward** and **backward** processes



Forward

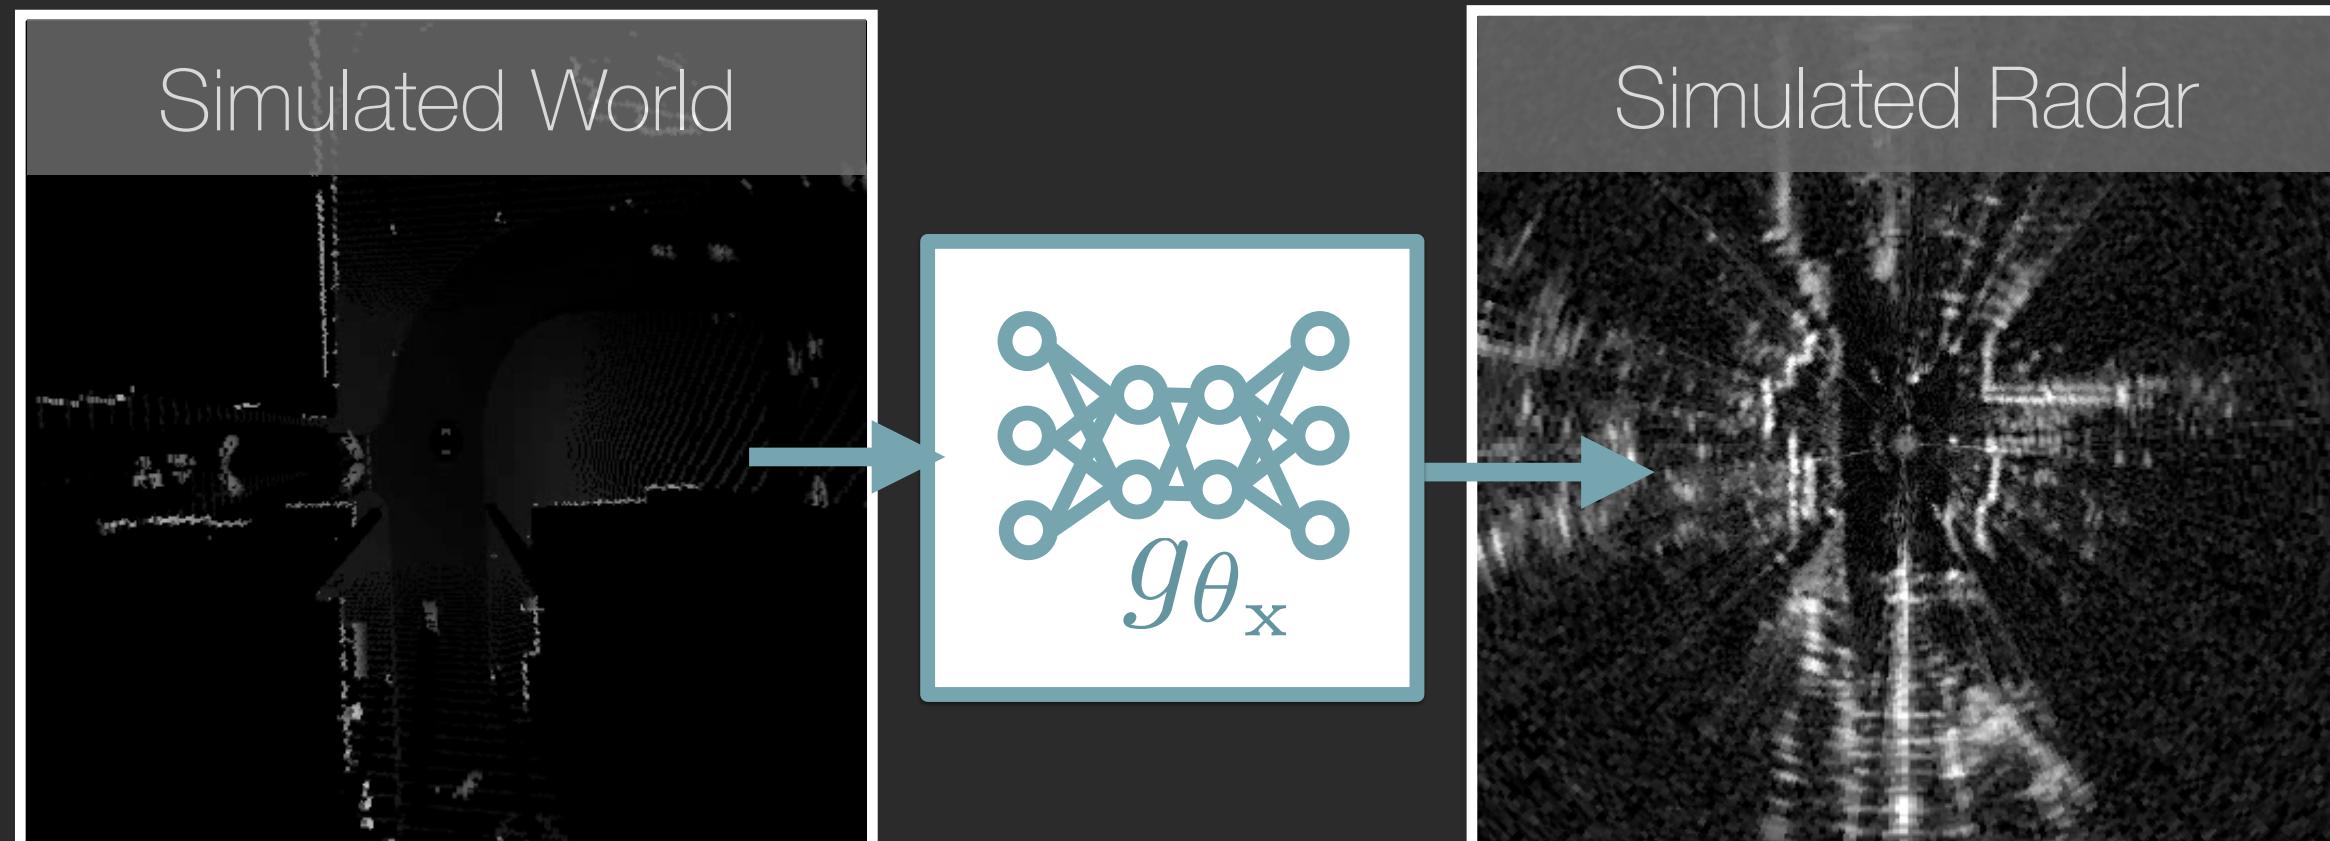


Backward

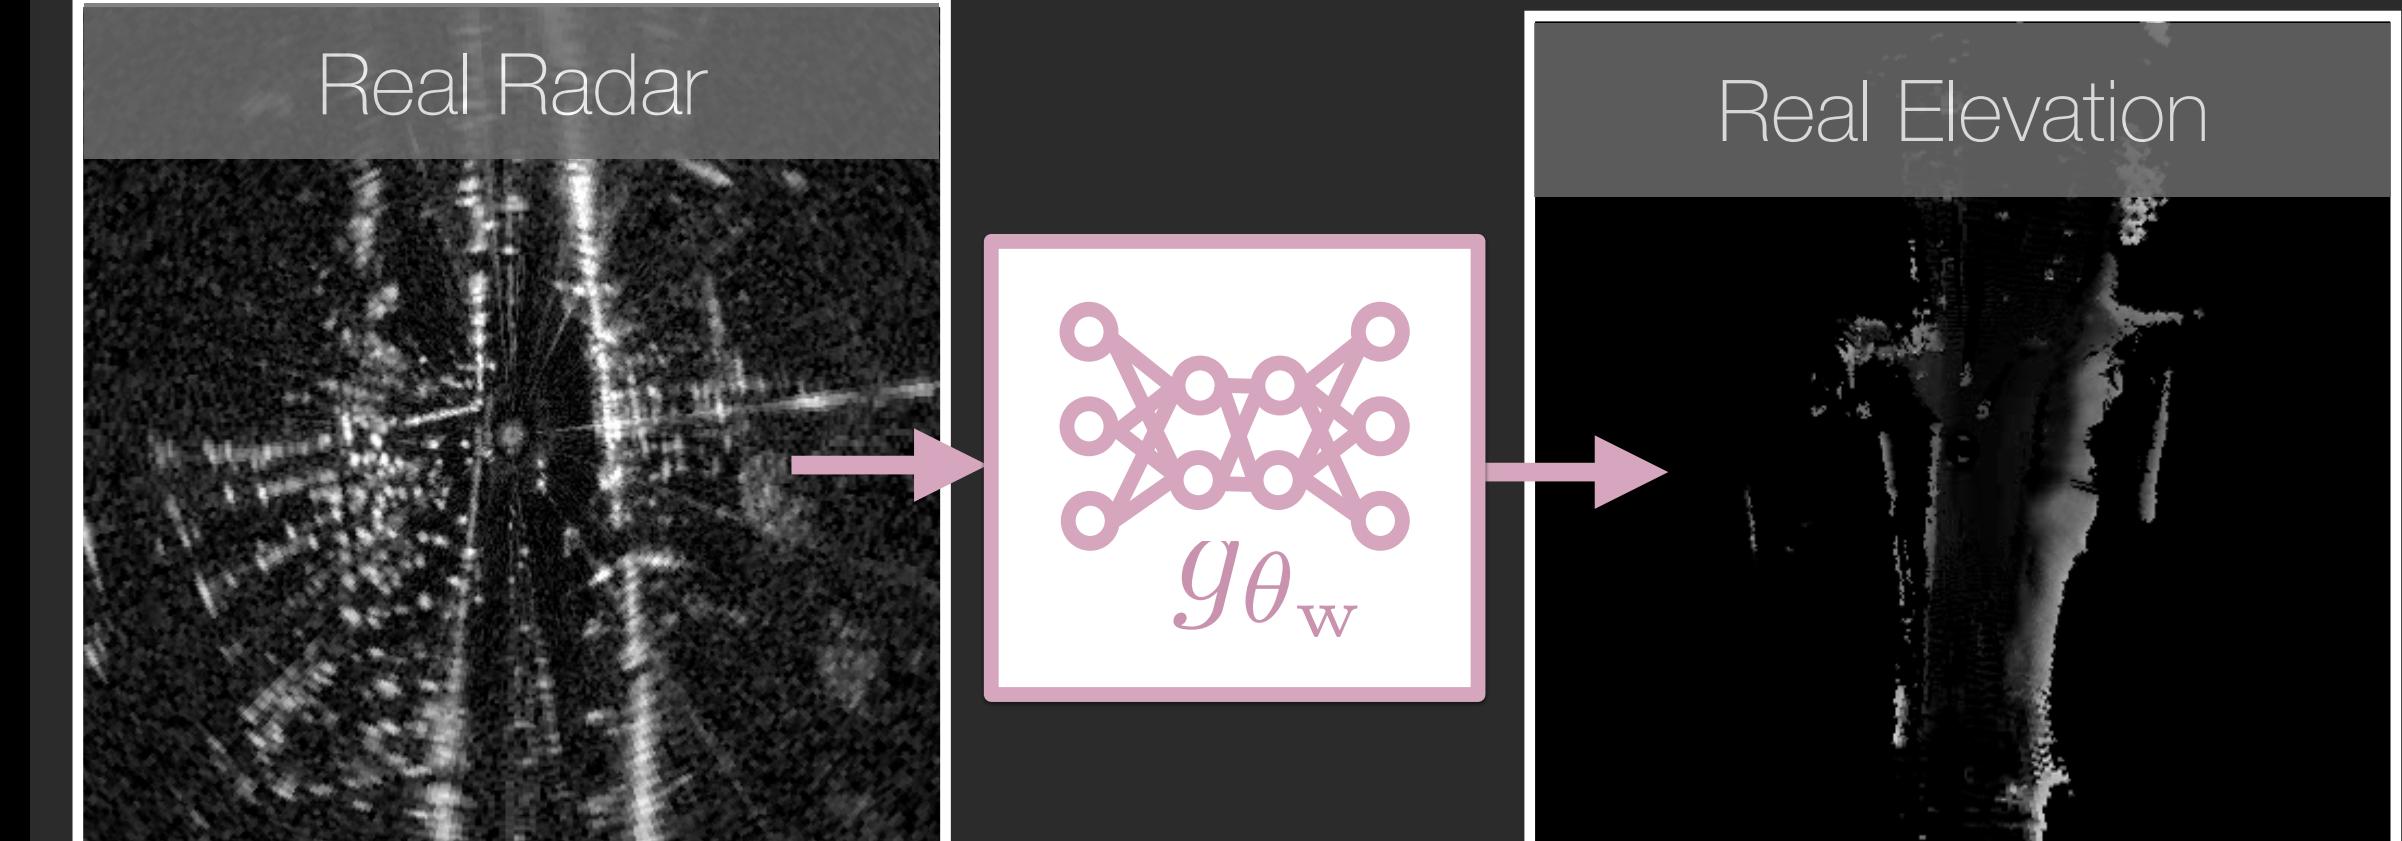
# Our Approach

Model both the **forward** and **backward** processes

**Forward:** radar from simulated world state

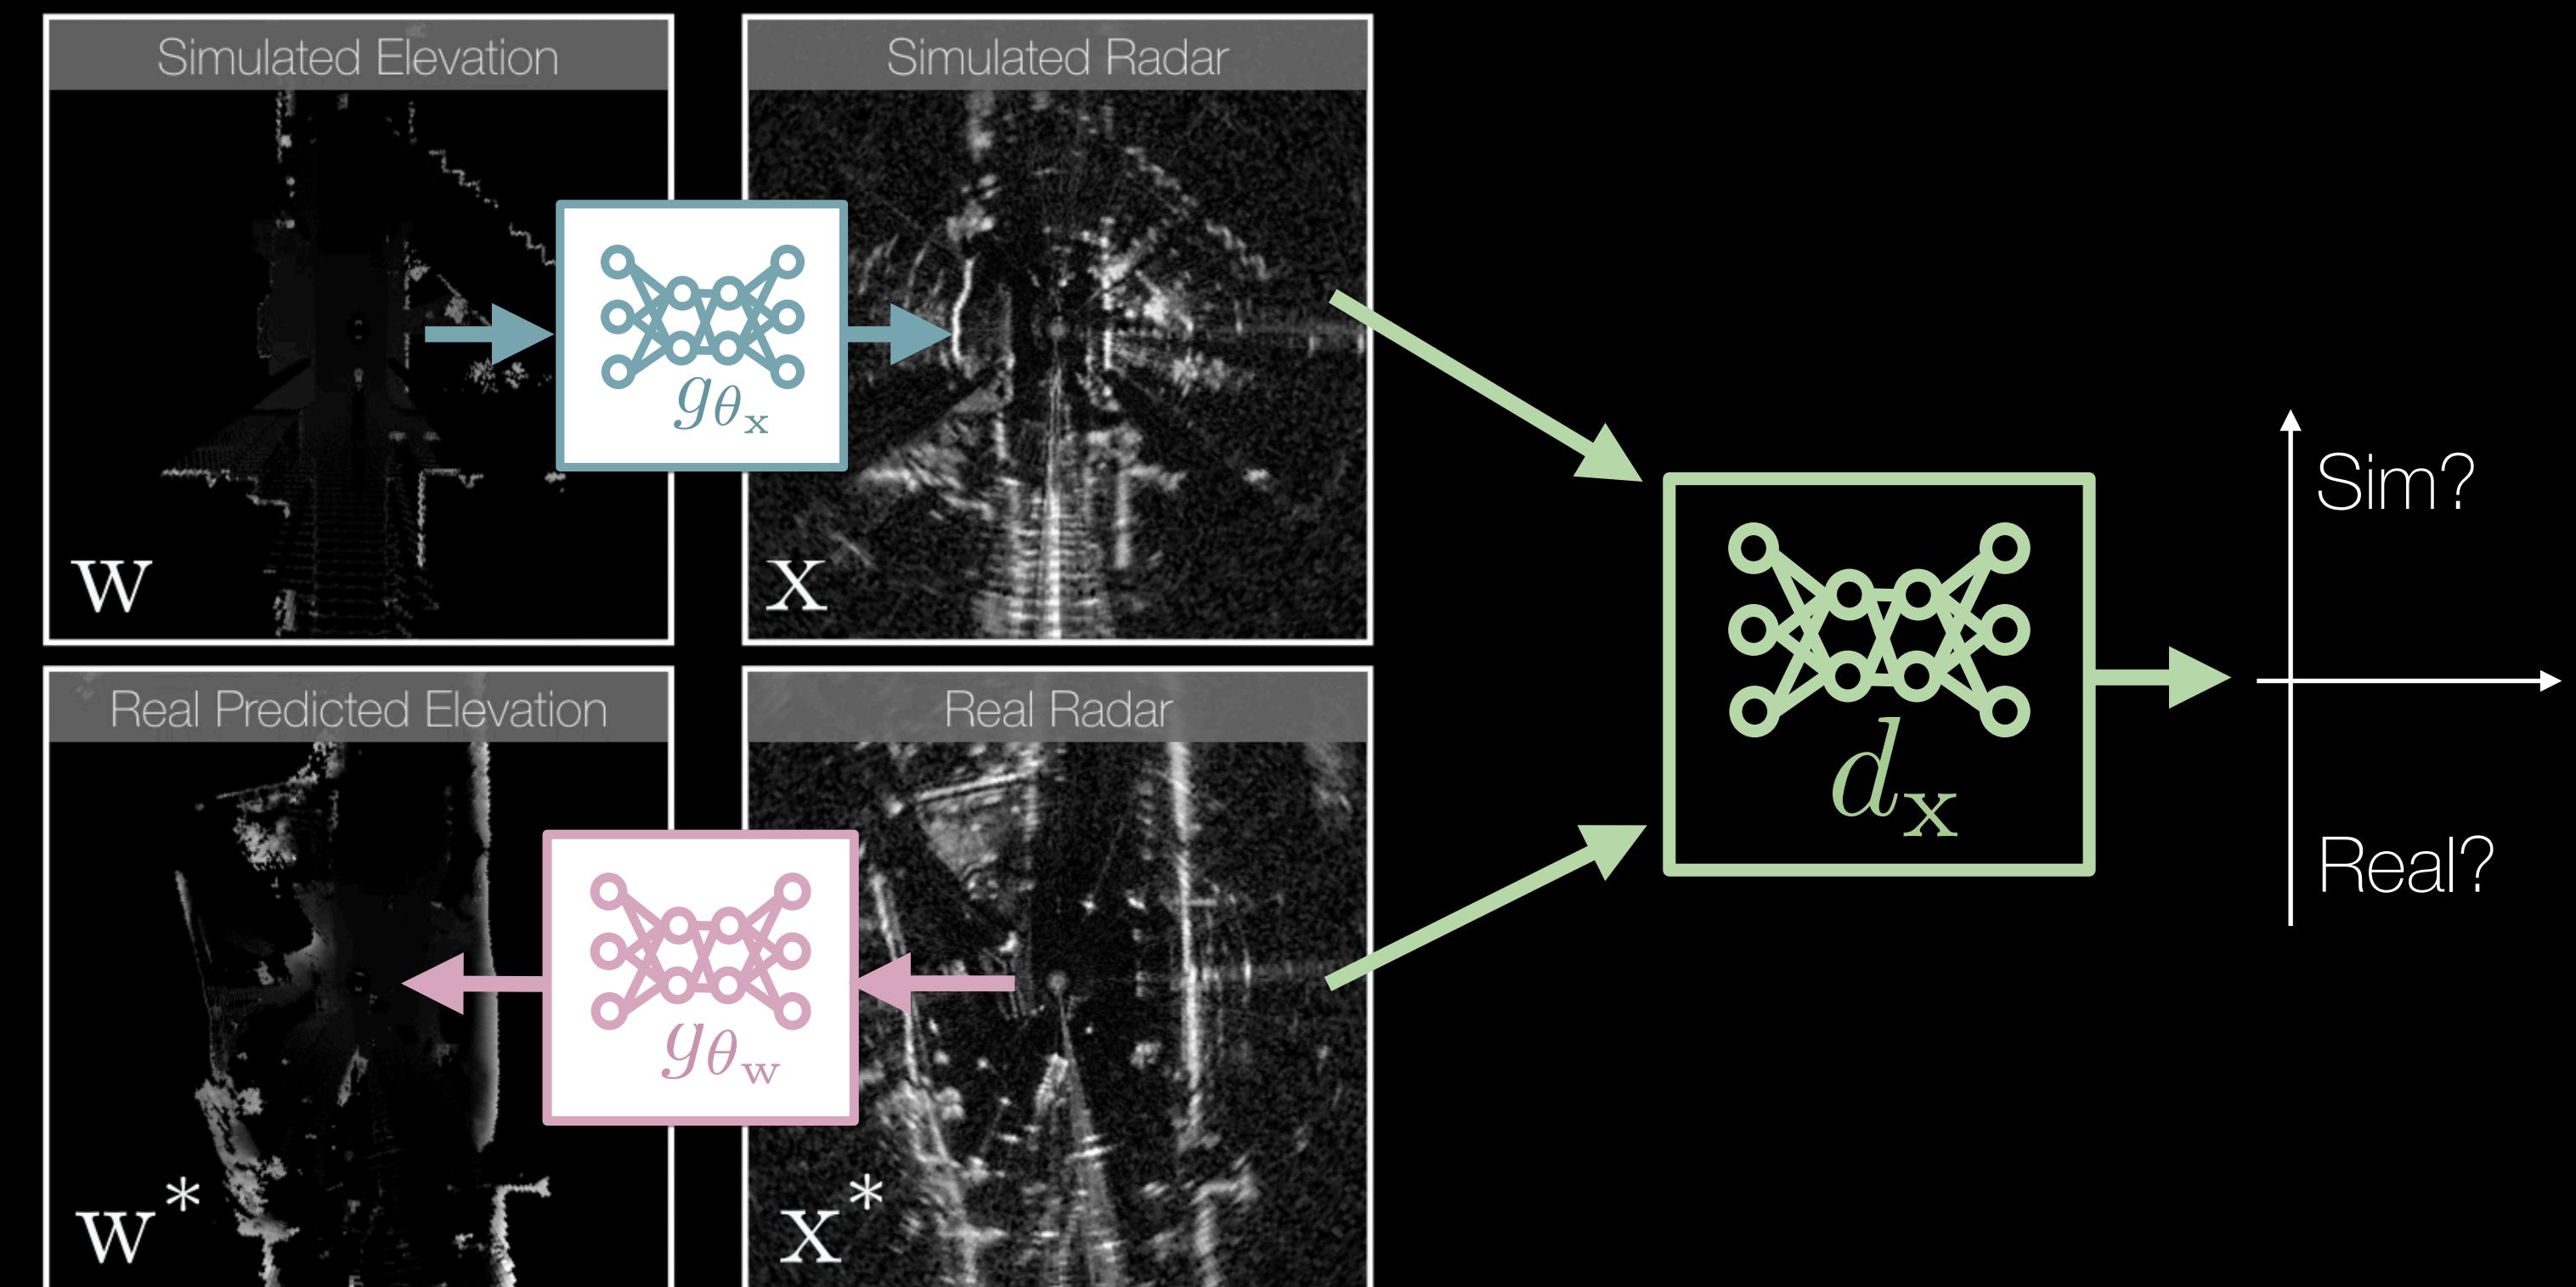


**Backward:** elevation from real radar



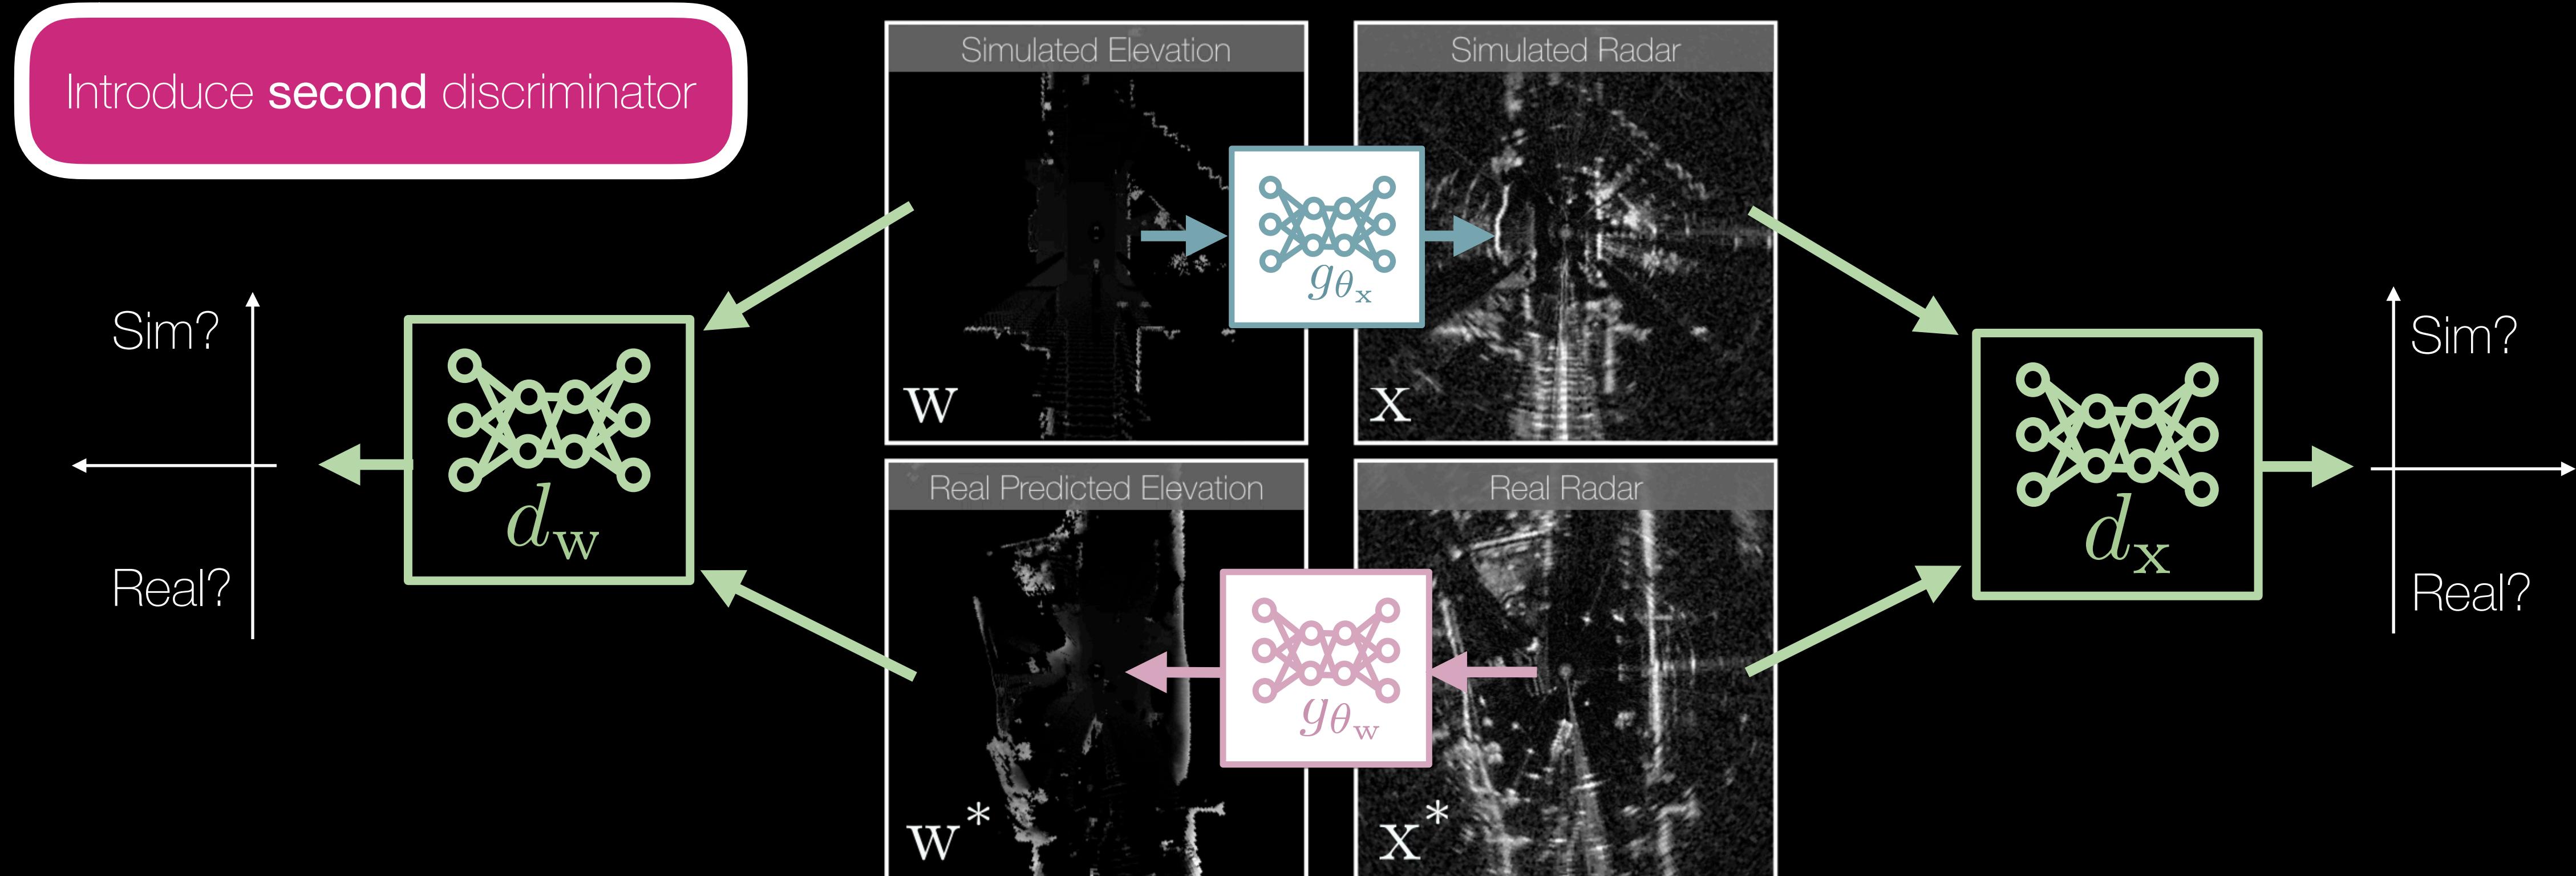
# Our Approach

Model both the **forward** and **backward** processes



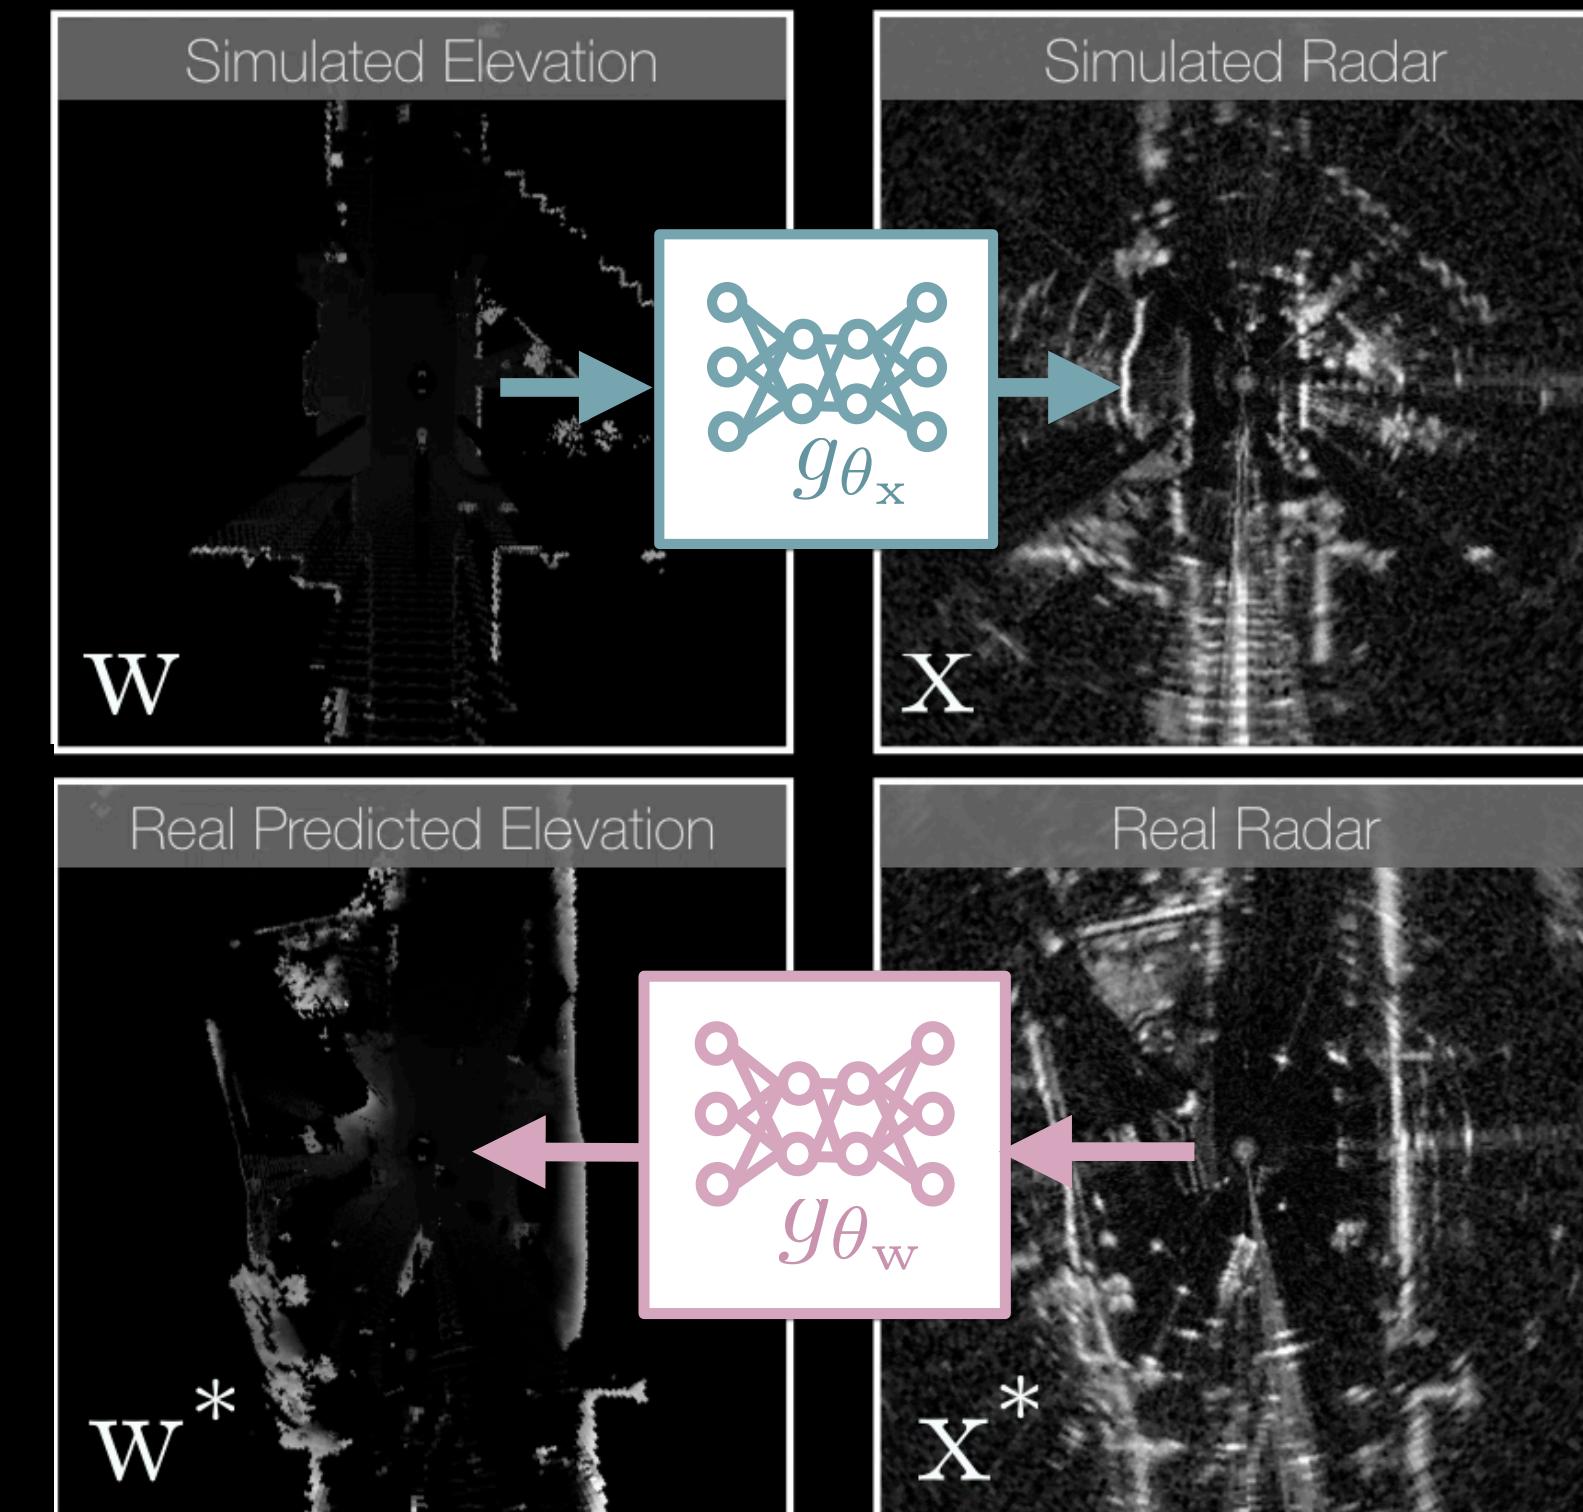
# Our Approach

Model both the **forward** and **backward** processes



# Our Approach

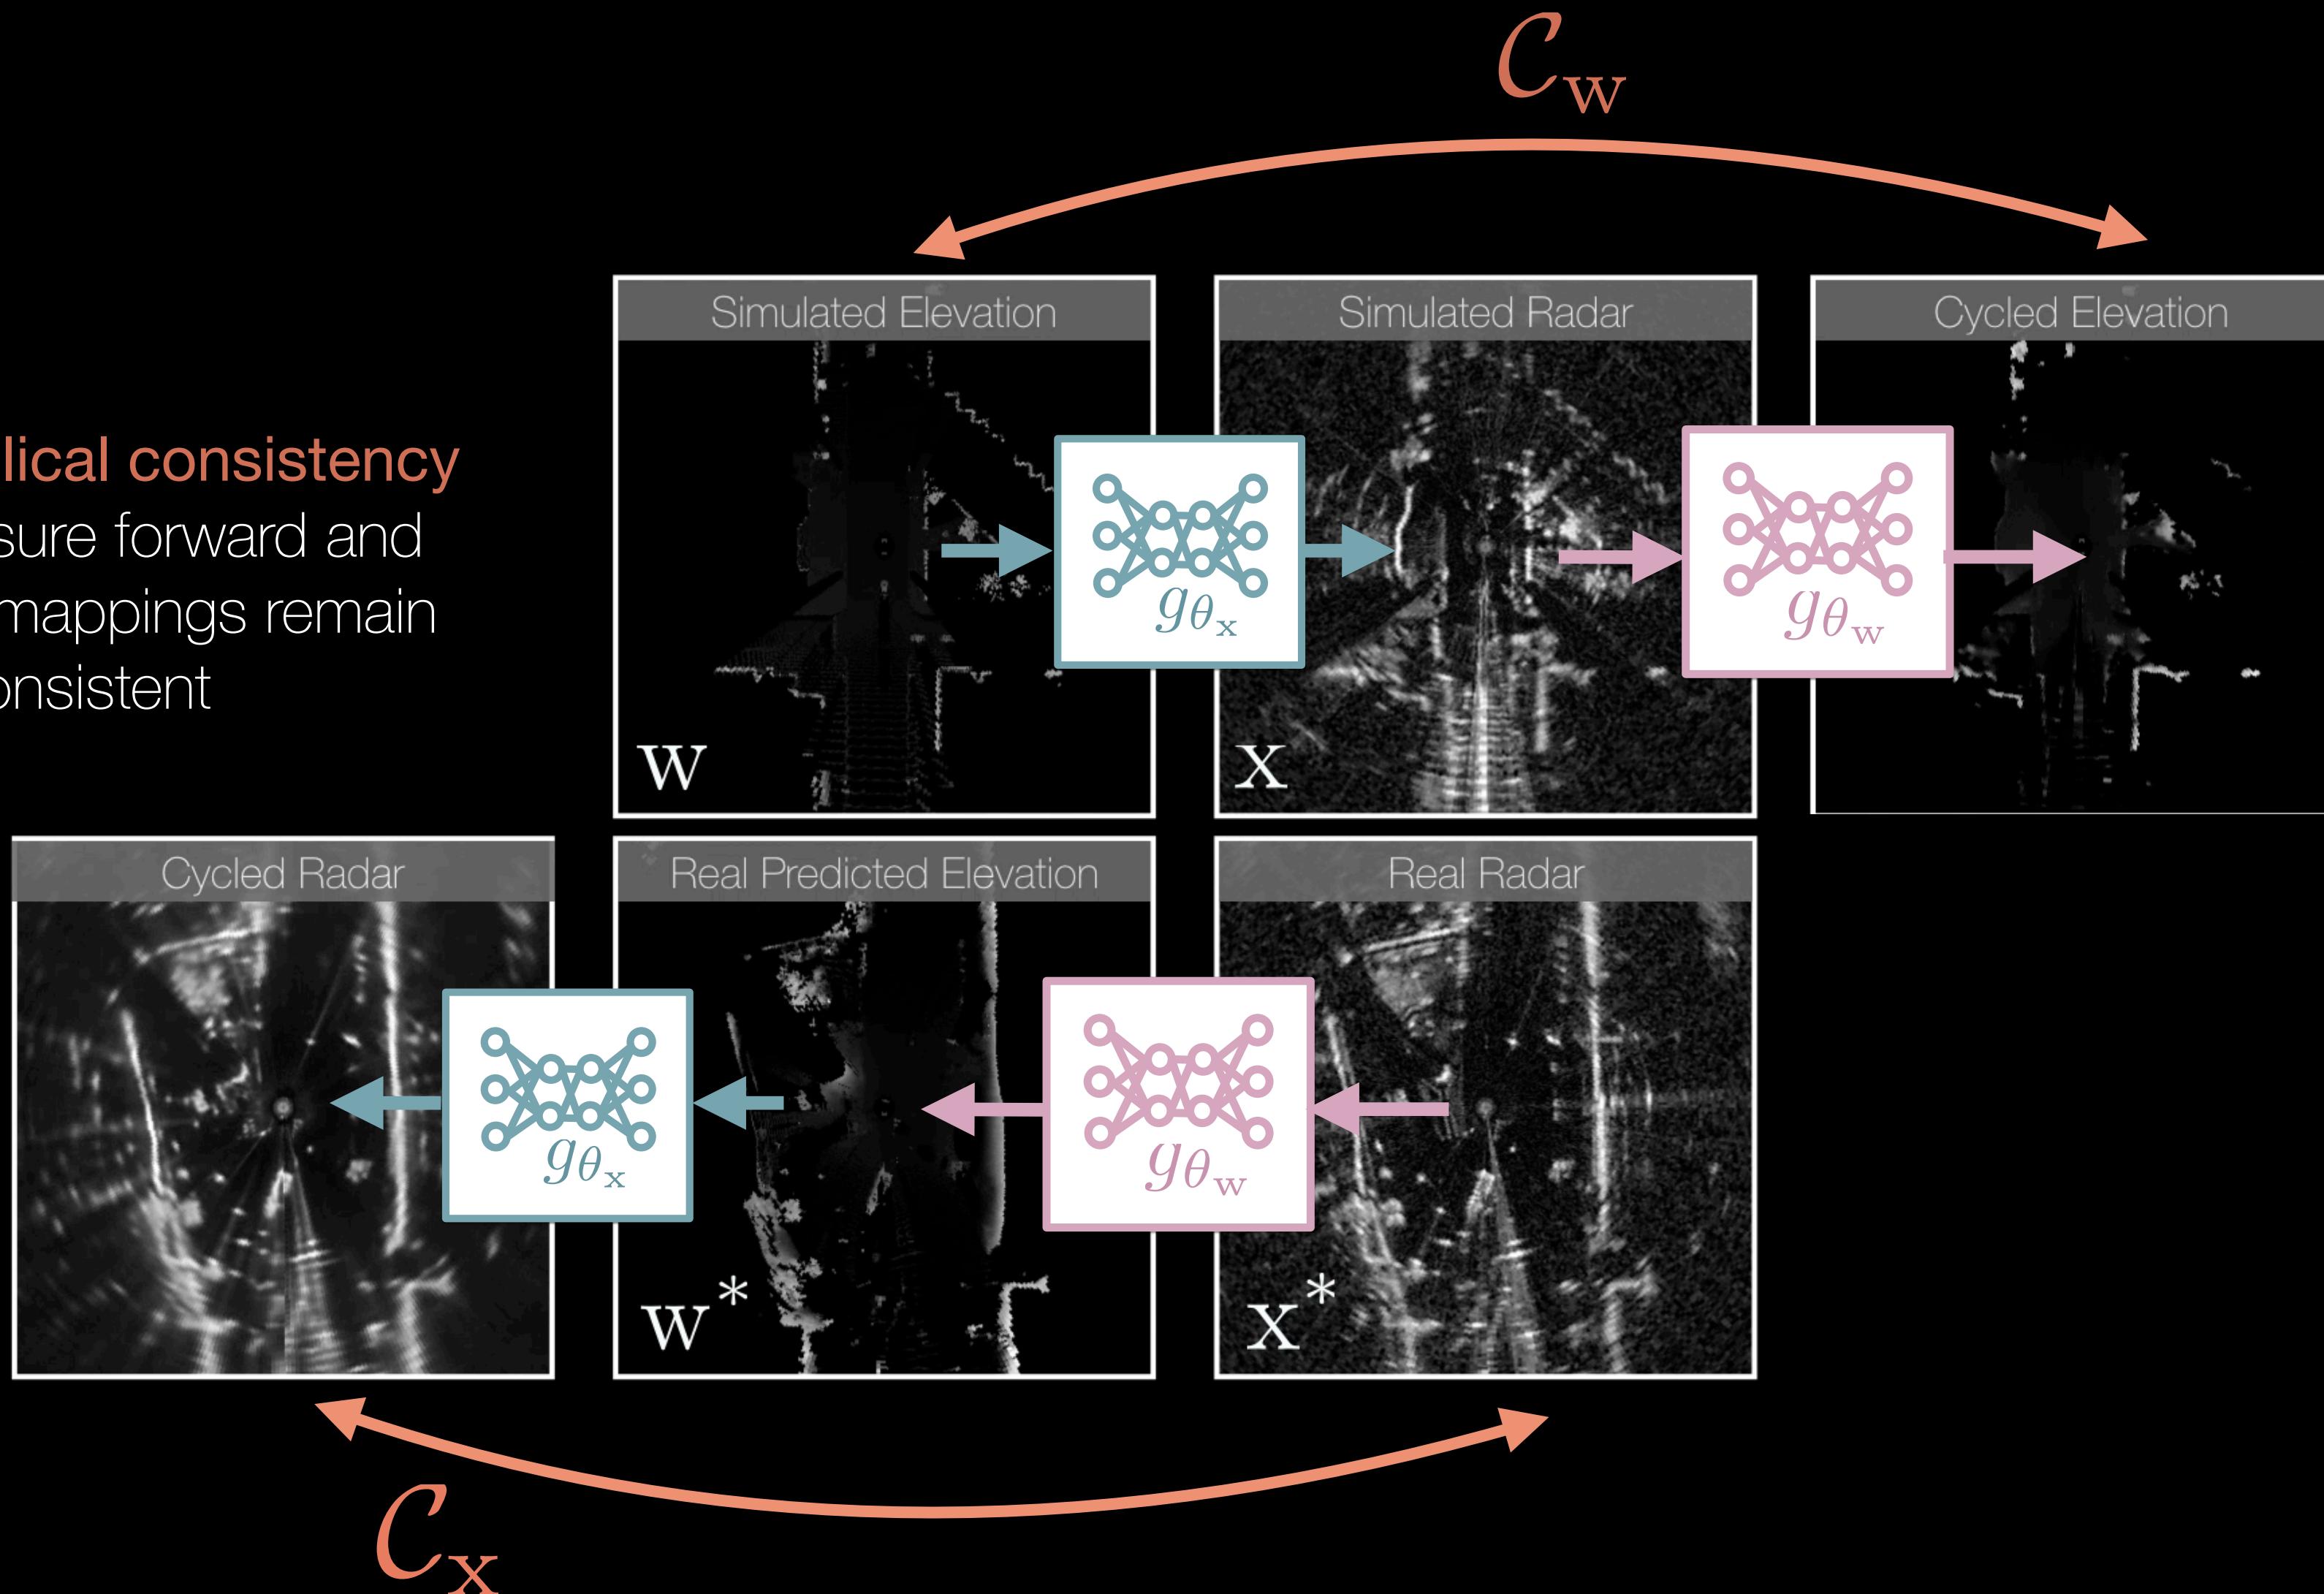
Further constrain training with **cyclical consistency** and **real world alignment**



# Our Approach

Further constrain training with **cyclical consistency** and **real world alignment**

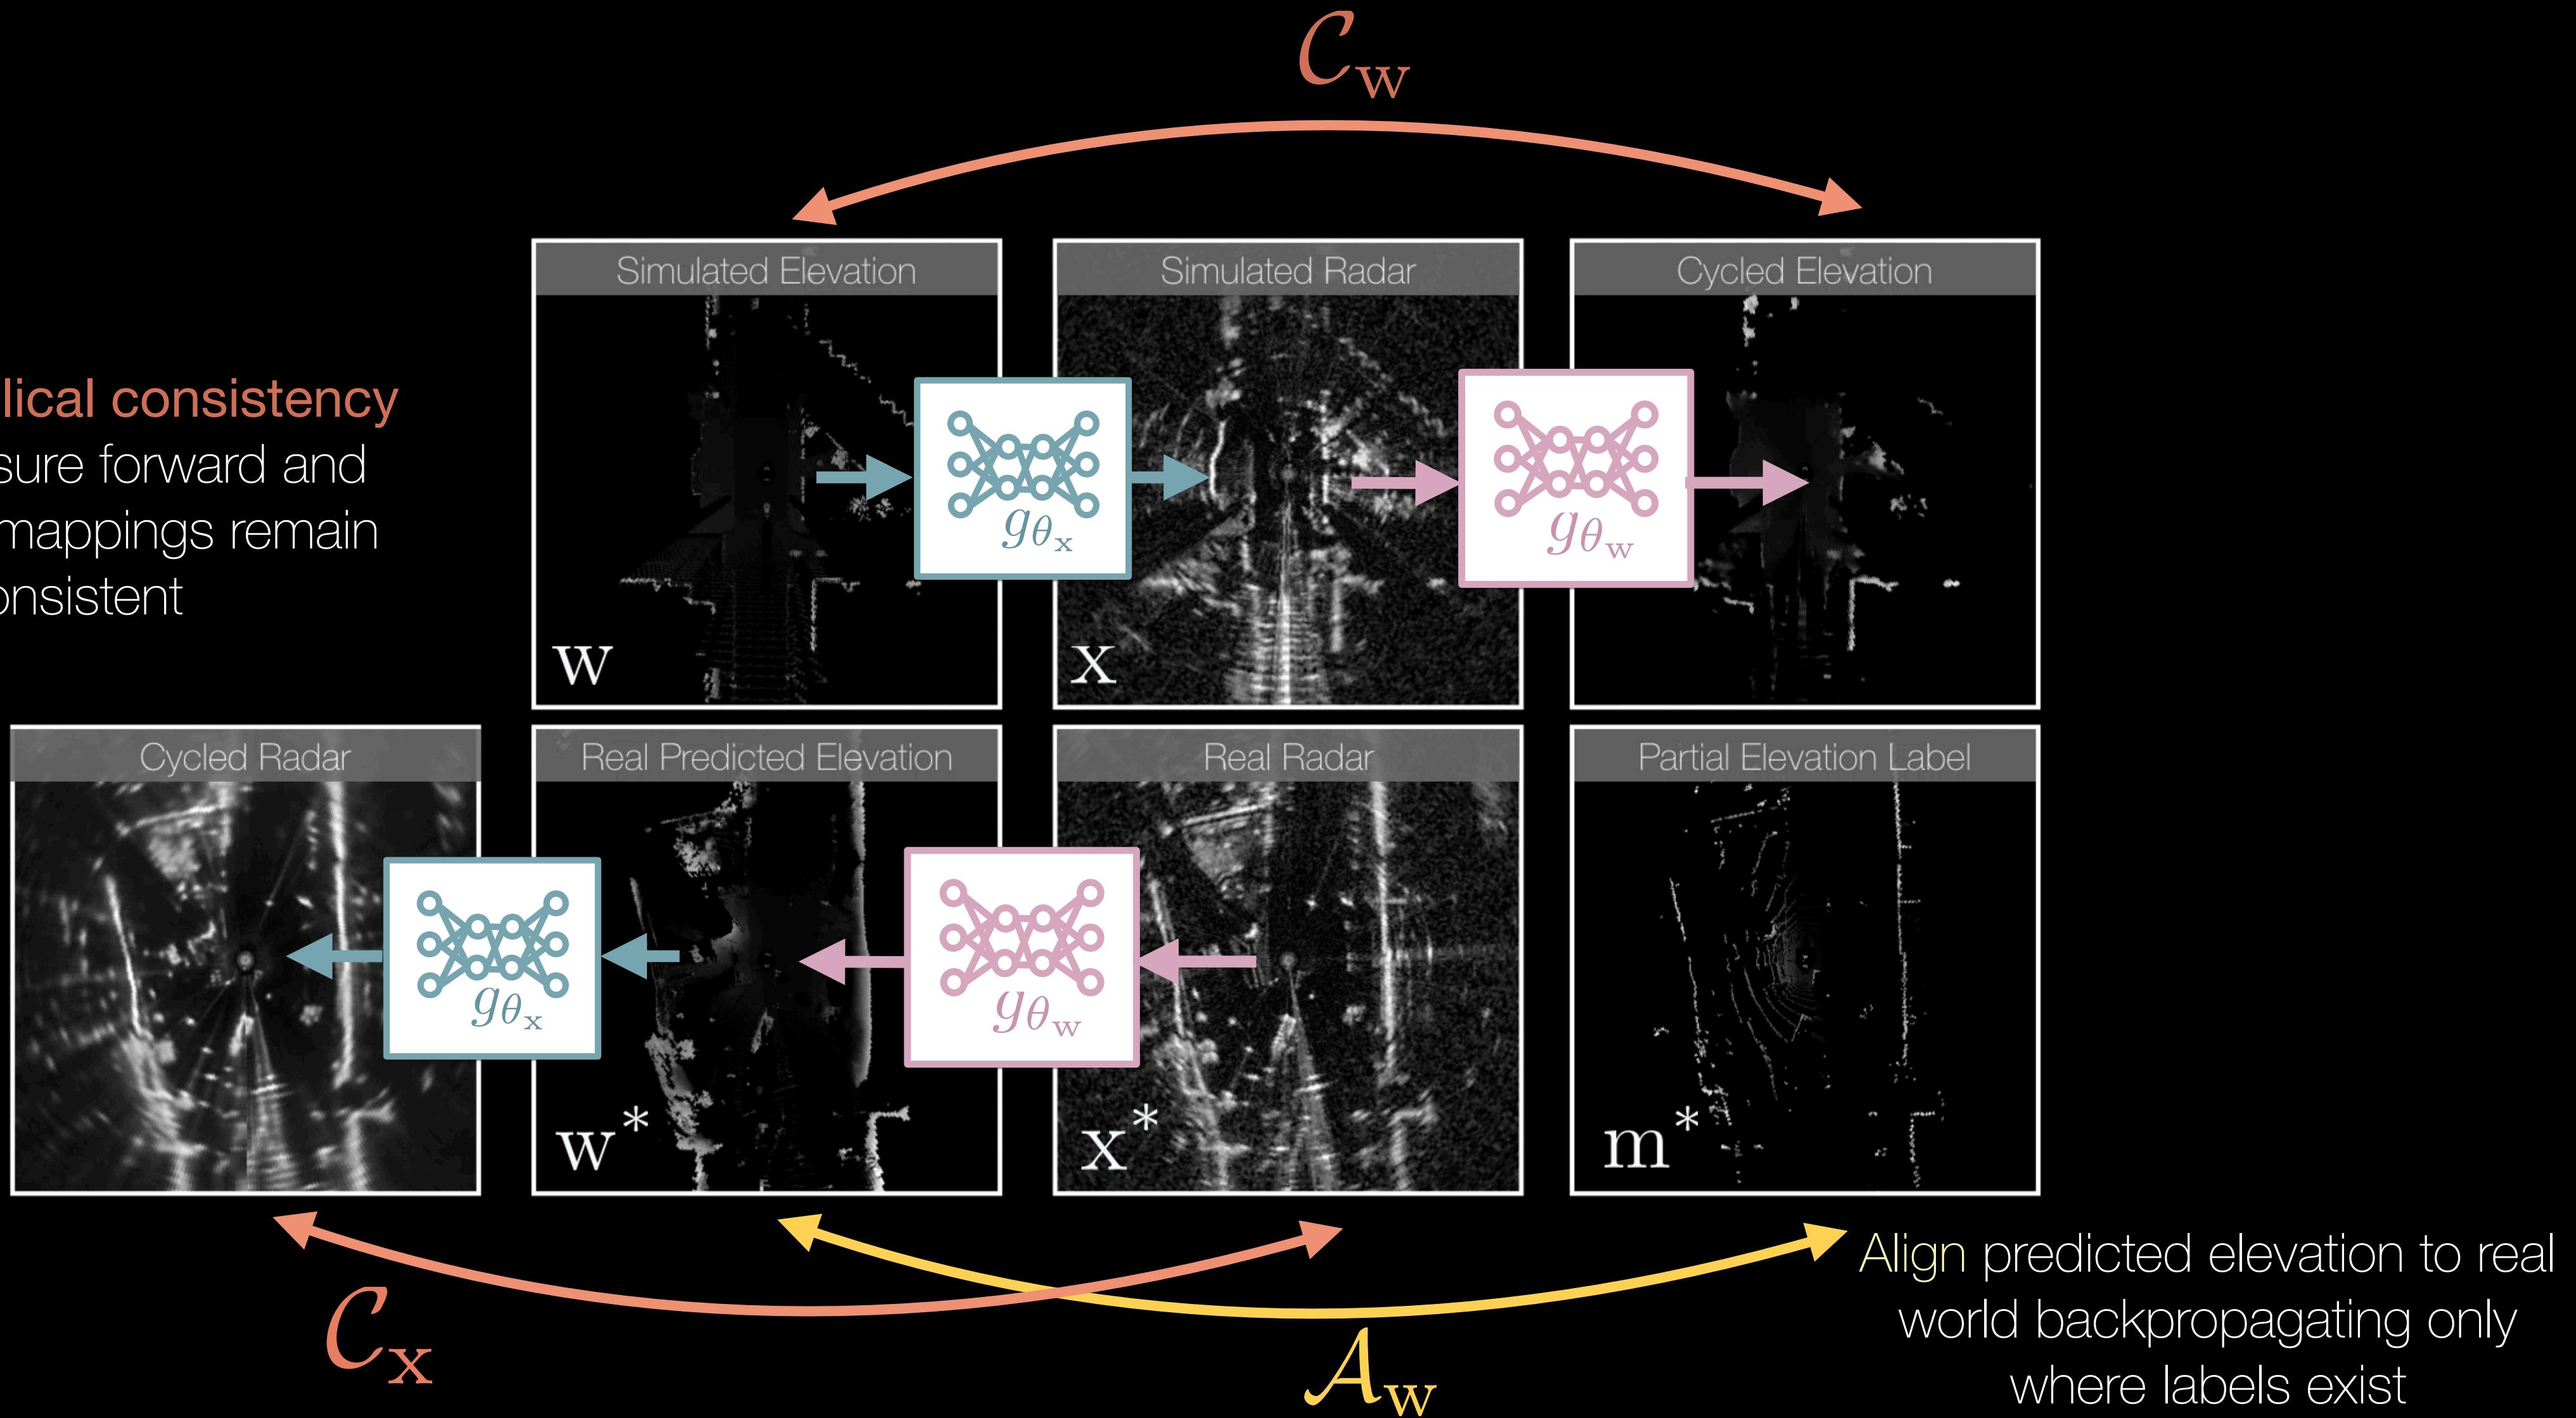
Use L1 **cyclical consistency**  
loss to ensure forward and  
backward mappings remain  
consistent



# Our Approach

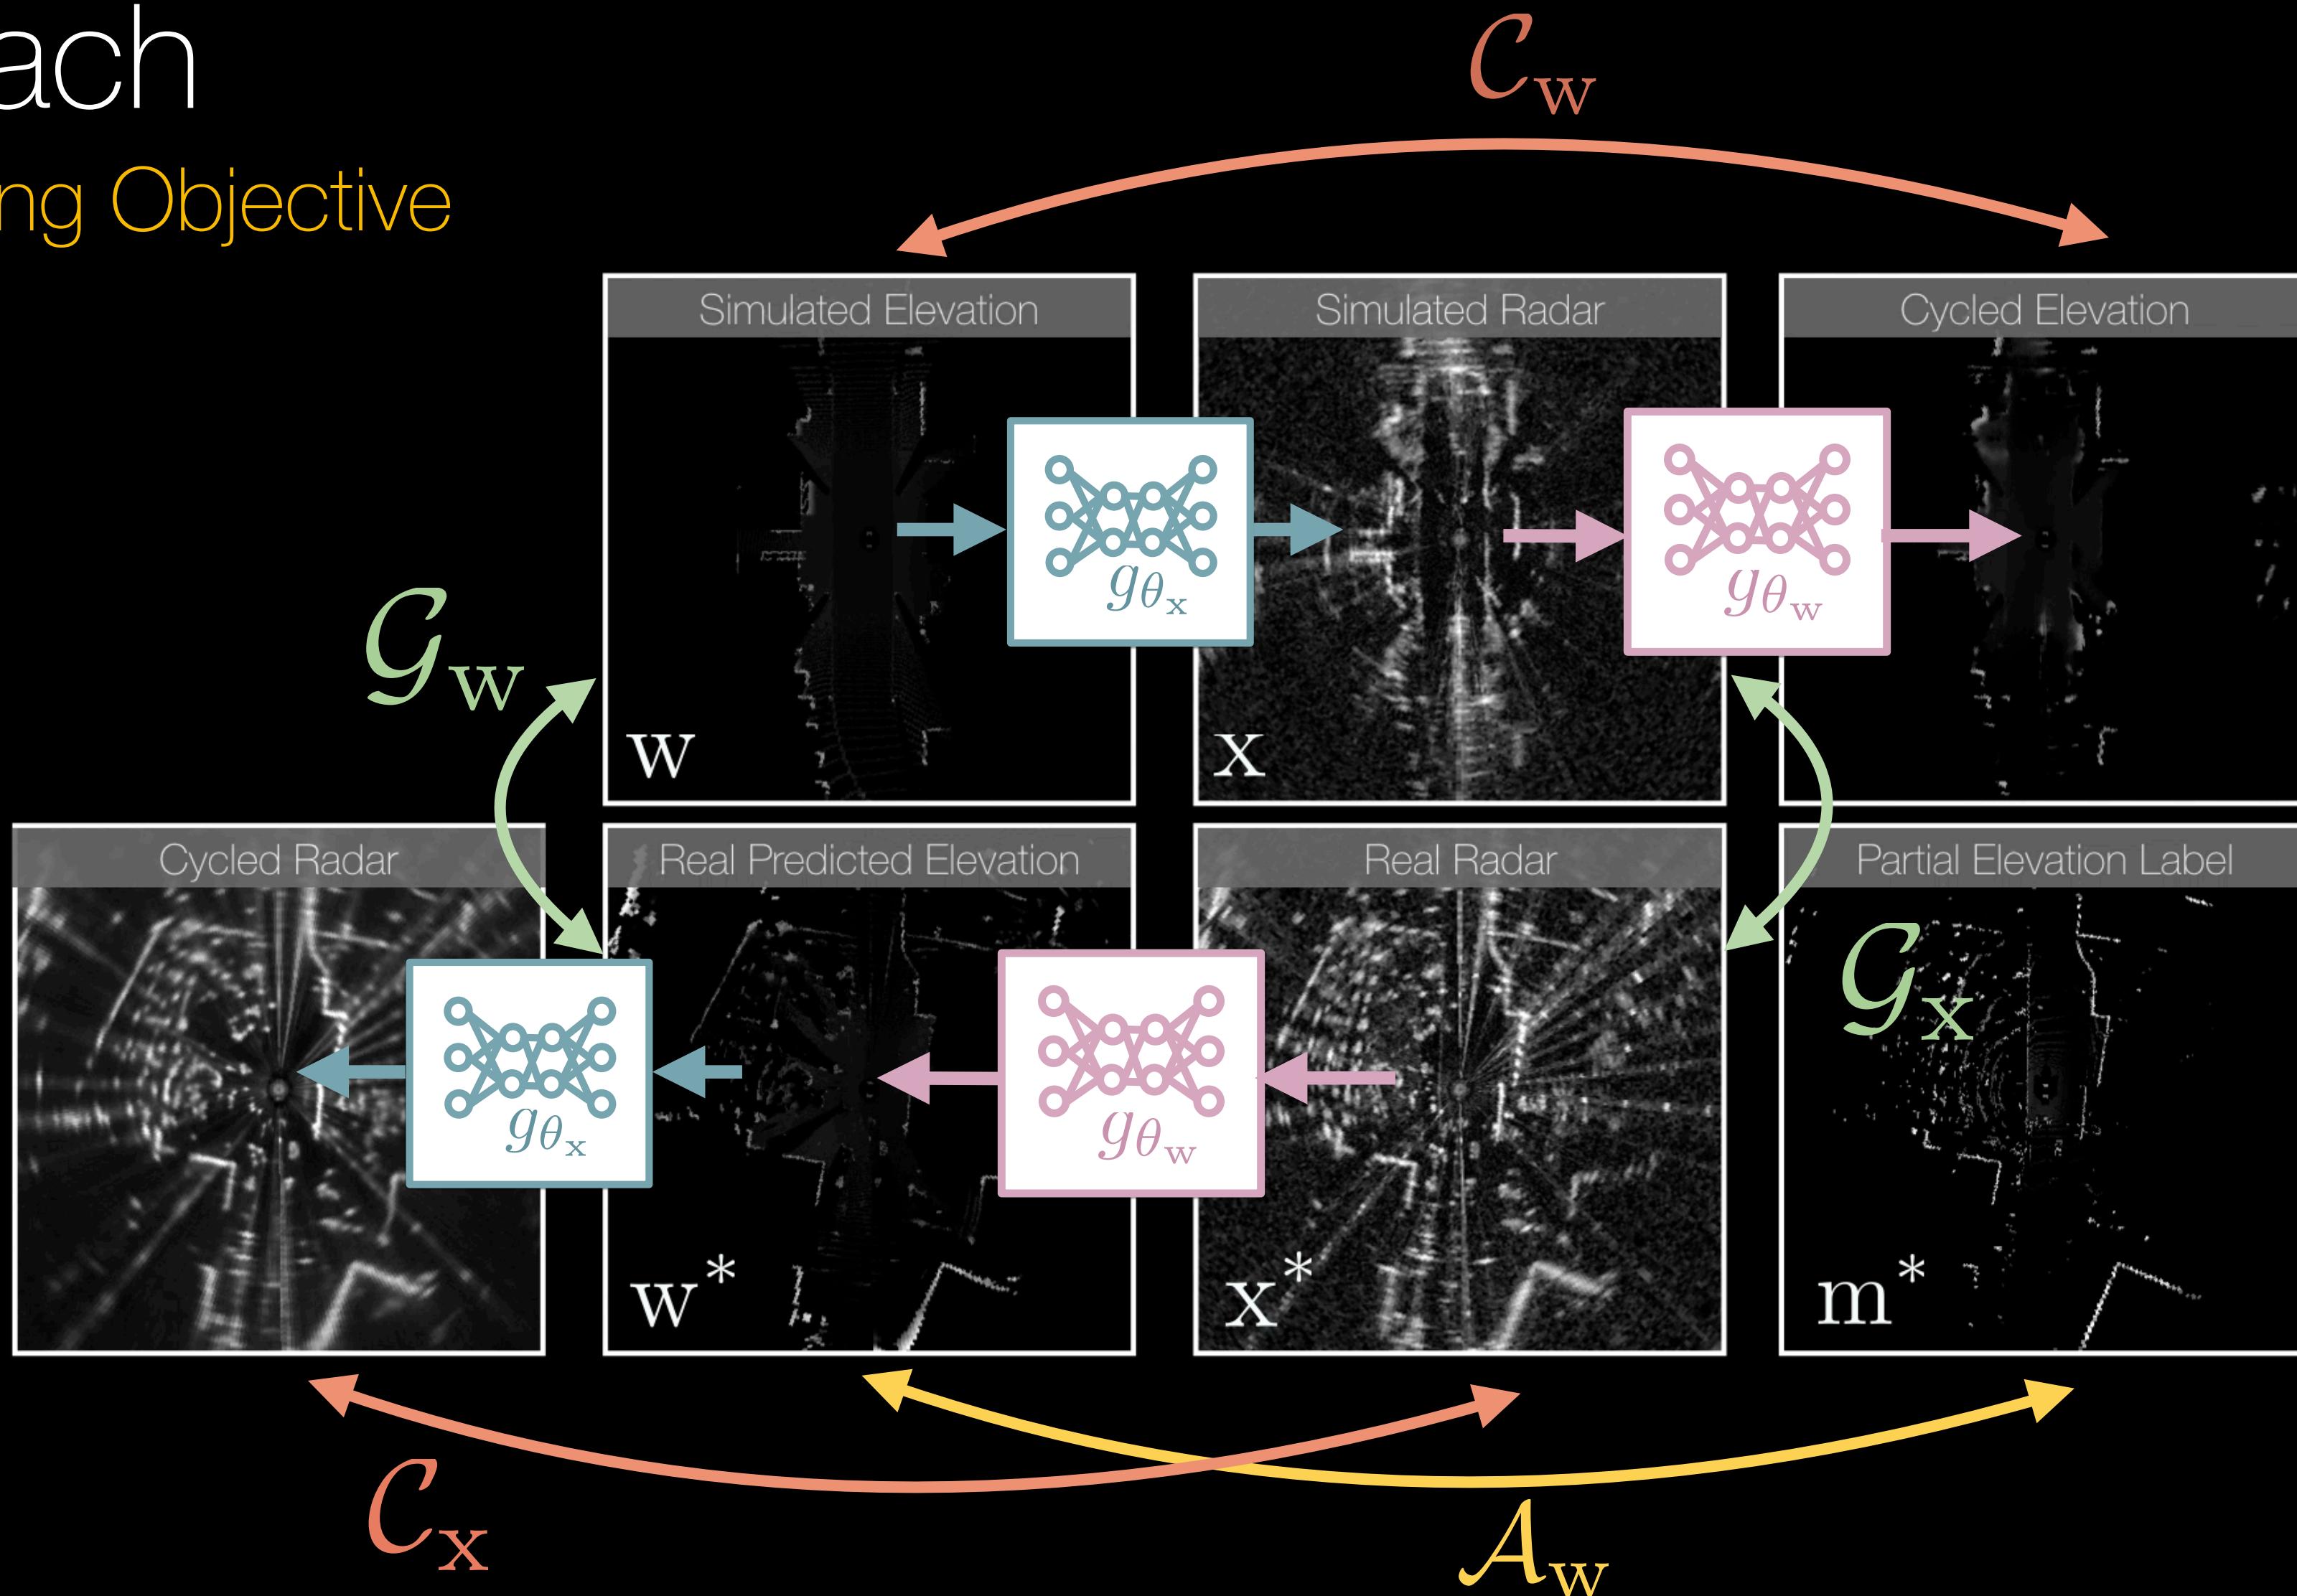
Further constrain training with **cyclical consistency** and **real world alignment**

Use L1 **cyclical consistency**  
loss to ensure forward and  
backward mappings remain  
consistent



# Approach

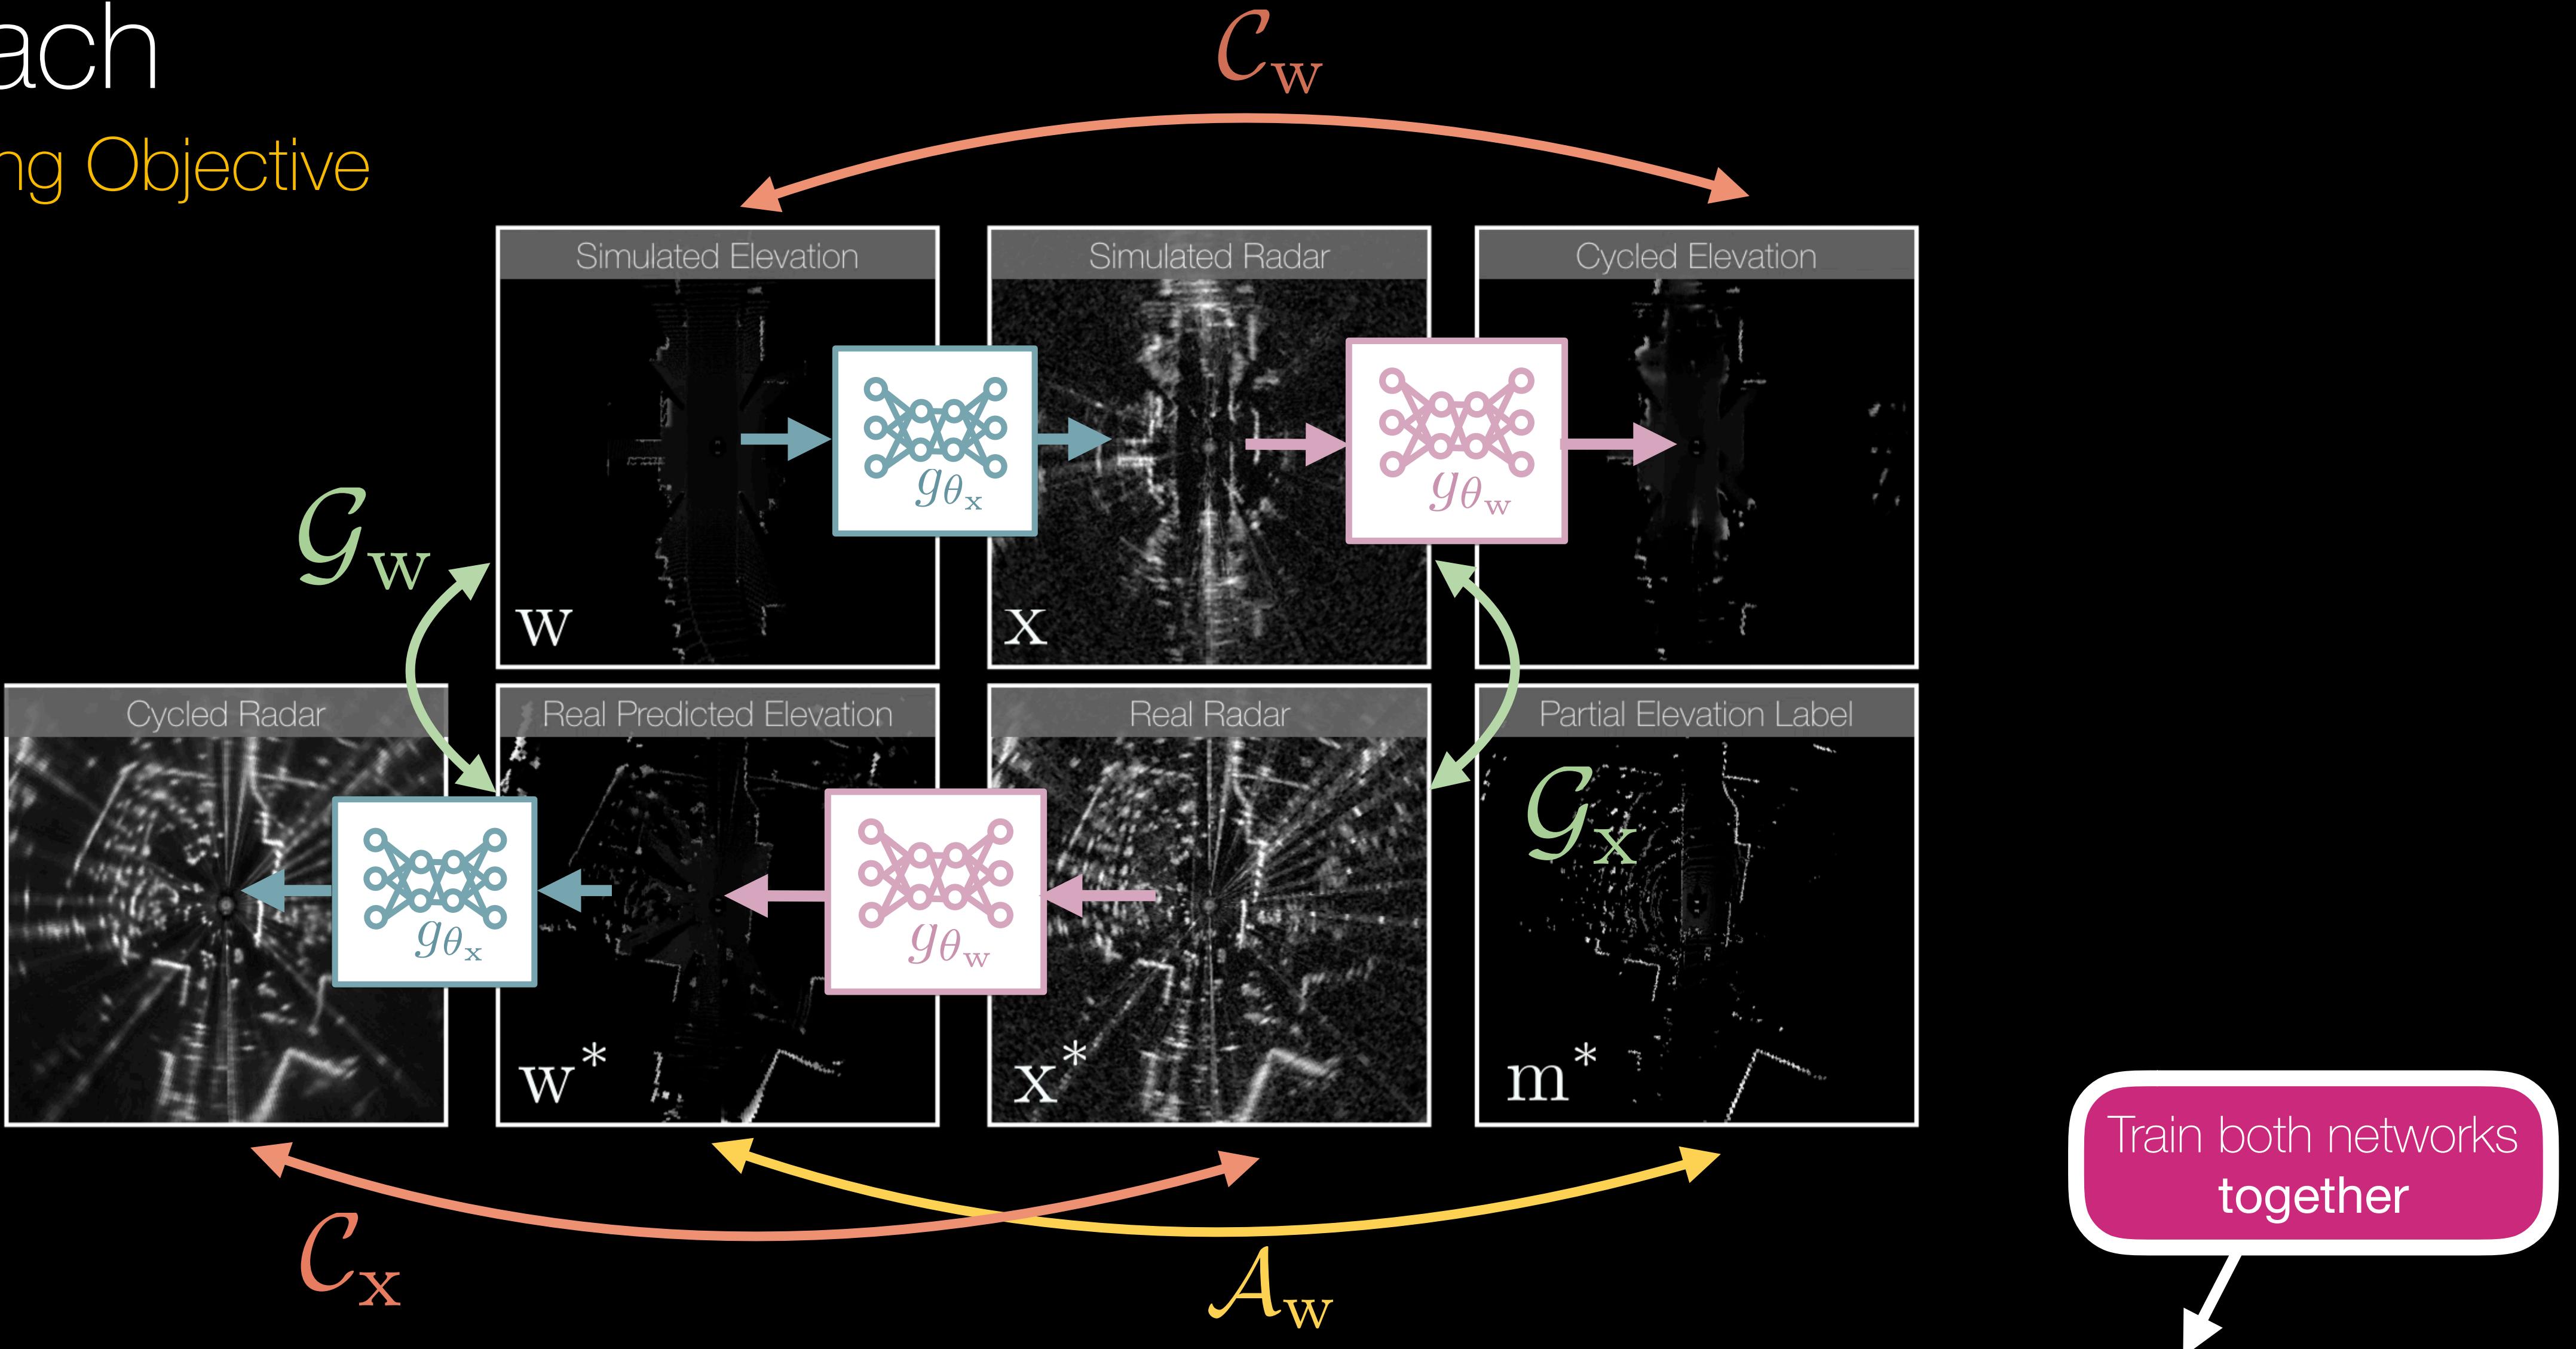
## Final Training Objective



$$\mathcal{L}(\theta_w, \theta_x) = \mathcal{G}_w + \mathcal{G}_x + \mathcal{C}_x + \mathcal{C}_w + \mathcal{A}_w$$

# Approach

## Final Training Objective



$$\mathcal{L}(\theta_w, \theta_x) = \mathcal{G}_w + \mathcal{G}_x + \mathcal{C}_x + \mathcal{C}_w + \mathcal{A}_w$$

# Results

## Radar Simulation Evaluation Setup

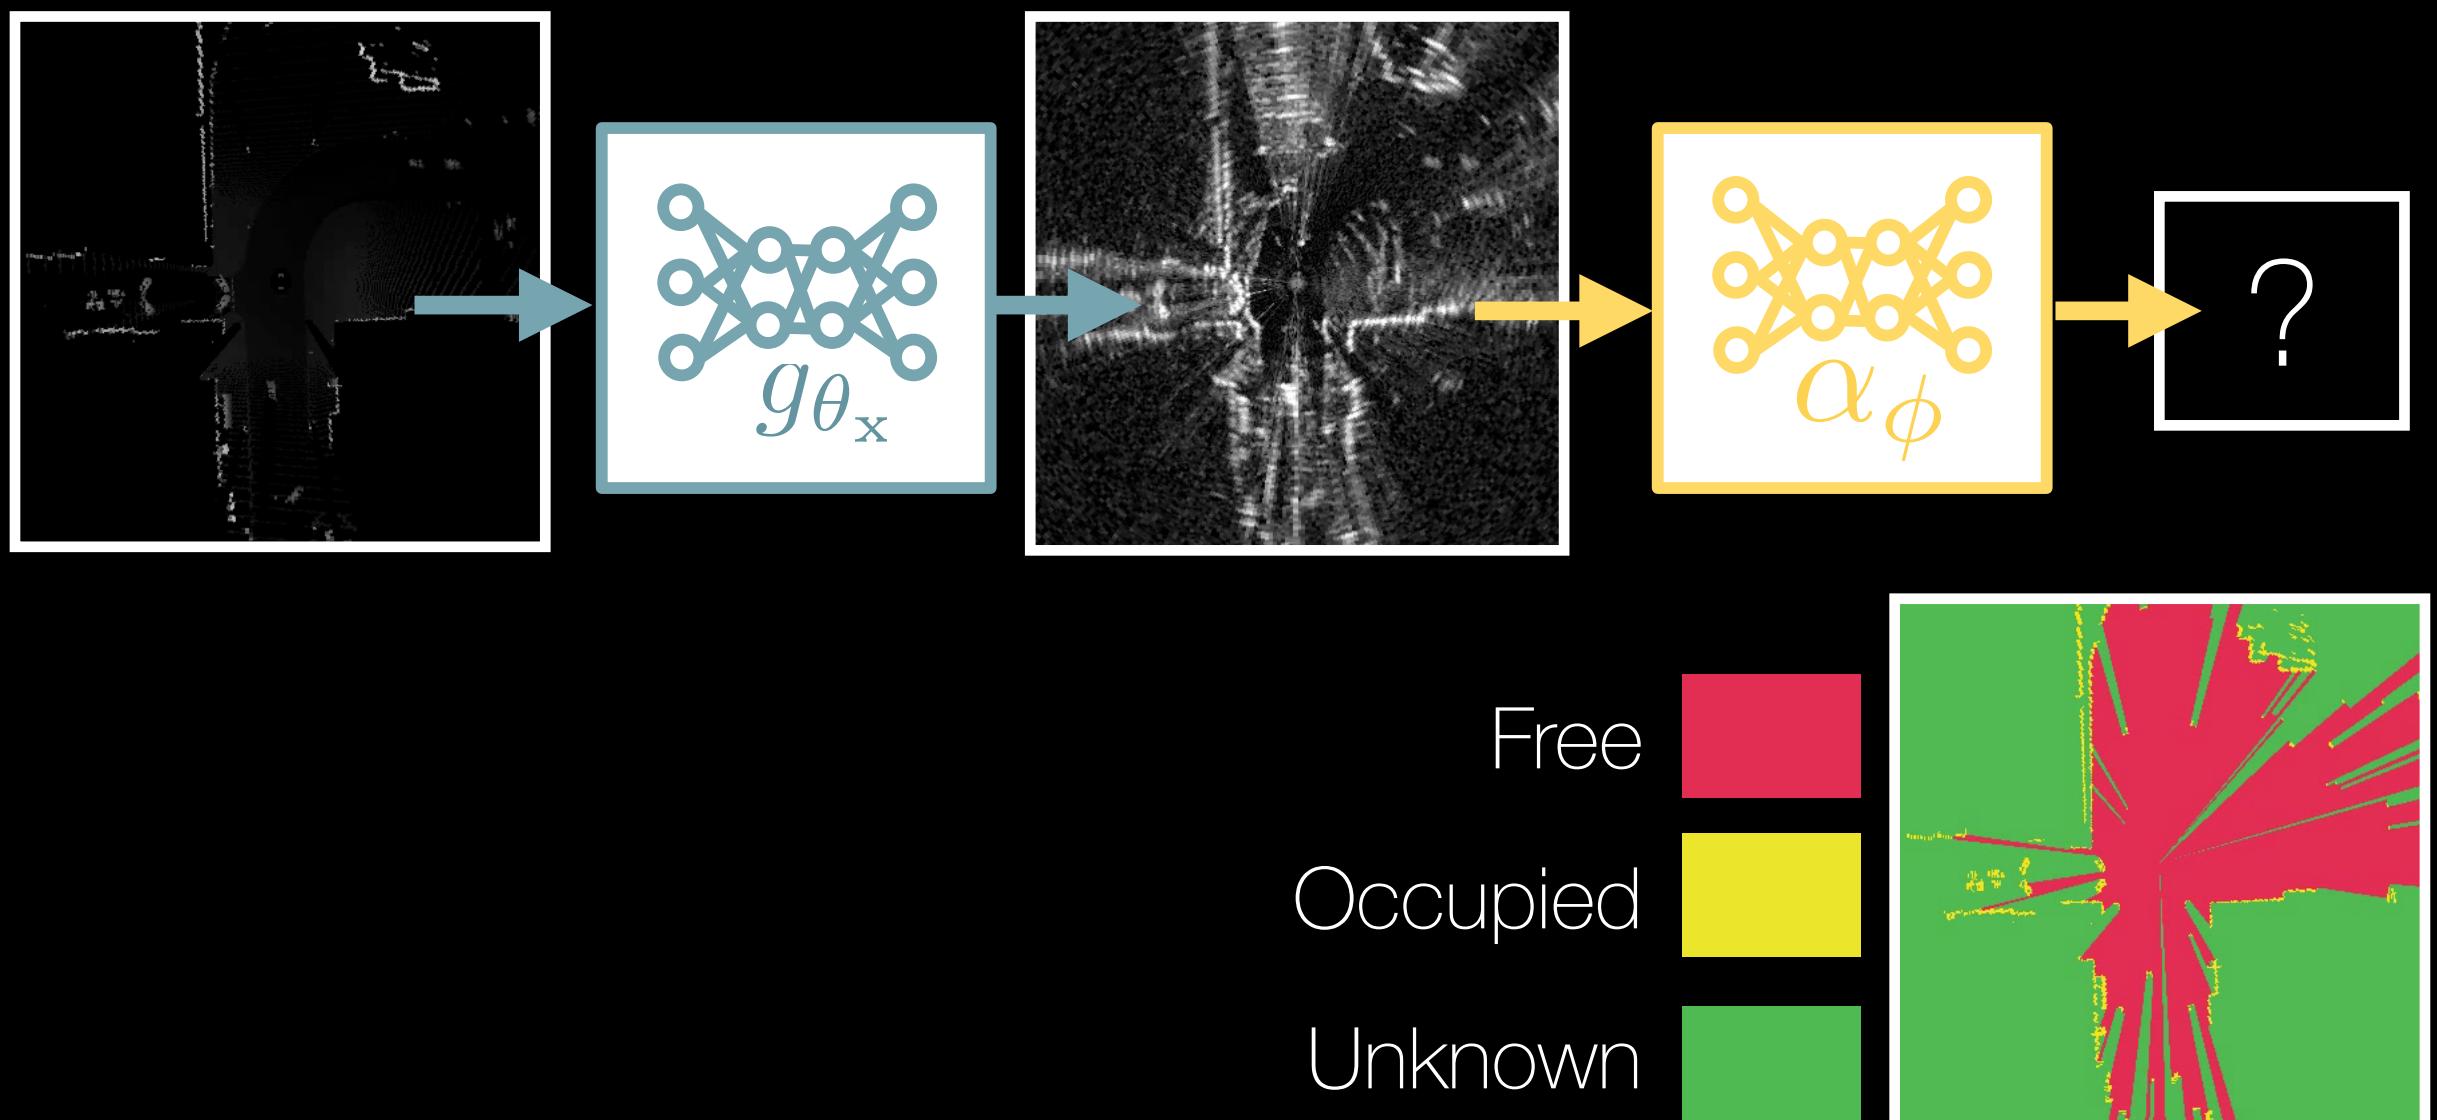
“Are we able to train **new** models in simulation?”

# Results

## Radar Simulation Evaluation Setup

“Are we able to train **new** models in simulation?”

(1) **Train** segmentation model in **simulation**

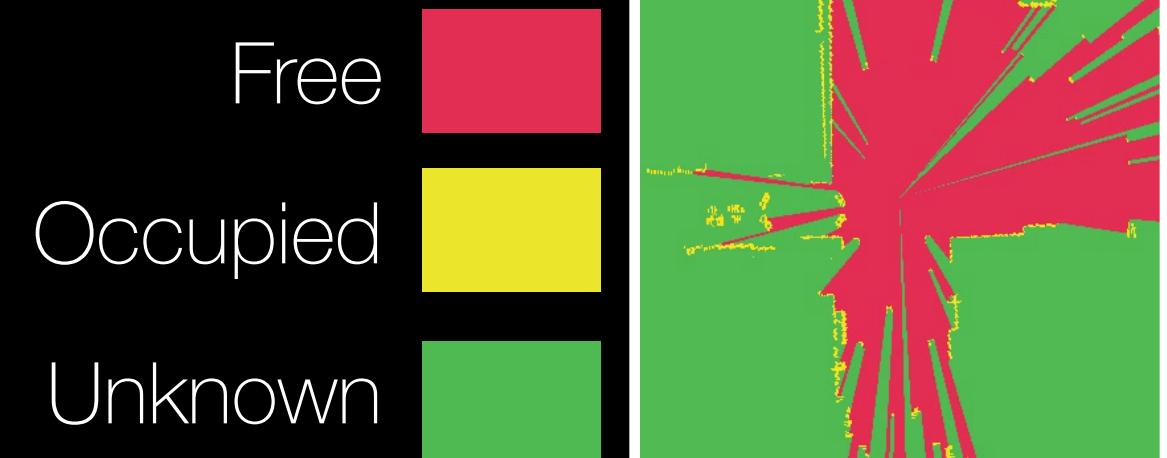
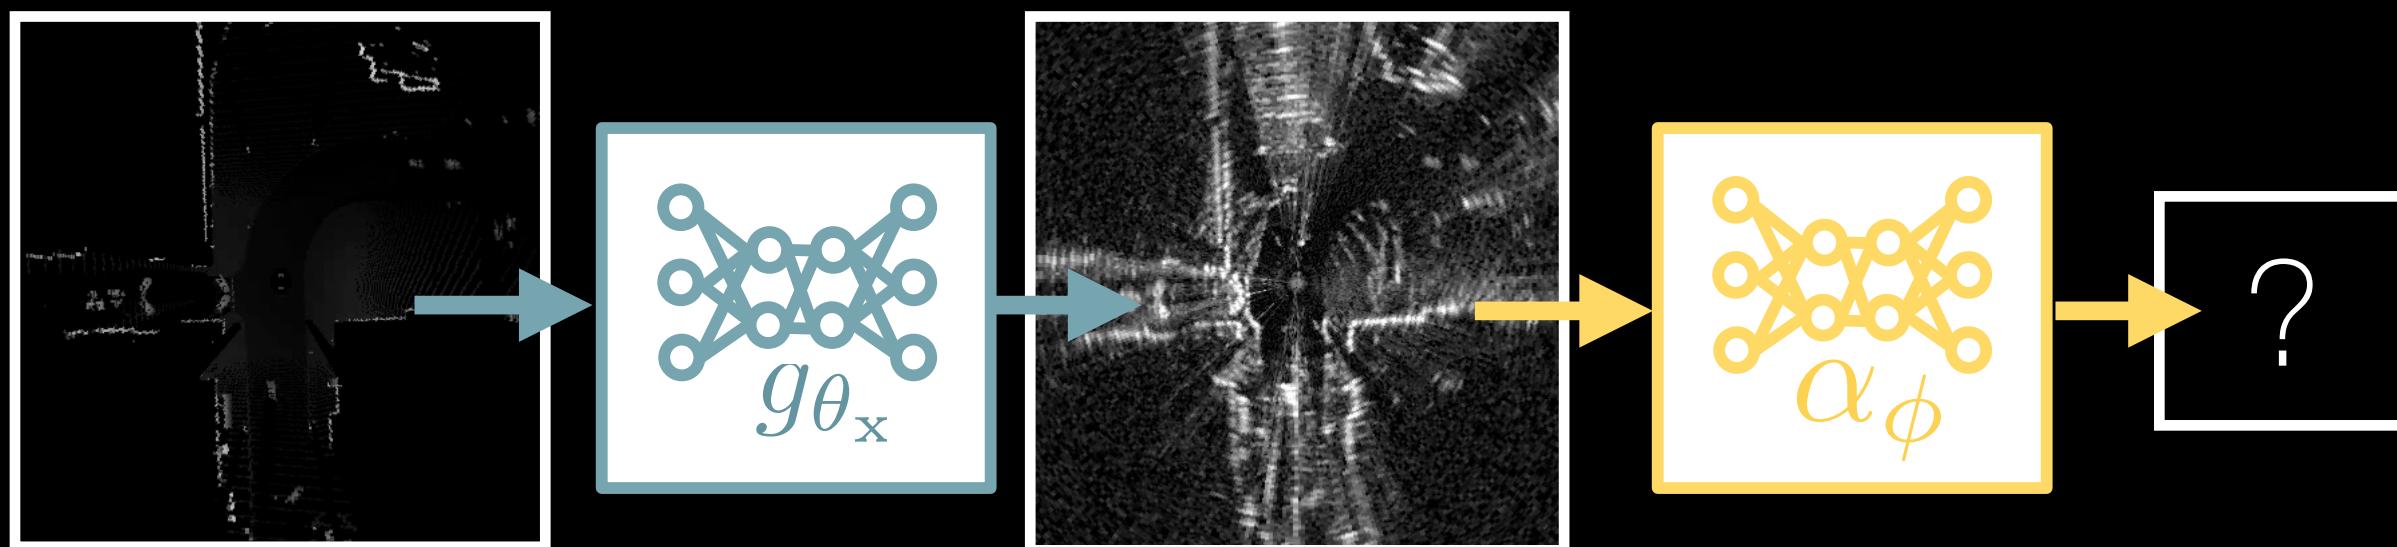


# Results

## Radar Simulation Evaluation Setup

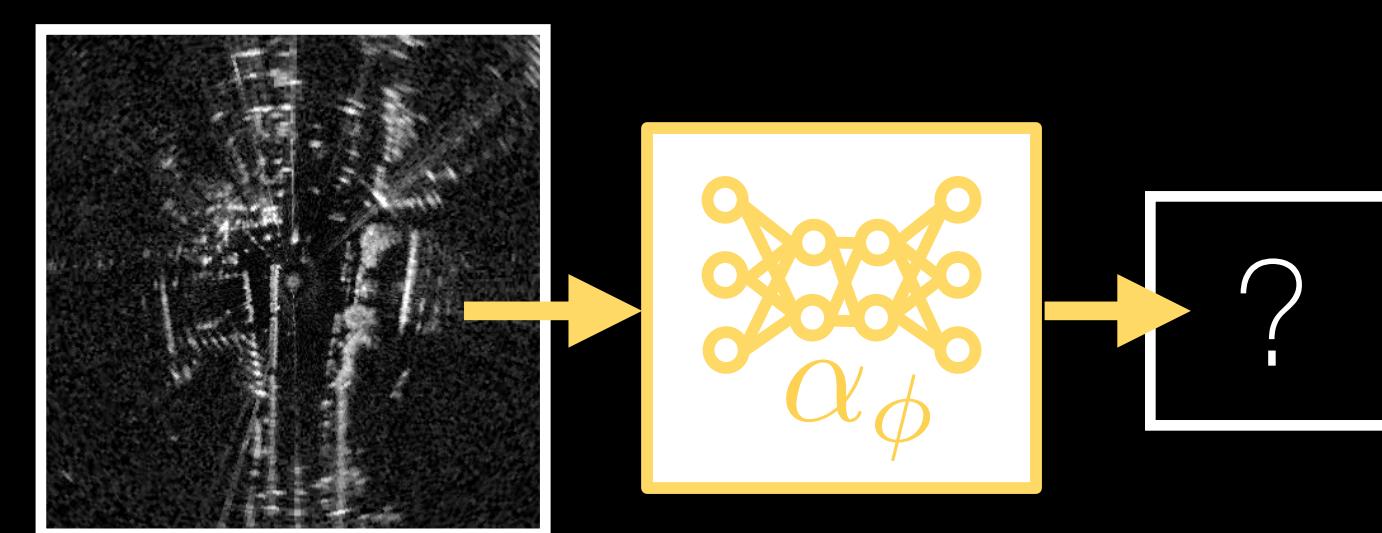
“Are we able to train **new** models in simulation?”

(1) **Train** segmentation model in **simulation**

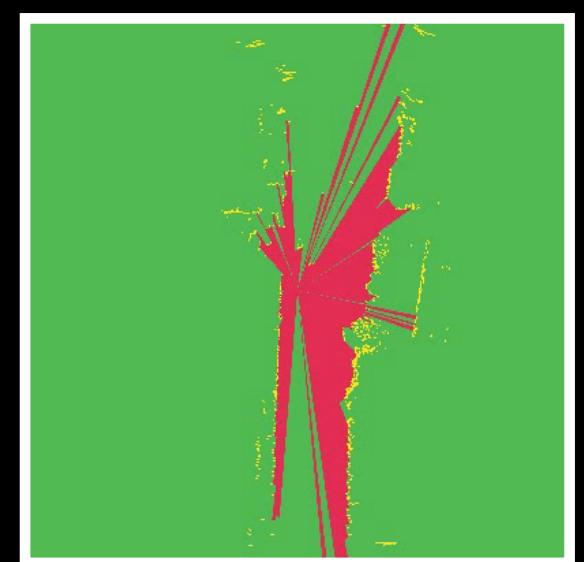


Ground truth training  
labels from simulation

(2) **Test** Model in **real world**



Real world  
ground truth



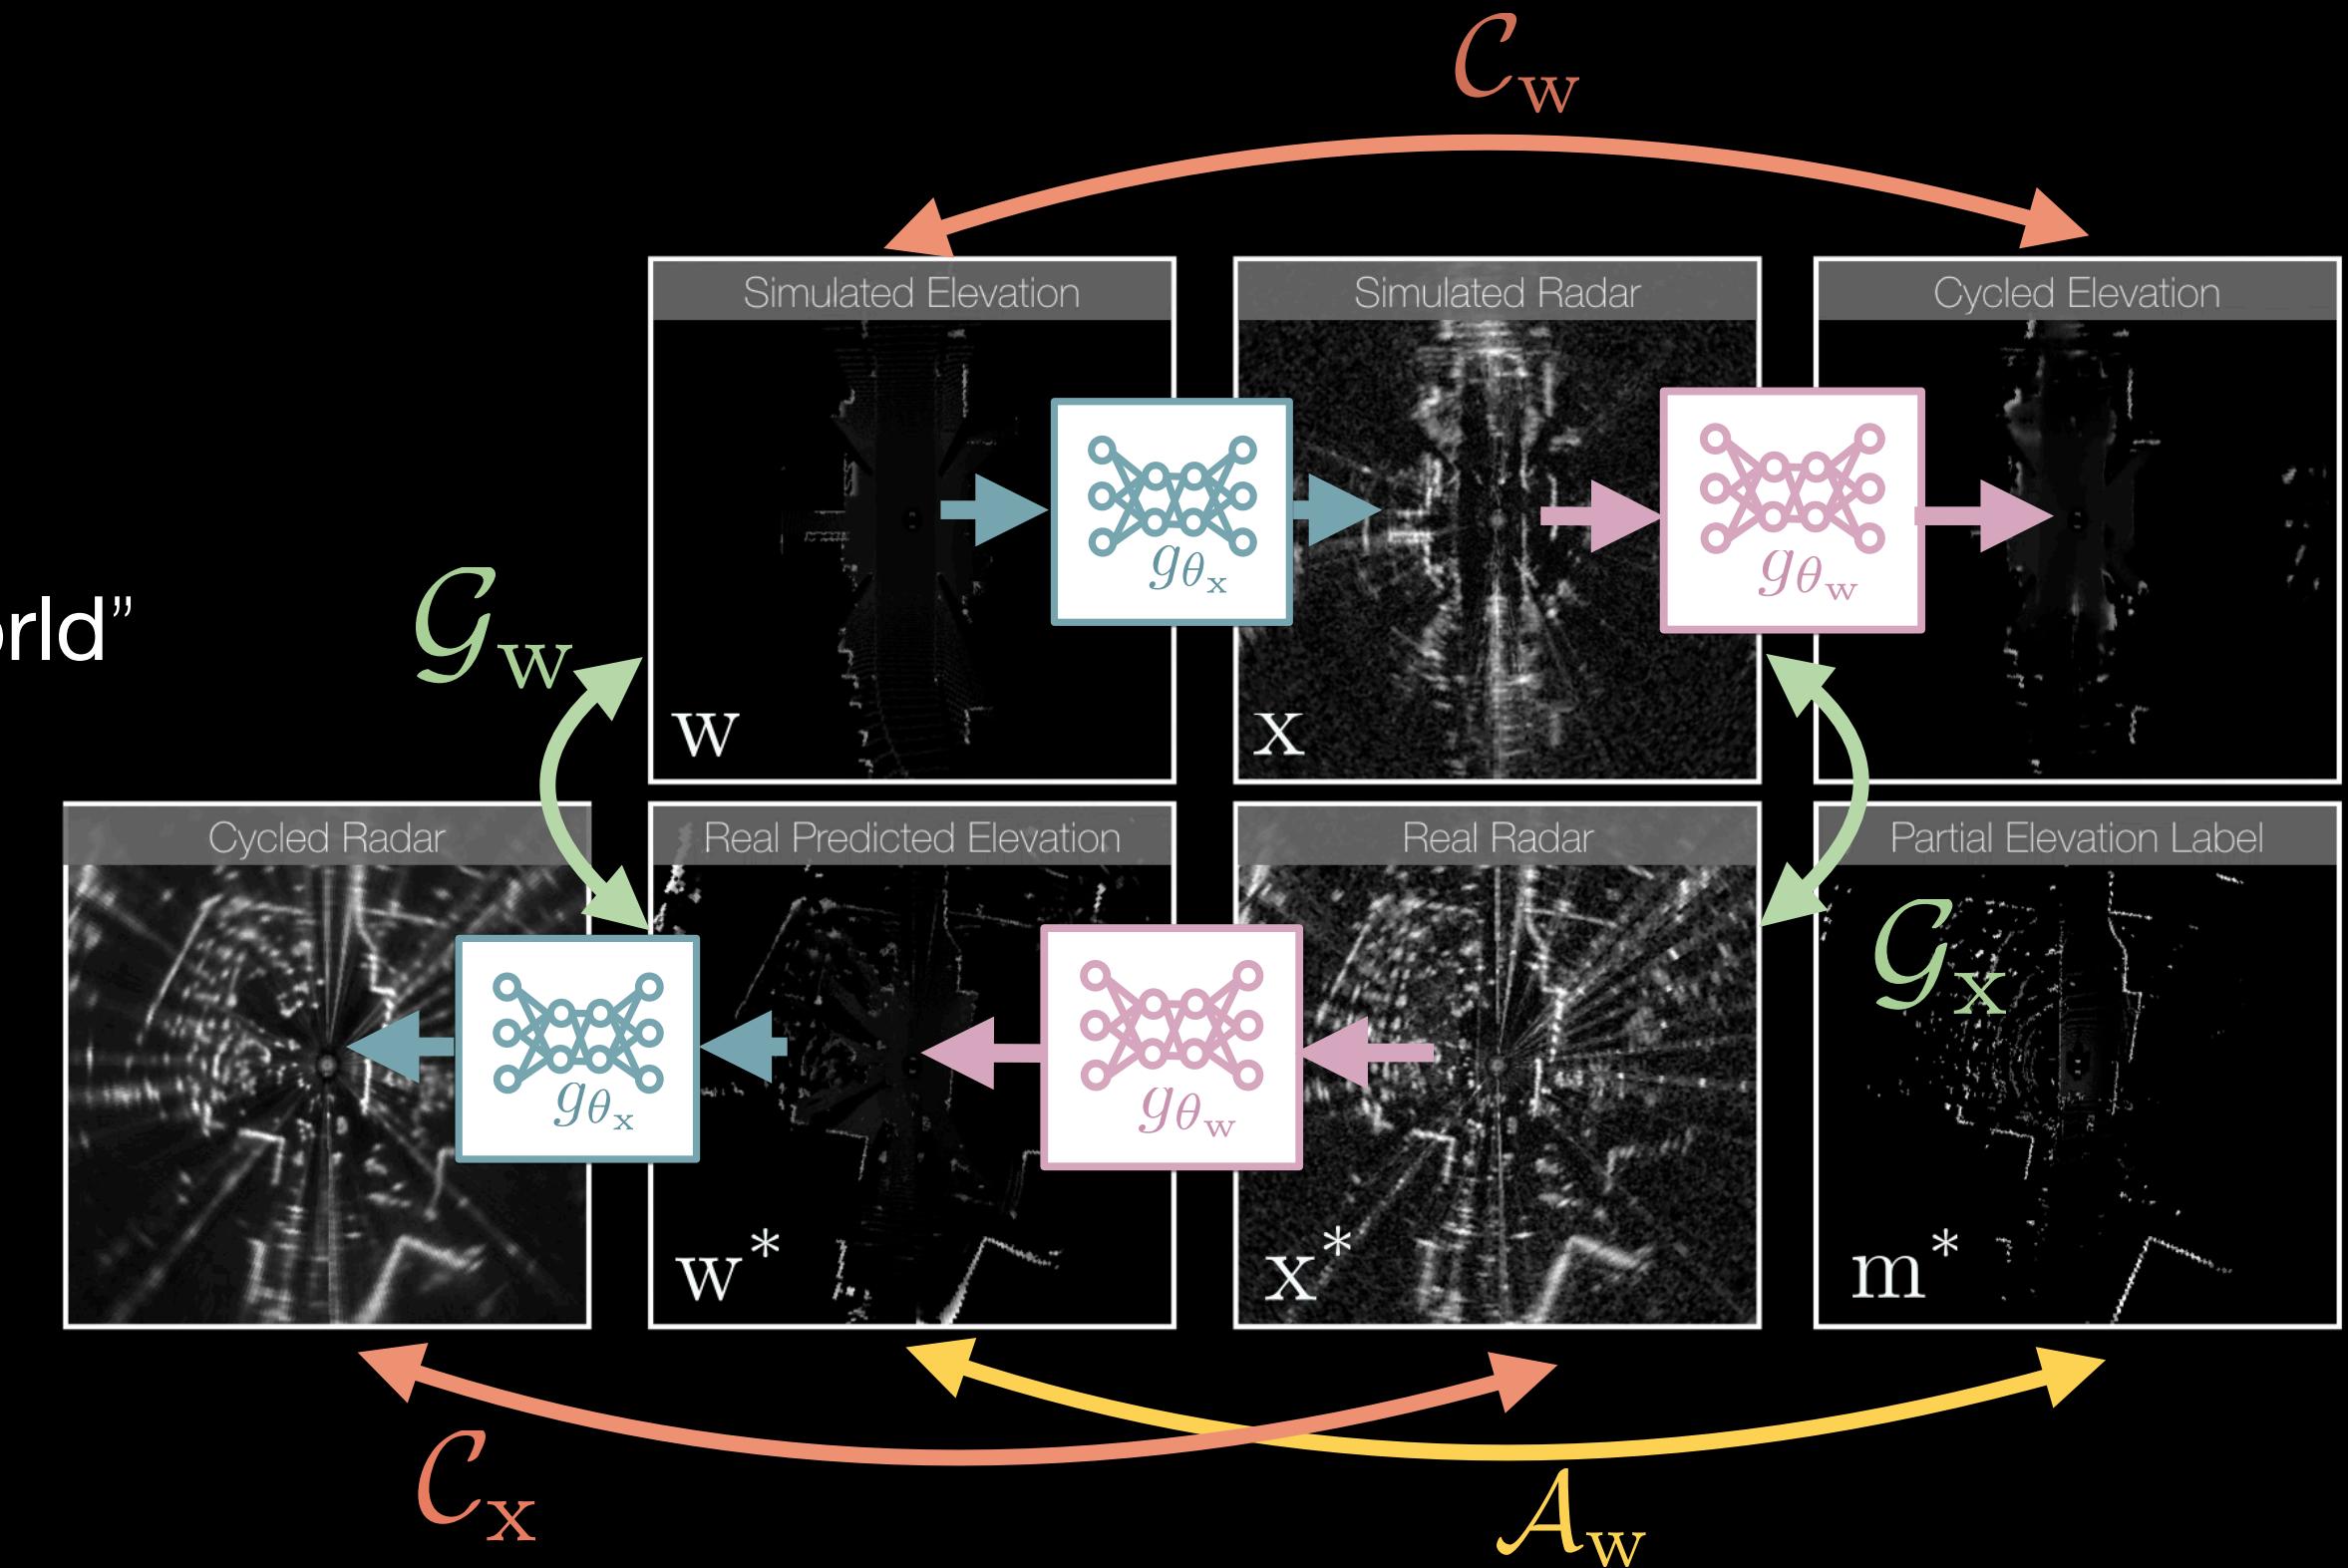
$$\mathcal{M}(\theta_x) = \frac{1}{2} \sum_c \left[ \frac{\text{TP}(c)}{\text{FP}(c) + \text{FN}(c) + \text{TP}(c)} \right]$$

Mean Intersection Over Union  
Metric (mIoU) over free and  
occupied

# Results

## Radar Simulation

- Benchmark = “segmentation model trained in the **real world**”
- (a) and (b) trained only on real elevation inputs
- (c) trained only with GAN loss
- (d) trained with GAN loss + cycle constraints
- (e) trained with **full** training criterion

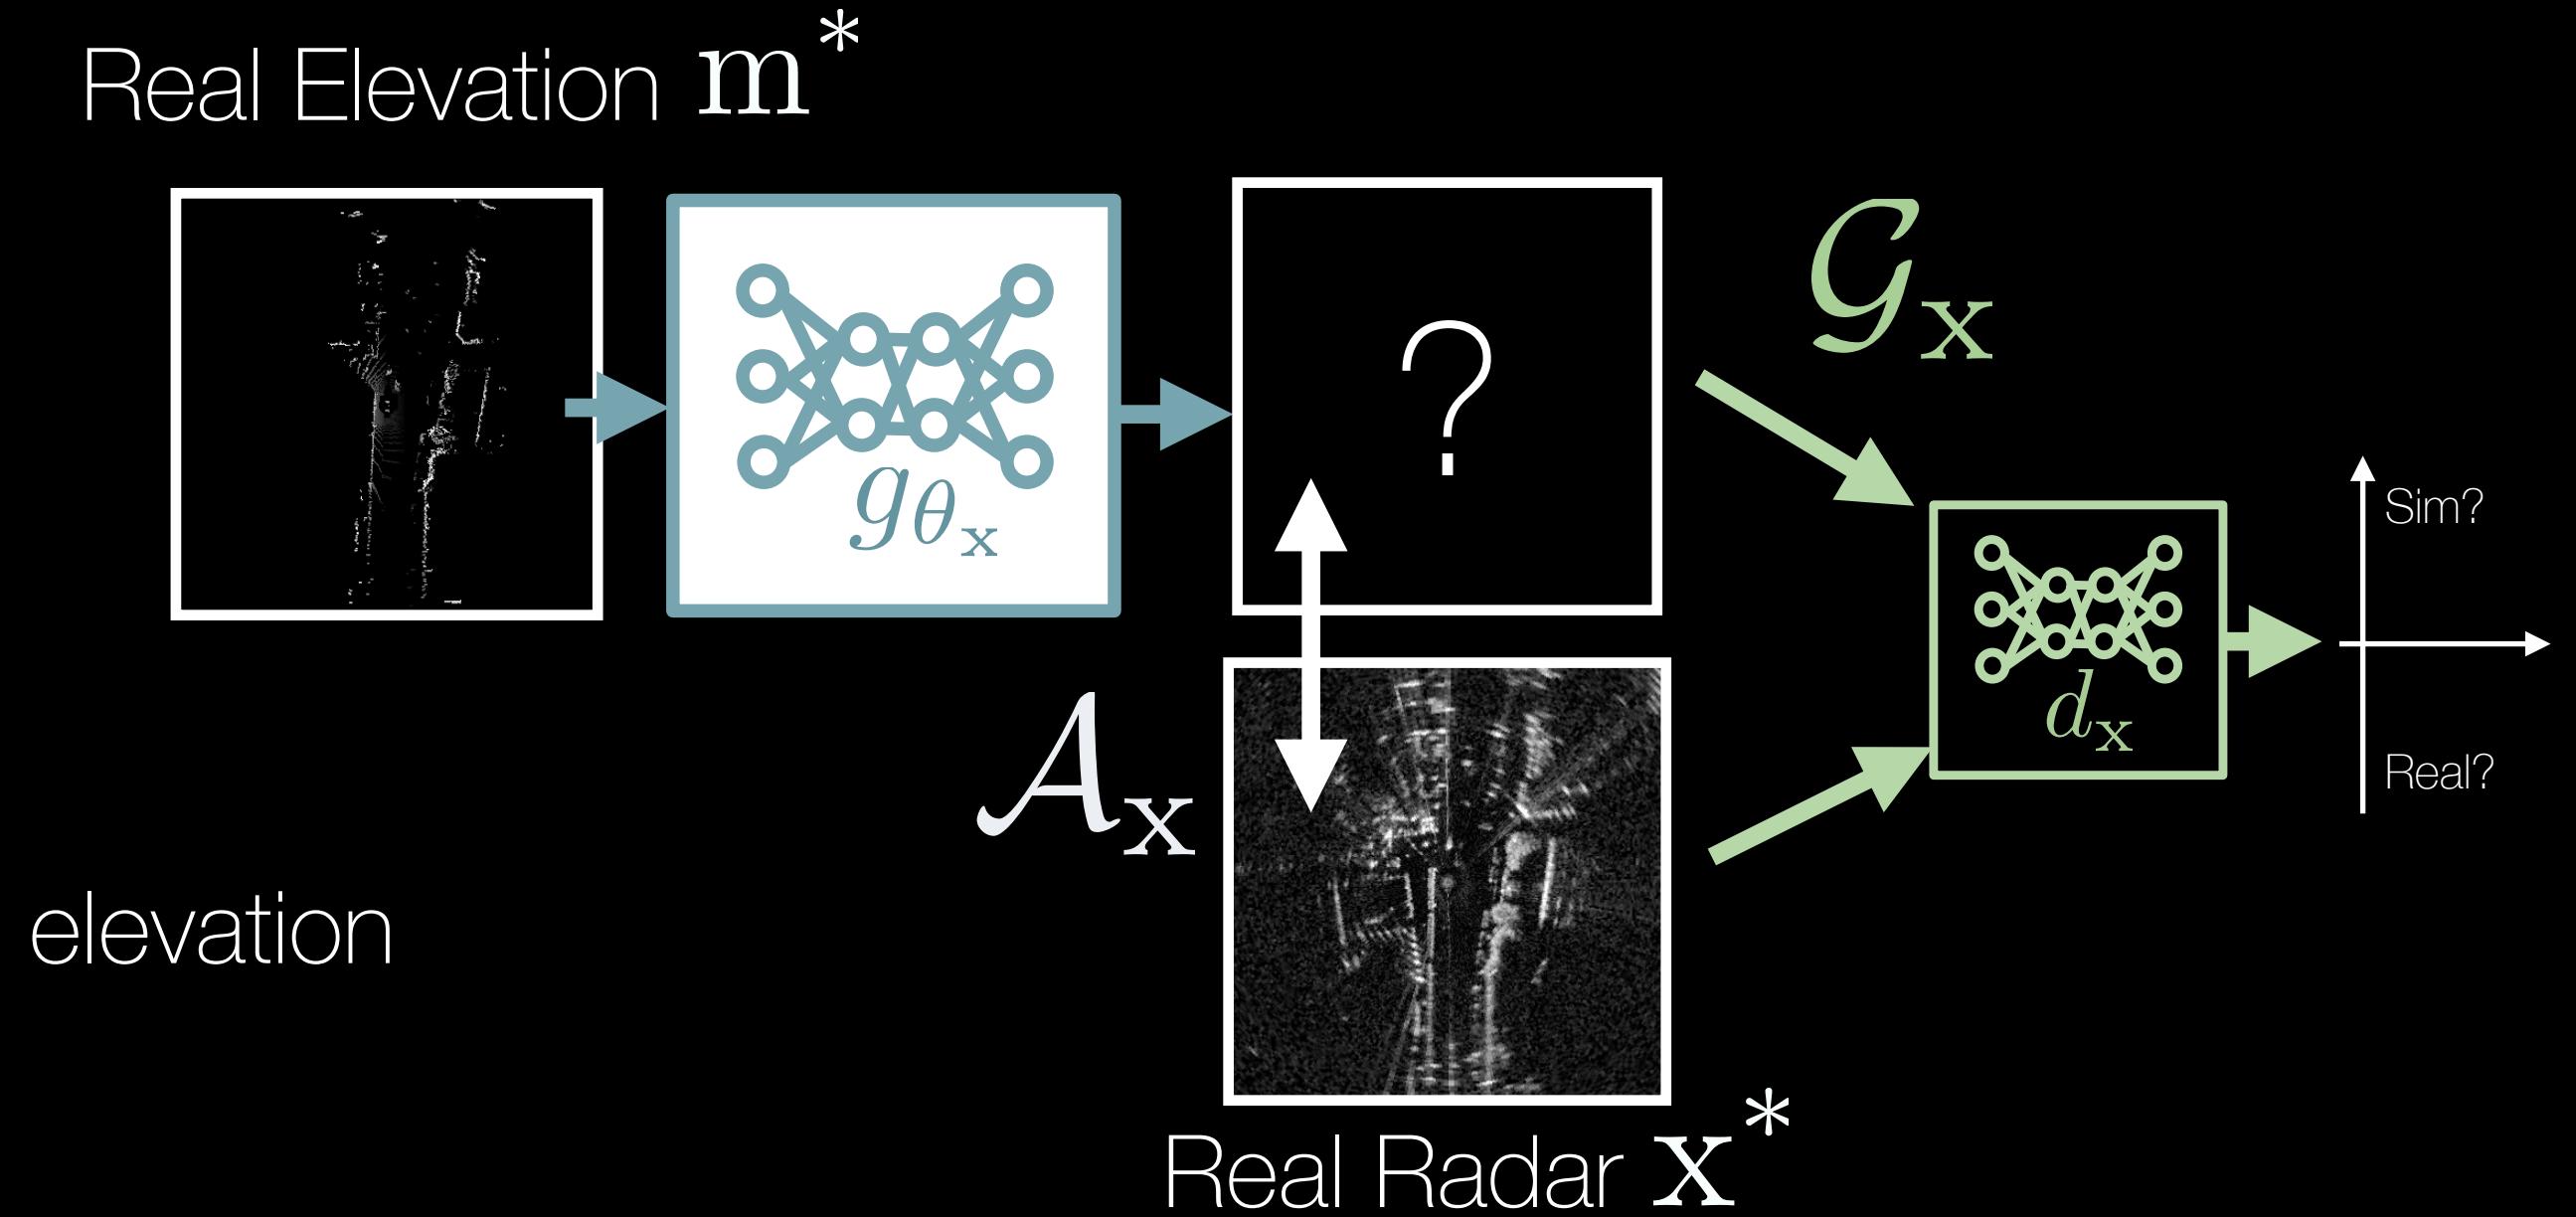


trained on	training objective						intersection over union		
	$\mathcal{A}_x$	$\mathcal{A}_w$	$\mathcal{G}_x$	$\mathcal{G}_w$	$\mathcal{C}_x$	$\mathcal{C}_w$	free	occ	mean
<b>benchmark</b>	-	-	-	-	-	-	0.856	0.553	0.705
<b>ours</b>									
(a)	x*	m*	-	✓	-	-	0.396 (0.00)	0.275 (0.00)	0.335 (0.00)
(b)	x*	m*	-	✓	-	✓	0.558 (0.11)	0.221 (0.03)	0.389 (0.07)
(c)	x*	-	w	-	-	✓	0.385 (0.07)	0.148 (0.01)	0.266 (0.04)
(d)	x*	-	w	-	-	✓	0.845 (0.02)	0.262 (0.02)	0.553 (0.02)
(e)	x*	m*	w	-	✓	✓	<b>0.872 (0.01)</b>	<b>0.455 (0.01)</b>	<b>0.664 (0.01)</b>

# Results

## Radar Simulation

- (a) and (b) simulators trained only on **real** data
  - (a) trained predicting radar from **real-world** elevation
  - (b) with **additional** GAN loss

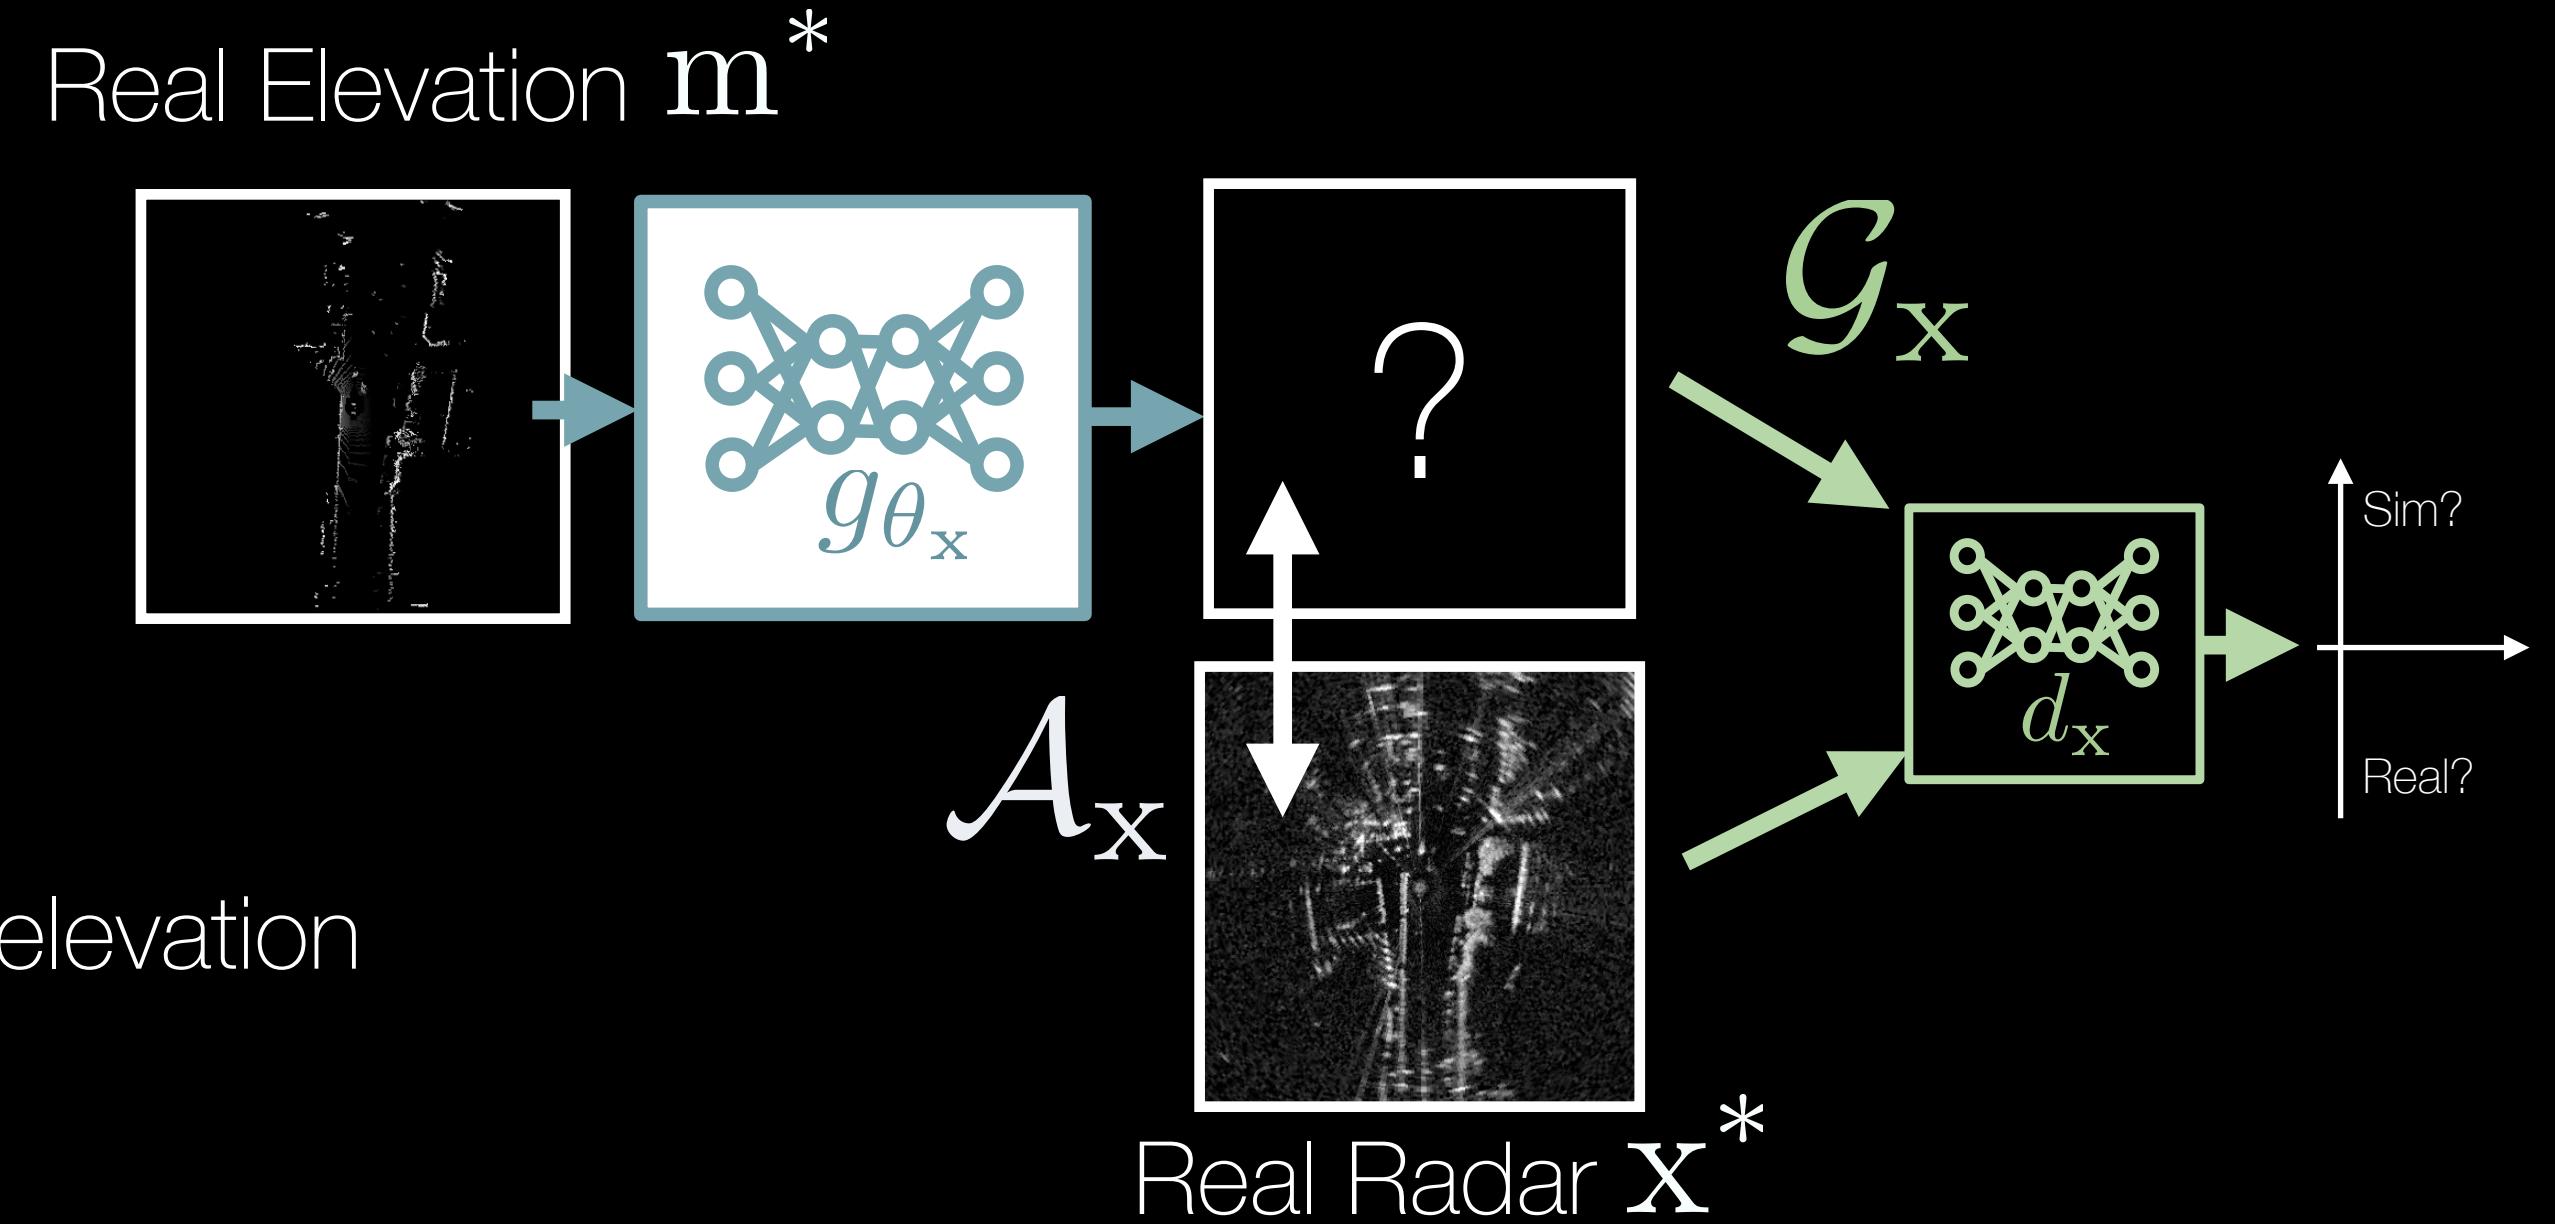


benchmark	trained on	training objective						intersection over union		
		$\mathcal{A}_x$	$\mathcal{A}_w$	$\mathcal{G}_x$	$\mathcal{G}_w$	$\mathcal{C}_x$	$\mathcal{C}_w$	free	occ	mean
real world	-	-	-	-	-	-	-	0.856	0.553	0.705
<b>ours</b>										
(a)	$x^*$	$m^*$	-	✓	-	-	-	0.396 (0.00)	0.275 (0.00)	0.335 (0.00)
(b)	$x^*$	$m^*$	-	✓	-	✓	-	0.558 (0.11)	0.221 (0.03)	0.389 (0.07)
(c)	$x^*$	-	w	-	-	✓	-	0.385 (0.07)	0.148 (0.01)	0.266 (0.04)
(d)	$x^*$	-	w	-	-	✓	✓	0.845 (0.02)	0.262 (0.02)	0.553 (0.02)
(e)	$x^*$	$m^*$	w	-	✓	✓	✓	<b>0.872 (0.01)</b>	<b>0.455 (0.01)</b>	<b>0.664 (0.01)</b>

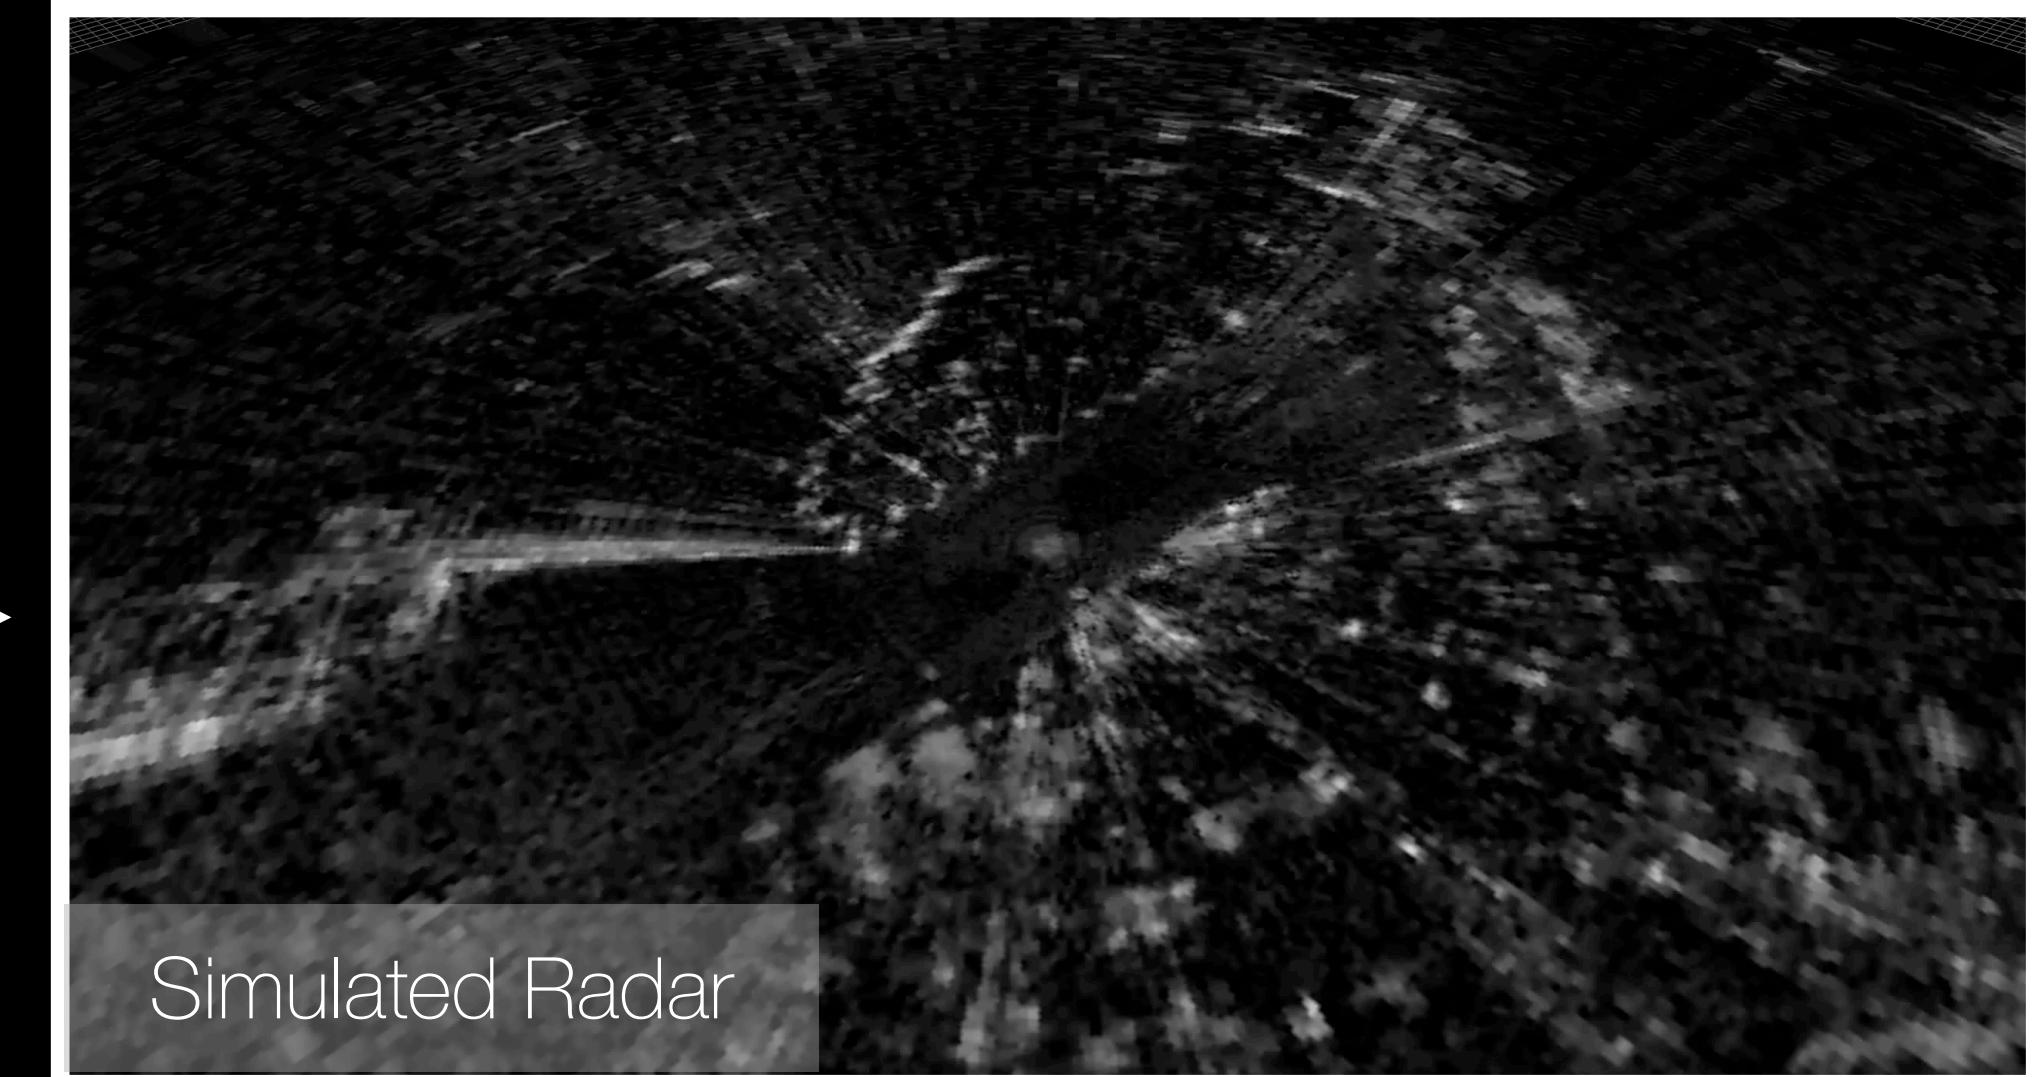
# Results

## Radar Simulation

- (a) and (b) simulators trained only on **real** data
  - (a) trained predicting radar from **real-world** elevation
  - (b) with **additional** GAN loss



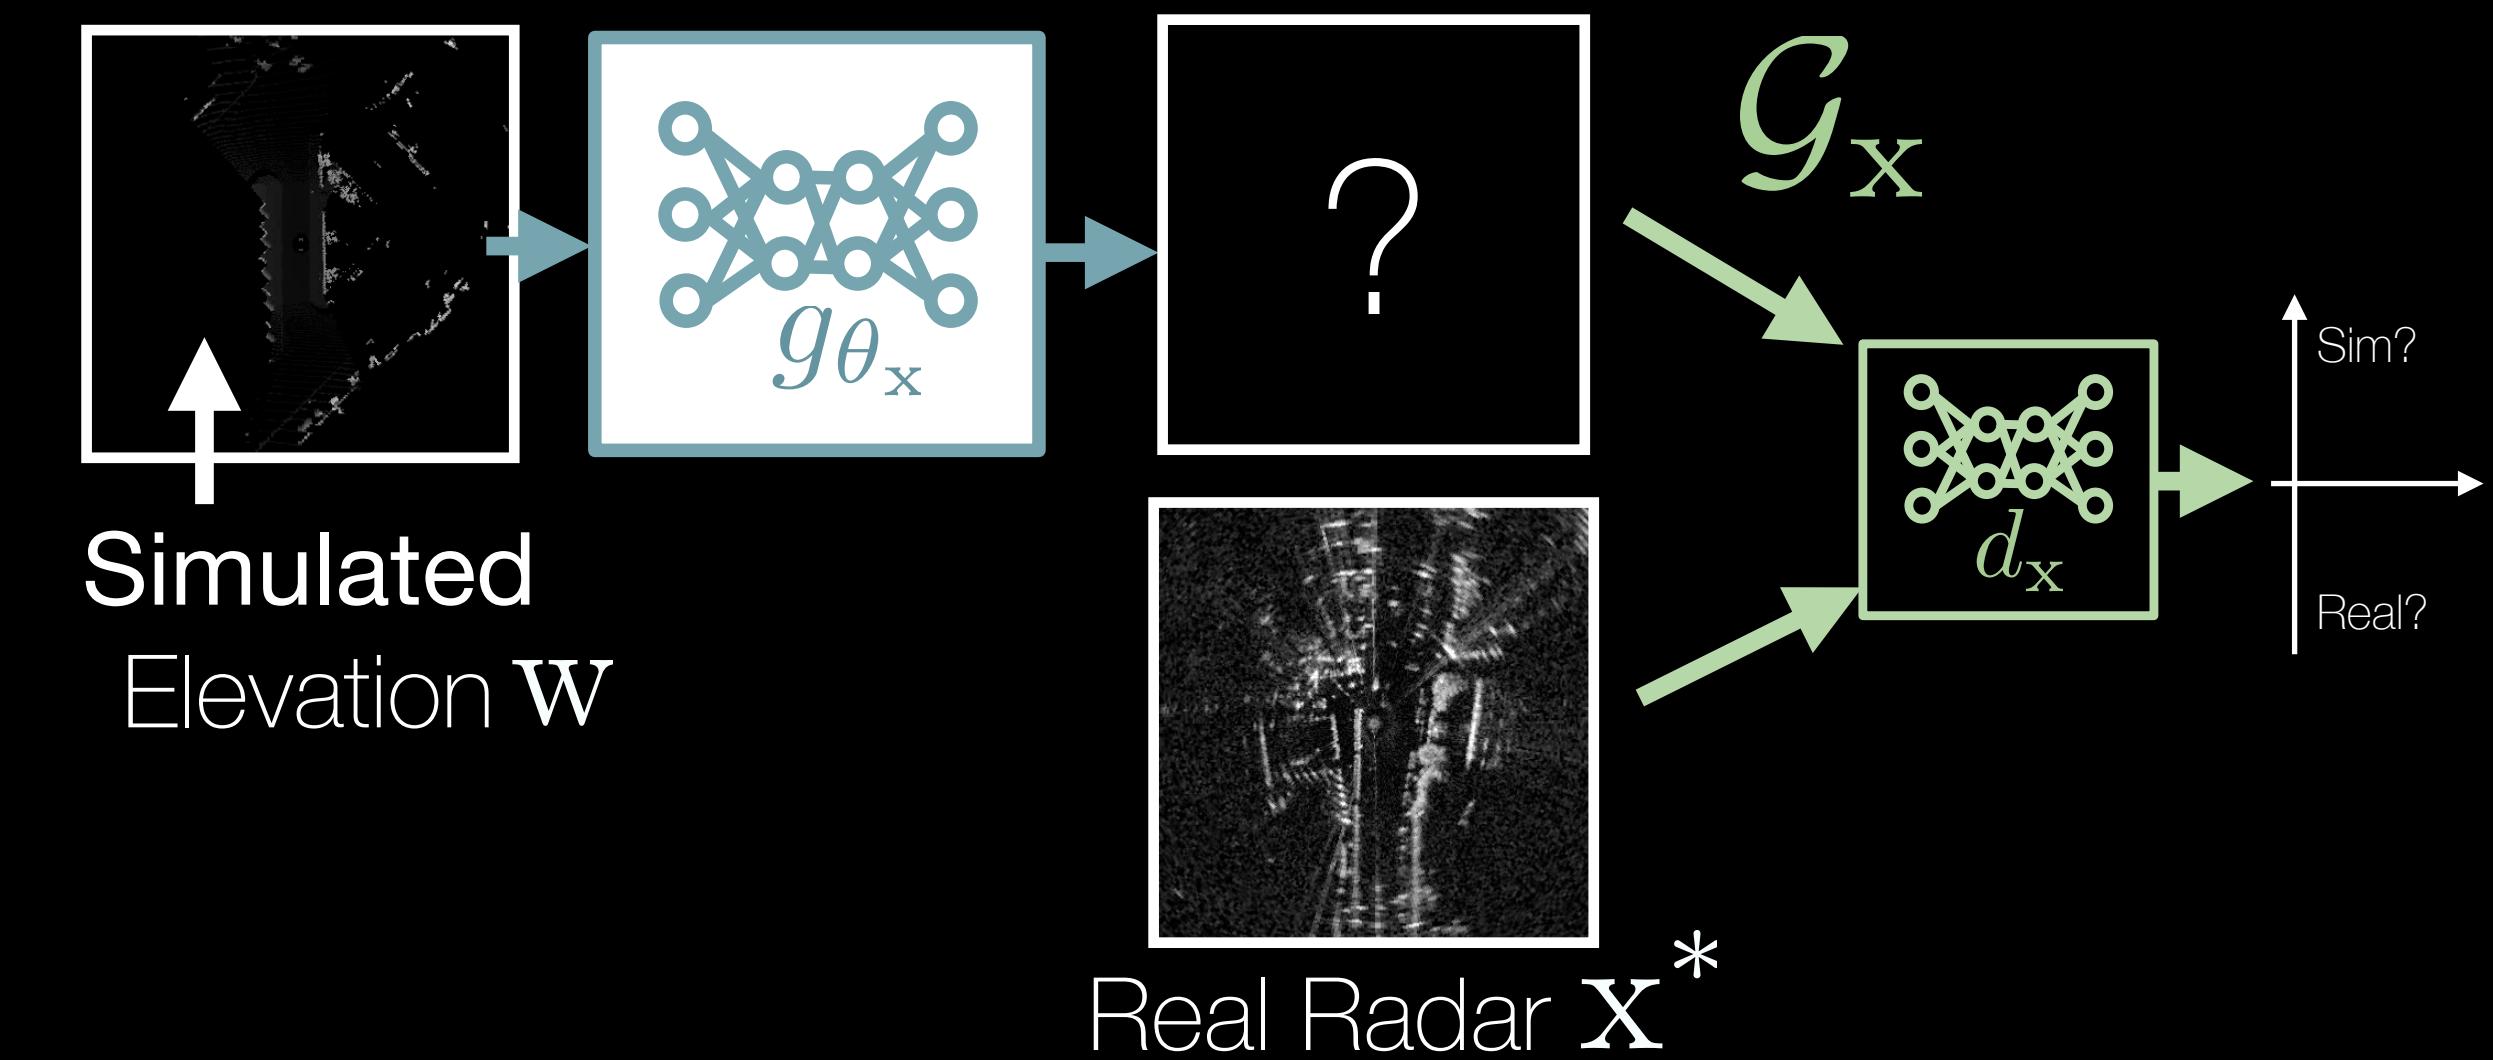
Performs poorly on simulated elevation maps



# Results

## Radar Simulation

- (c) trained to predict radar from **simulated** data using only **only** GAN loss

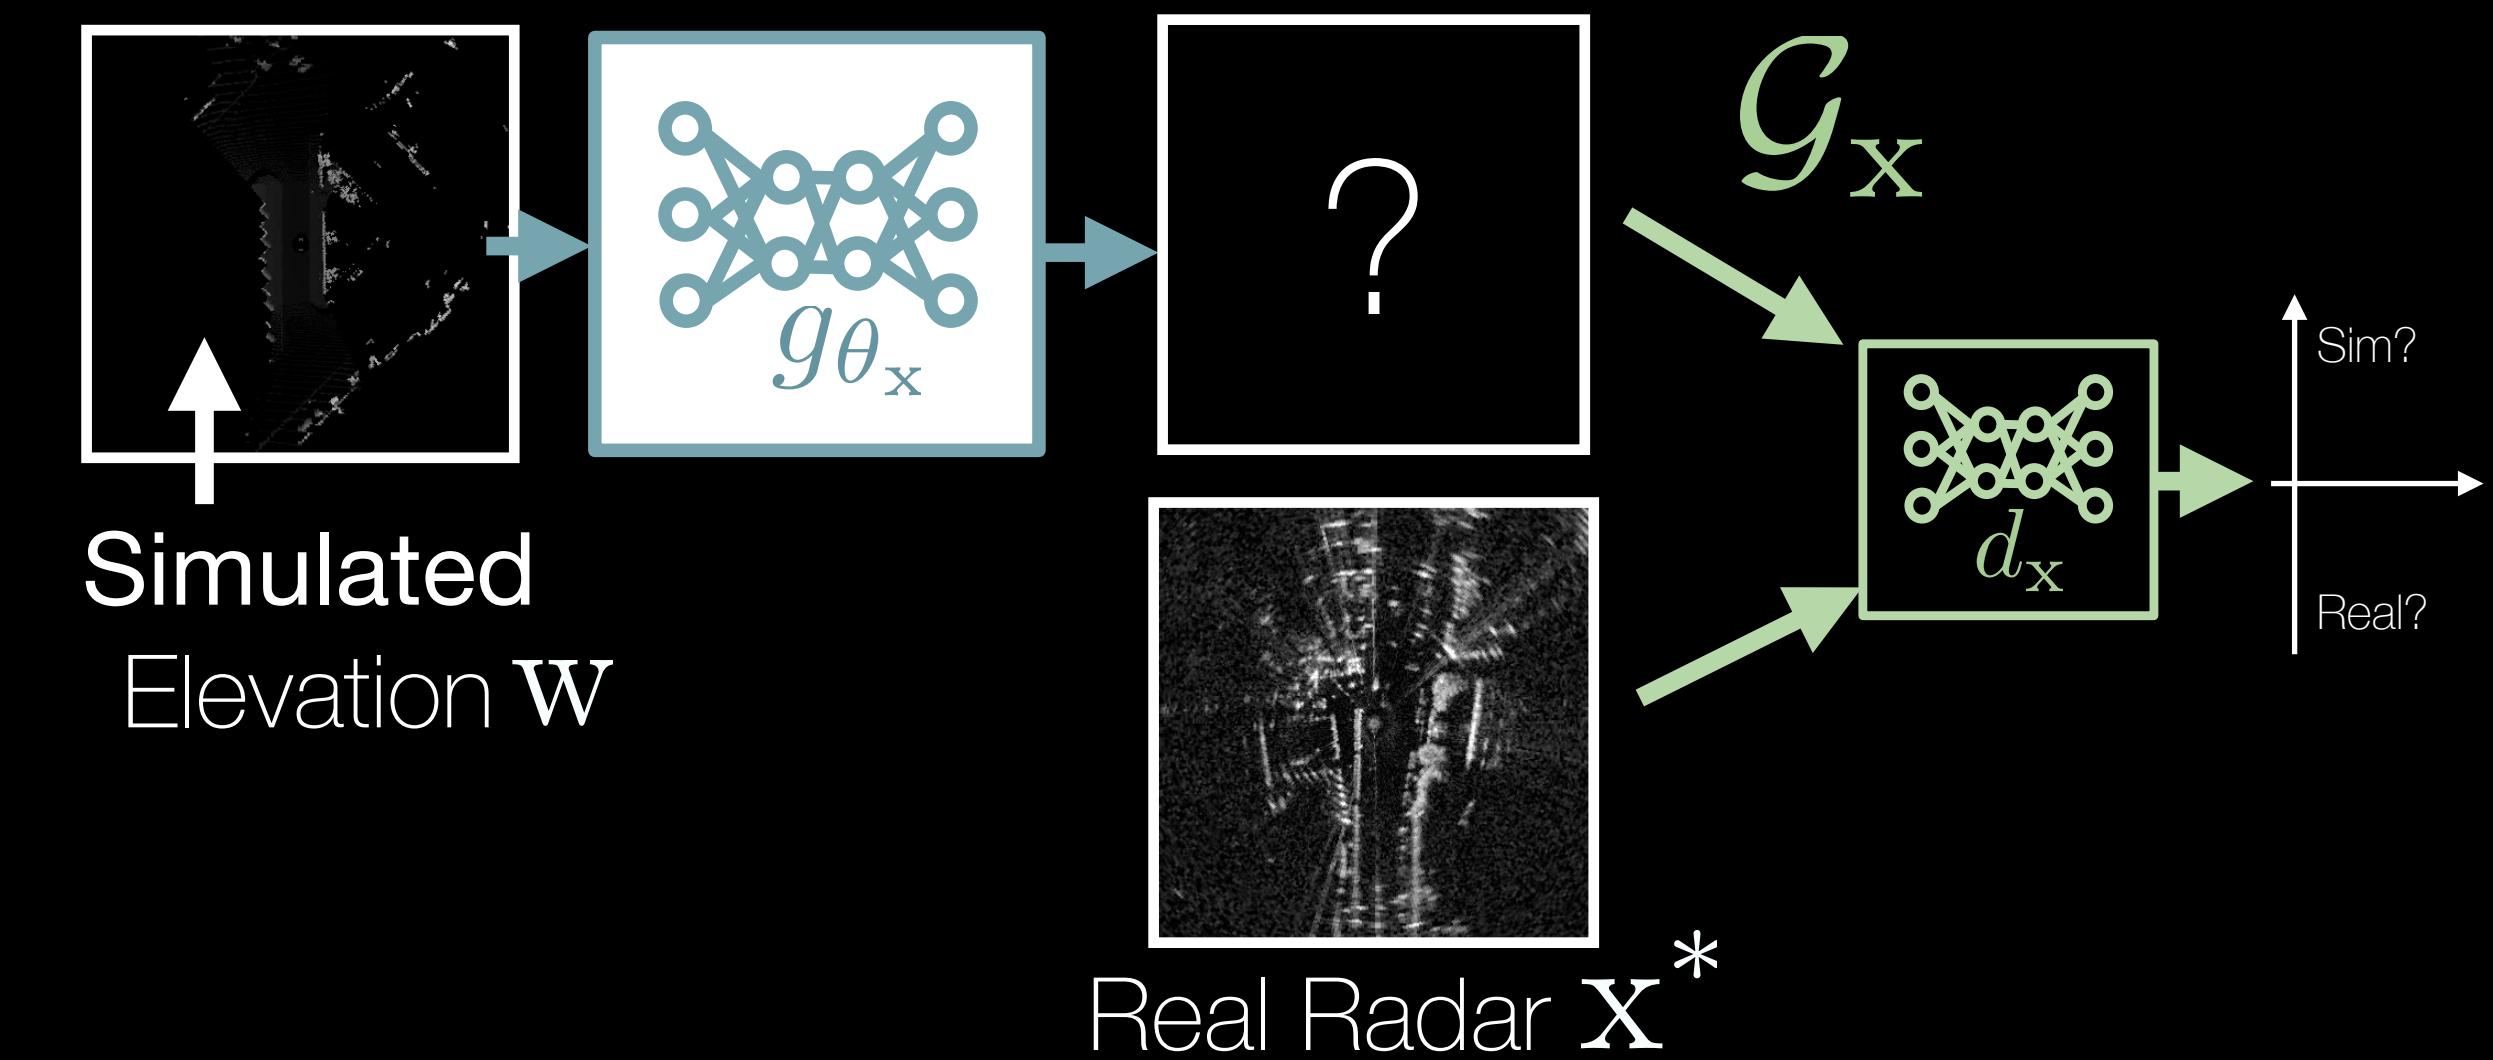


benchmark	trained on	training objective						intersection over union		
		$\mathcal{A}_x$	$\mathcal{A}_w$	$\mathcal{G}_x$	$\mathcal{G}_w$	$\mathcal{C}_x$	$\mathcal{C}_w$	free	occ	mean
real world	-	-	-	-	-	-	-	0.856	0.553	0.705
<b>ours</b>										
(a)	$\mathbf{x}^*$	$\mathbf{m}^*$	-	✓	-	-	-	0.396 (0.00)	0.275 (0.00)	0.335 (0.00)
(b)	$\mathbf{x}^*$	$\mathbf{m}^*$	-	✓	-	✓	-	0.558 (0.11)	0.221 (0.03)	0.389 (0.07)
(c)	$\mathbf{x}^*$	-	$\mathbf{w}$	-	-	✓	-	0.385 (0.07)	0.148 (0.01)	0.266 (0.04)
(d)	$\mathbf{x}^*$	-	$\mathbf{w}$	-	-	✓	✓	0.845 (0.02)	0.262 (0.02)	0.553 (0.02)
(e)	$\mathbf{x}^*$	$\mathbf{m}^*$	$\mathbf{w}$	-	✓	✓	✓	✓	<b>0.872 (0.01)</b>	<b>0.455 (0.01)</b>
										<b>0.664 (0.01)</b>

# Results

## Radar Simulation

- (c) trained to predict radar from **simulated** data using only **only** GAN loss

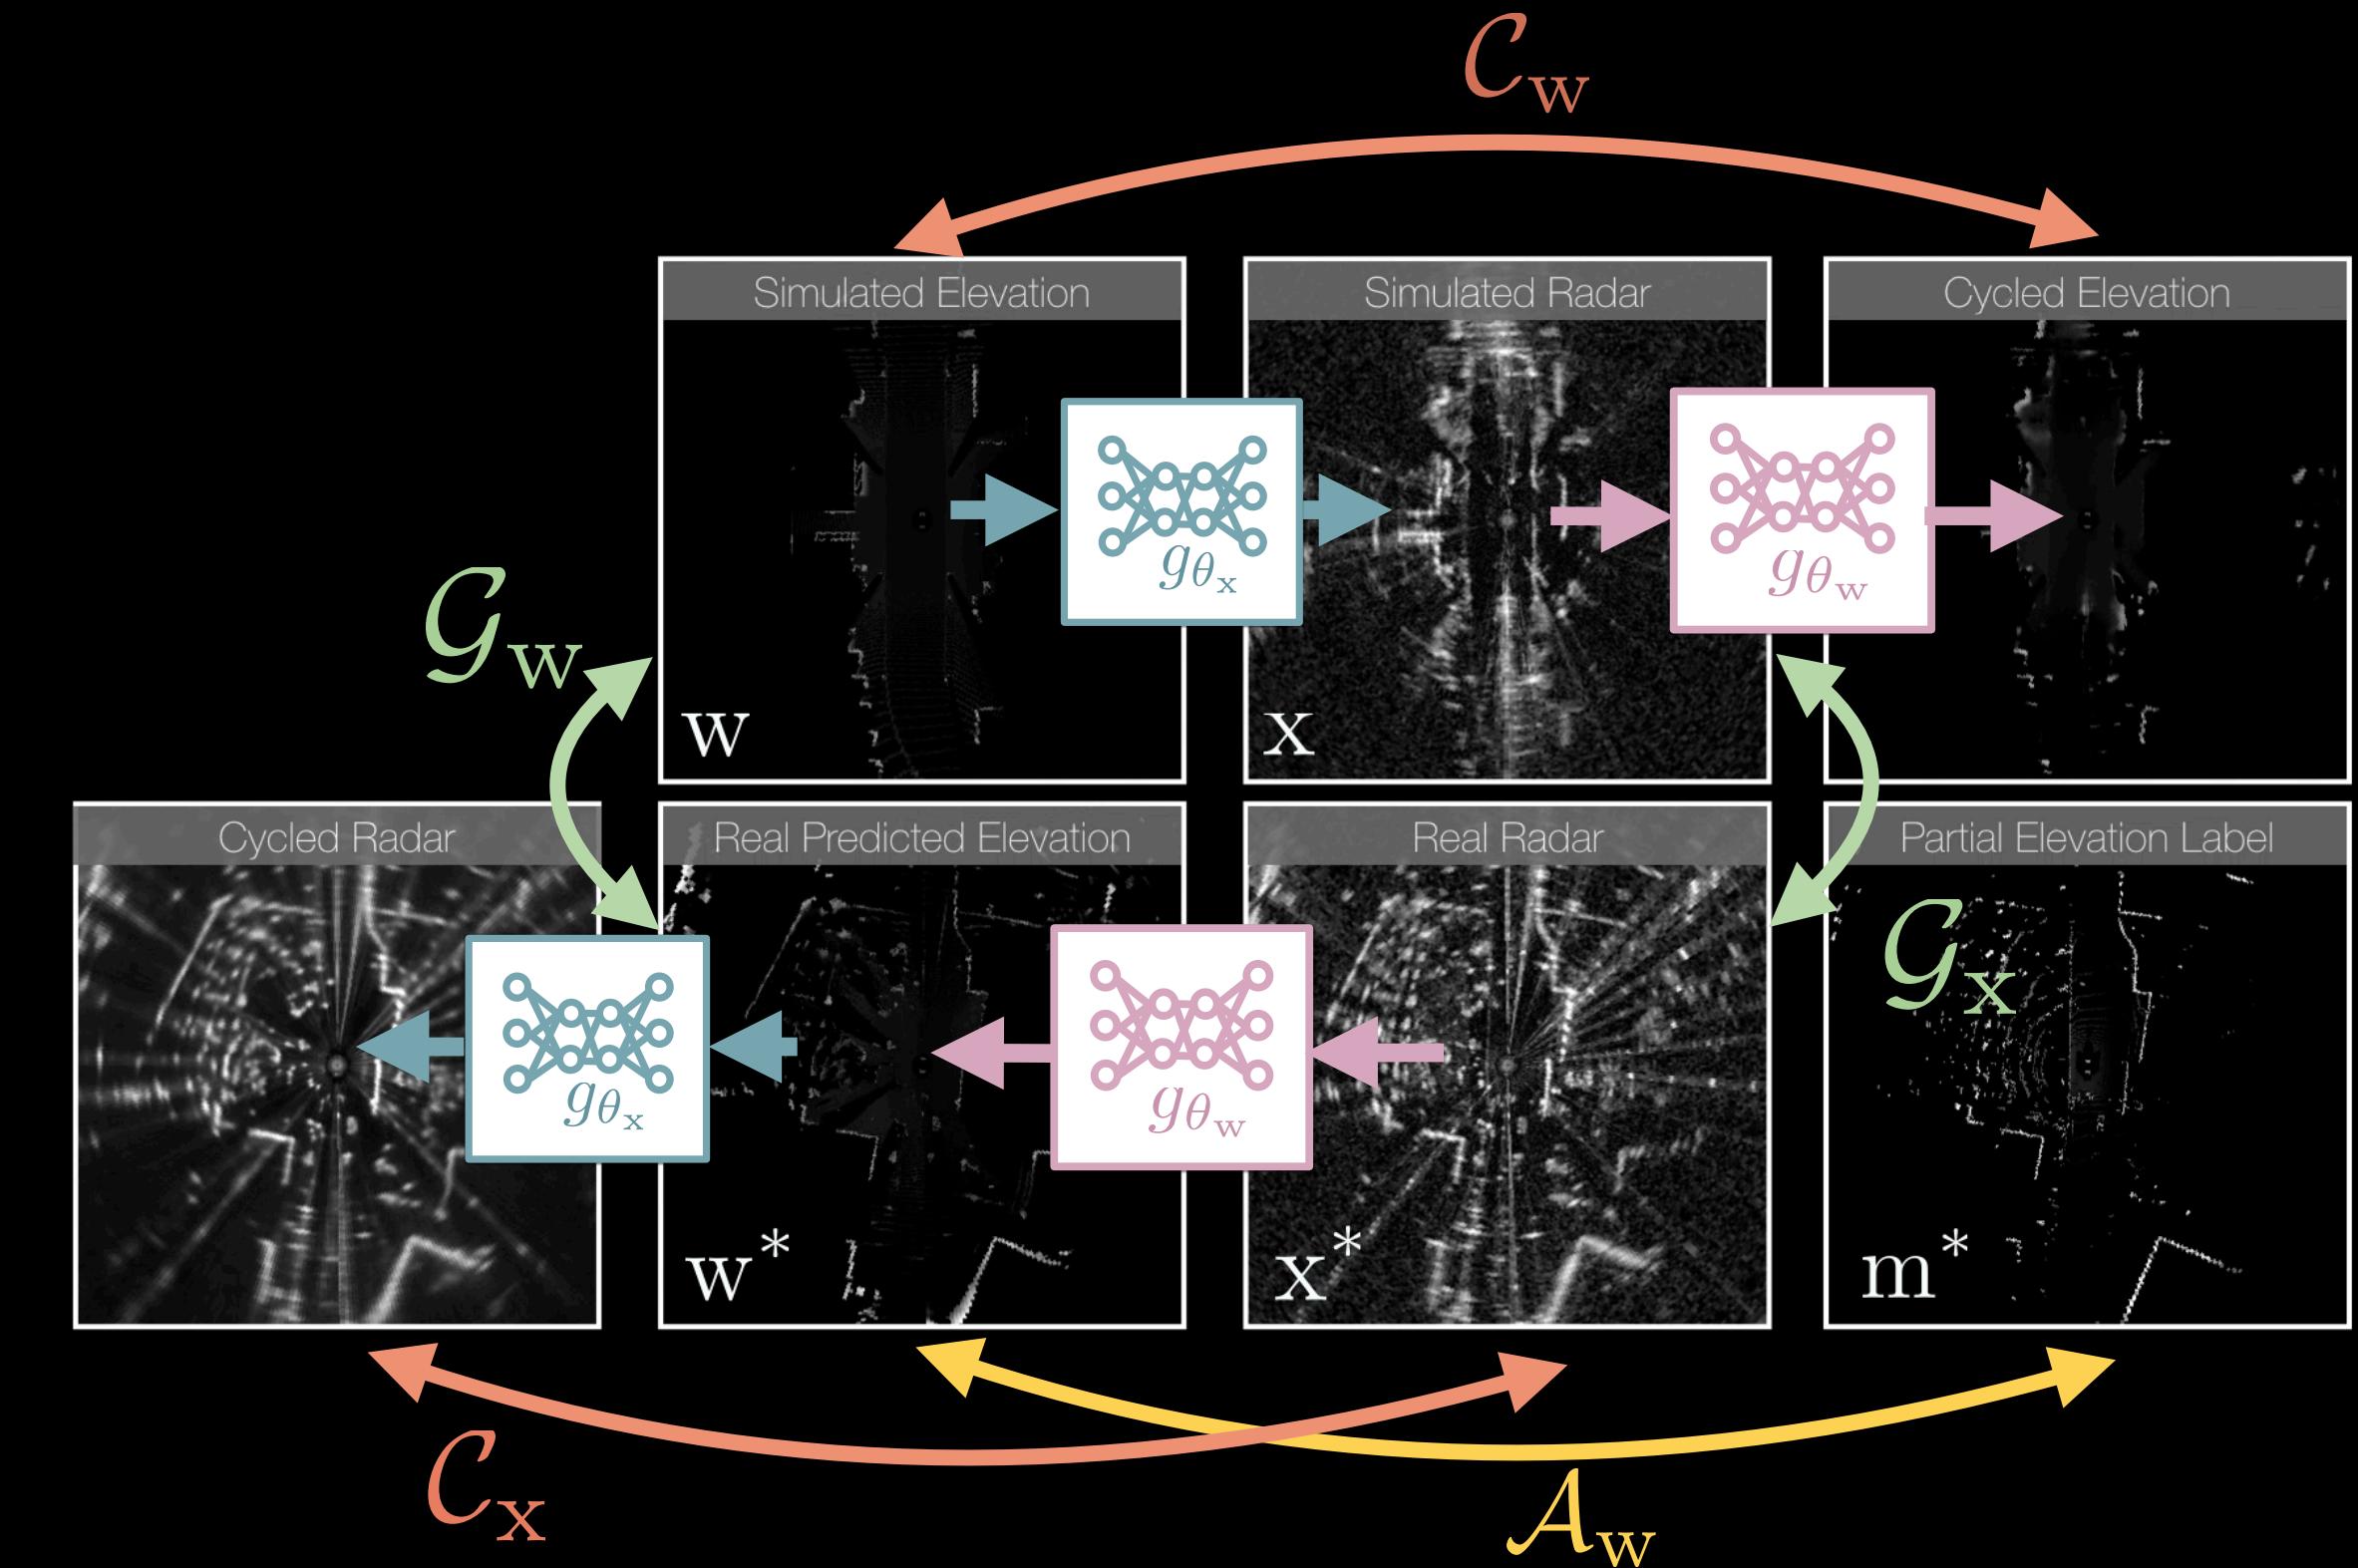


benchmark	trained on	training objective						intersection over union		
		$\mathcal{A}_x$	$\mathcal{A}_w$	$\mathcal{G}_x$	$\mathcal{G}_w$	$\mathcal{C}_x$	$\mathcal{C}_w$	free	occ	mean
real world	-	-	-	-	-	-	-	0.856	0.553	0.705
<b>ours</b>										
(a)	$\mathbf{x}^*$	$\mathbf{m}^*$	-	✓	-	-	-	0.396 (0.00)	0.275 (0.00)	0.335 (0.00)
(b)	$\mathbf{x}^*$	$\mathbf{m}^*$	-	✓	-	✓	-	0.558 (0.11)	0.221 (0.03)	0.389 (0.07)
(c)	$\mathbf{x}^*$	-	$\mathbf{w}$	-	-	✓	-	0.385 (0.07)	0.148 (0.01)	0.266 (0.04)
(d)	$\mathbf{x}^*$	-	$\mathbf{w}$	-	-	✓	✓	0.845 (0.02)	0.262 (0.02)	0.553 (0.02)
(e)	$\mathbf{x}^*$	$\mathbf{m}^*$	$\mathbf{w}$	-	✓	✓	✓	✓	<b>0.872 (0.01)</b>	<b>0.455 (0.01)</b>
										<b>0.664 (0.01)</b>

# Results

## Radar Simulation

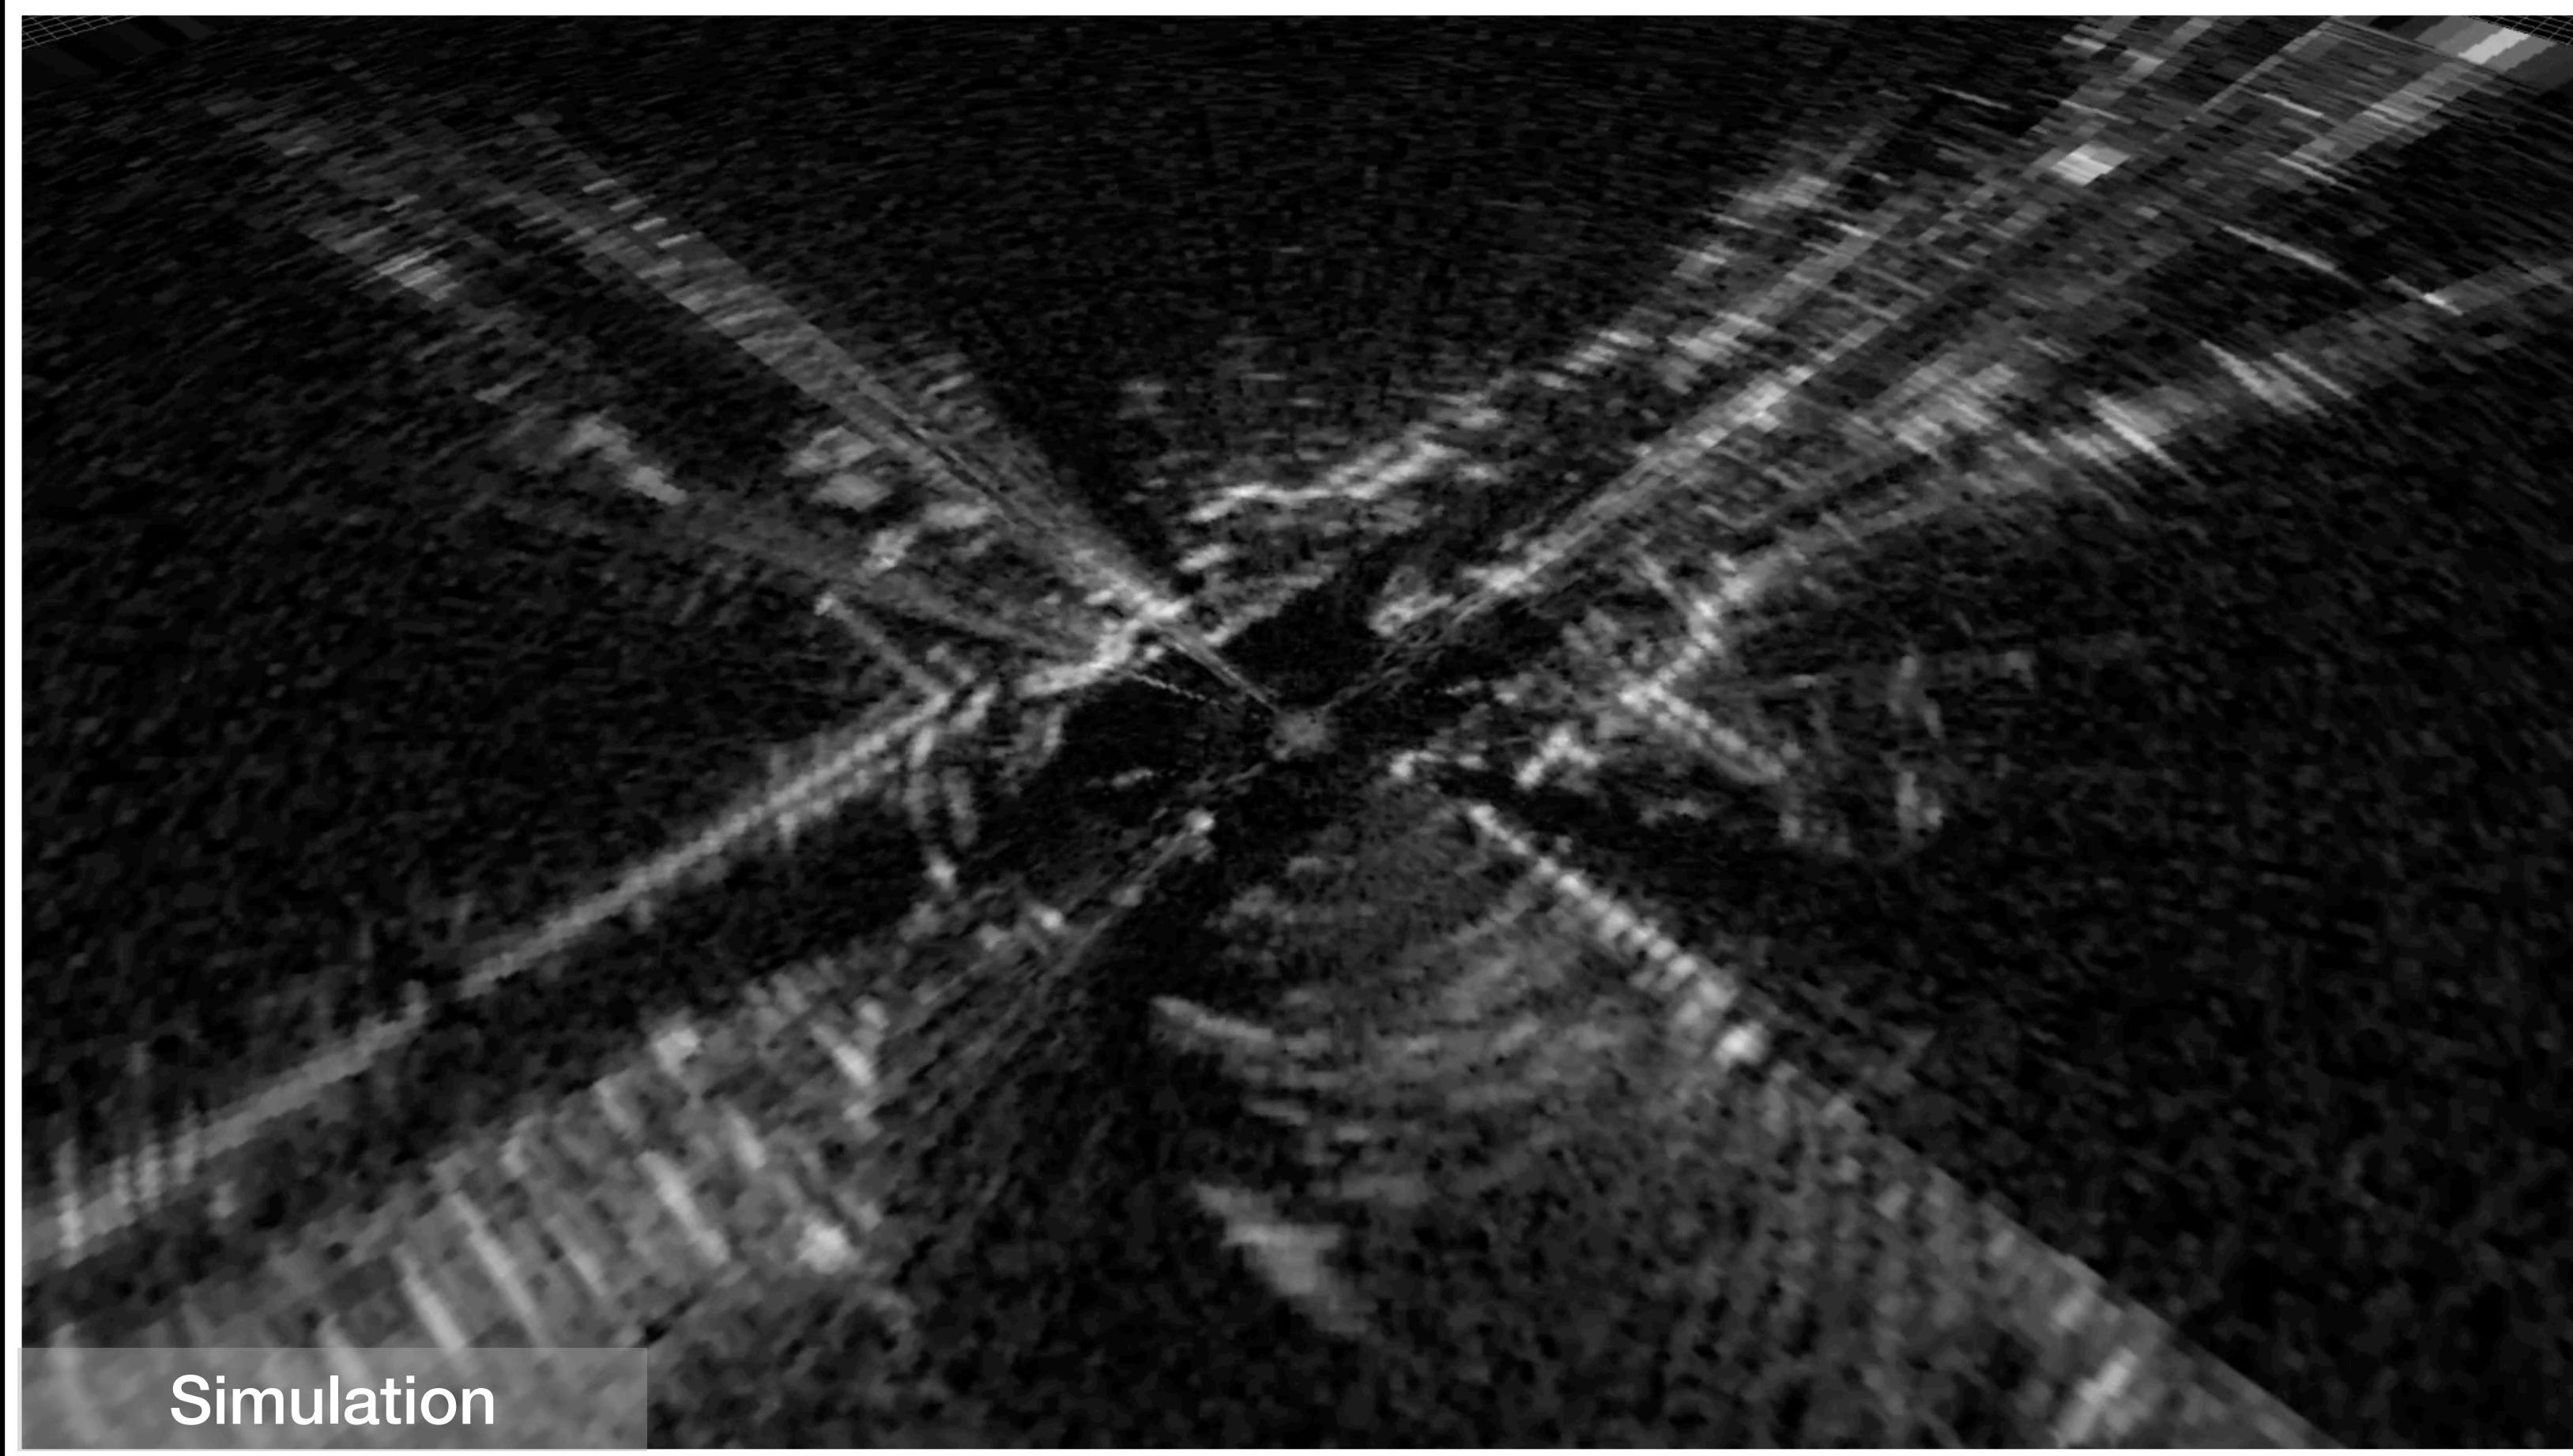
- (d) trained with GAN loss + cycle constraints
- (e) trained with **full** training criterion
  - Best performing model!
  - Only 5% percentage points off benchmark



trained on	training objective						intersection over union		
	$\mathcal{A}_x$	$\mathcal{A}_w$	$\mathcal{G}_x$	$\mathcal{G}_w$	$\mathcal{C}_x$	$\mathcal{C}_w$	free	occ	mean
<b>benchmark</b>	-	-	-	-	-	-	0.856	0.553	0.705
real world	-	-	-	-	-	-			
<b>ours</b>									
(a)	x*	m*	-	✓	-	-	0.396 (0.00)	0.275 (0.00)	0.335 (0.00)
(b)	x*	m*	-	✓	-	✓	0.558 (0.11)	0.221 (0.03)	0.389 (0.07)
(c)	x*	-	w	-	-	✓	0.385 (0.07)	0.148 (0.01)	0.266 (0.04)
(d)	x*	-	w	-	-	✓	0.845 (0.02)	0.262 (0.02)	0.553 (0.02)
(e)	x*	m*	w	-	✓	✓	✓	✓	0.872 (0.01) 0.455 (0.01) 0.664 (0.01)

# Results

## Radar Simulation

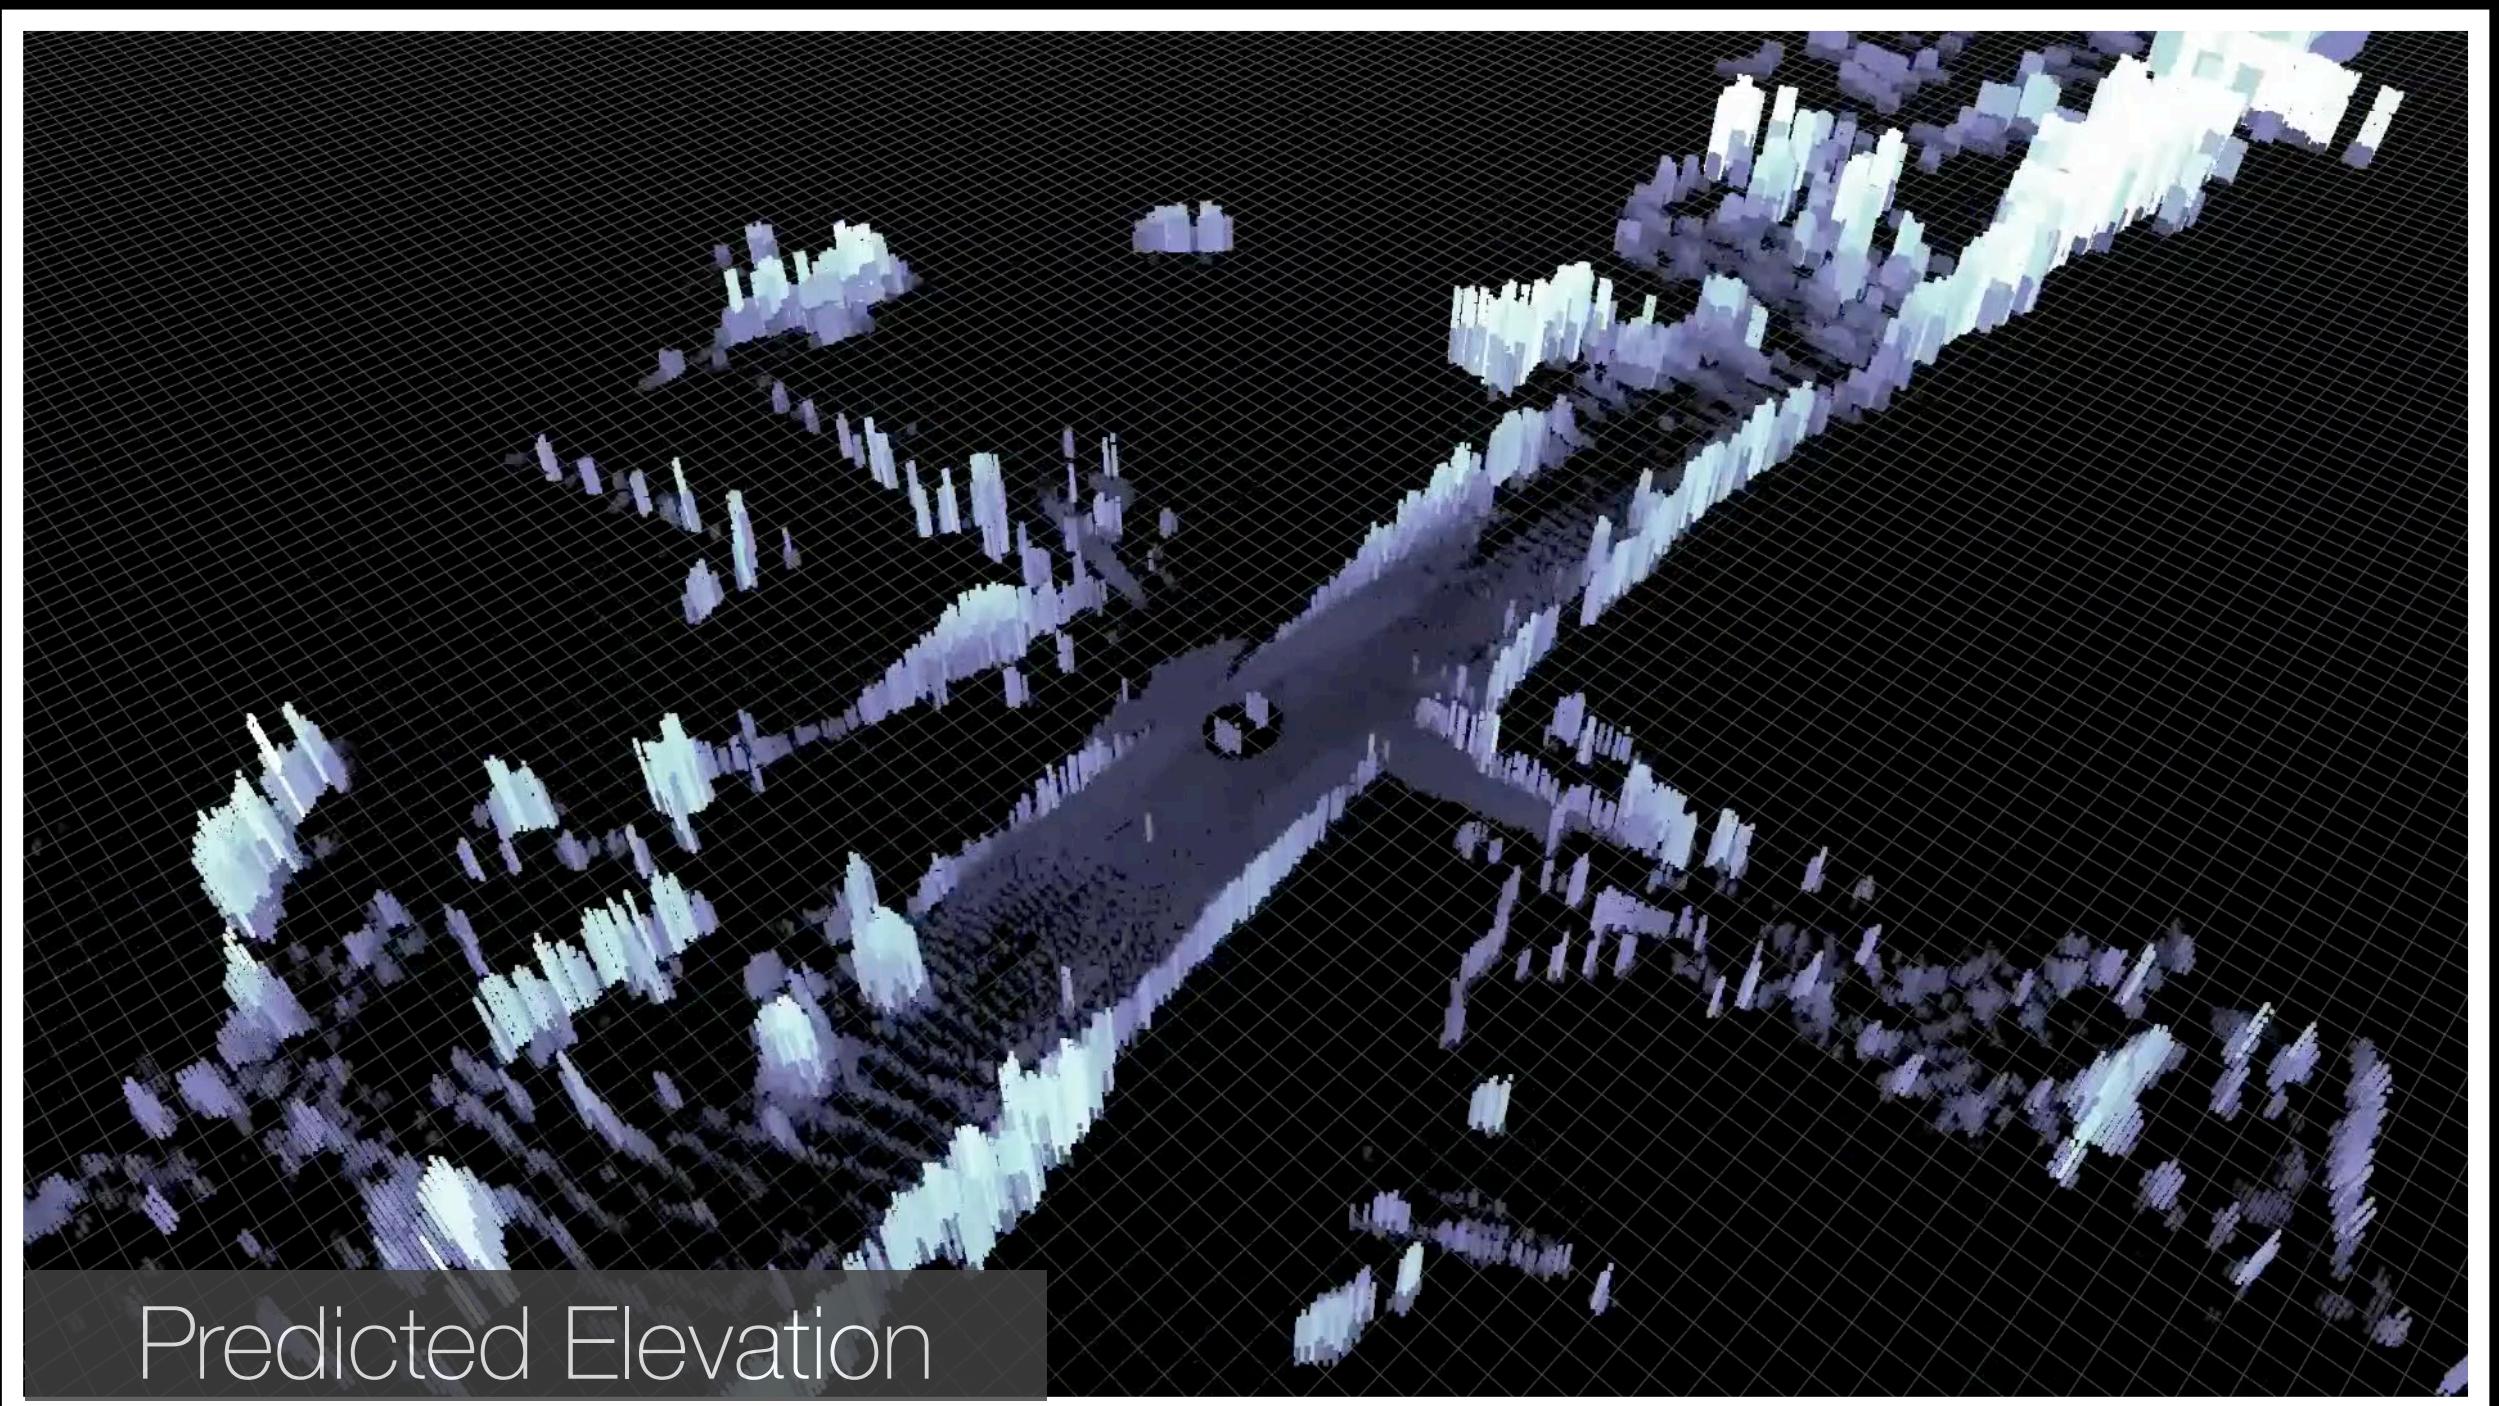
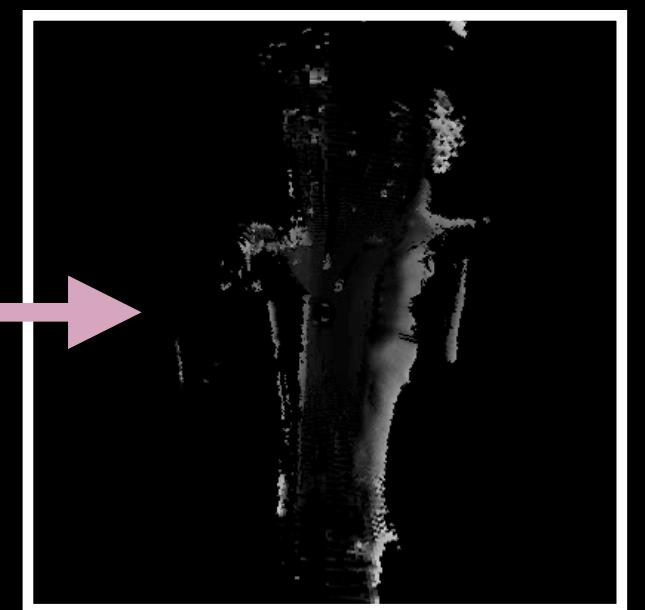
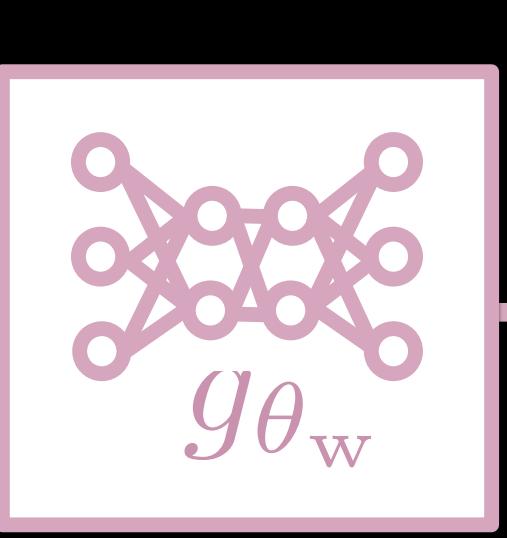
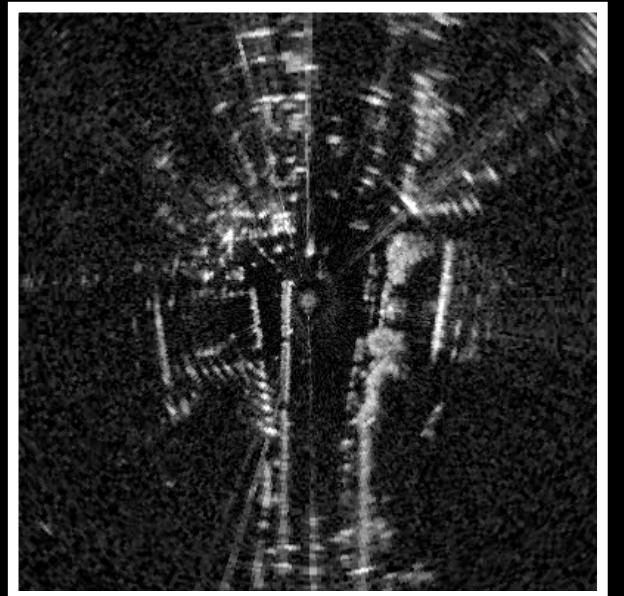


Successfully captures many of radars characteristic noise artefacts

# Results

## Backward Model

- Learnt as part of the same training setup
- Allows the elevation state to be inferred from real 2D radar data



# Results

## Backward Model Evaluation

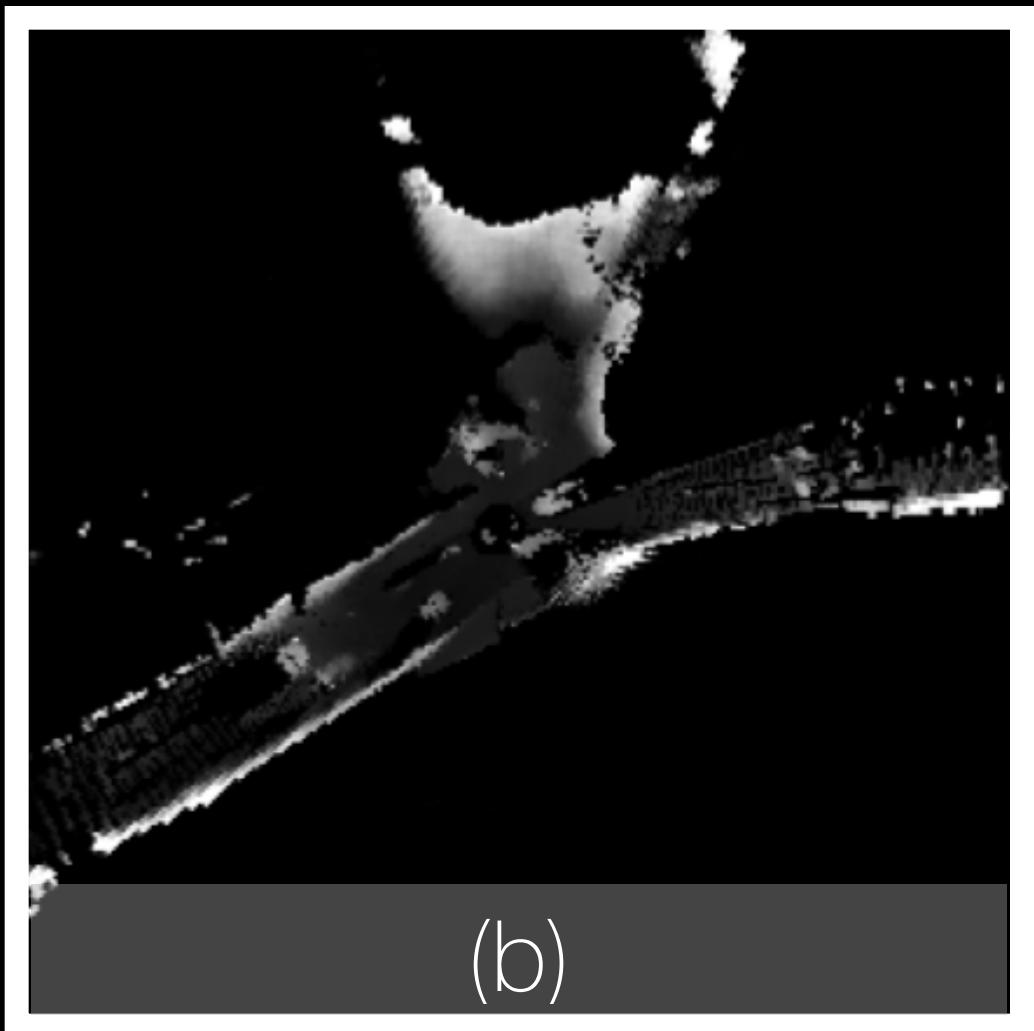
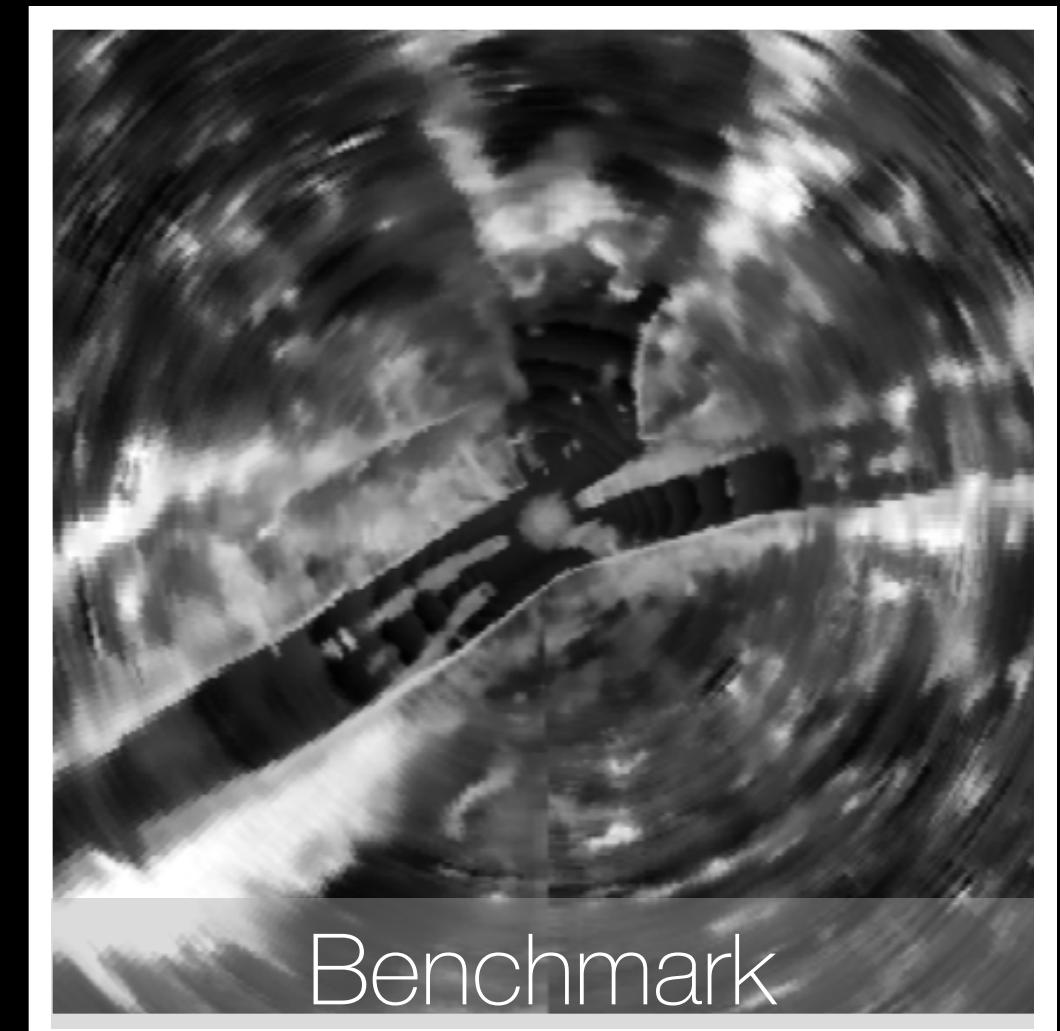
- Evaluated against ground truth elevation labels using **Mean Absolute Error**
- **Benchmark** = “model trained directly regressing to partial elevation labels”
- **Ours:**
  - (a) trained with **just** cycle consistency
  - (b) trained with cycle consistency **and** alignment

	data	$\mathcal{A}_z$	$\mathcal{G}_w$	$\mathcal{C}_w$	free	occ	mean
<b>Benchmark</b>							
direct regression	R	✓	-	-	18.7	7.5	13.1
<b>Ours</b>							
(a)	R, S	-	✓	✓	3.9 (0.4)	74.2 (1.8)	39.0 (0.9)
(b)	R, S	✓	✓	✓	4.1 (0.6)	41.1 (5.7)	22.6 (2.6)

# Results

## Backward Model Evaluation

- Evaluated against ground truth elevation labels using **Mean Absolute Error**
- Benchmark** = “model trained directly regressing to partial elevation labels”



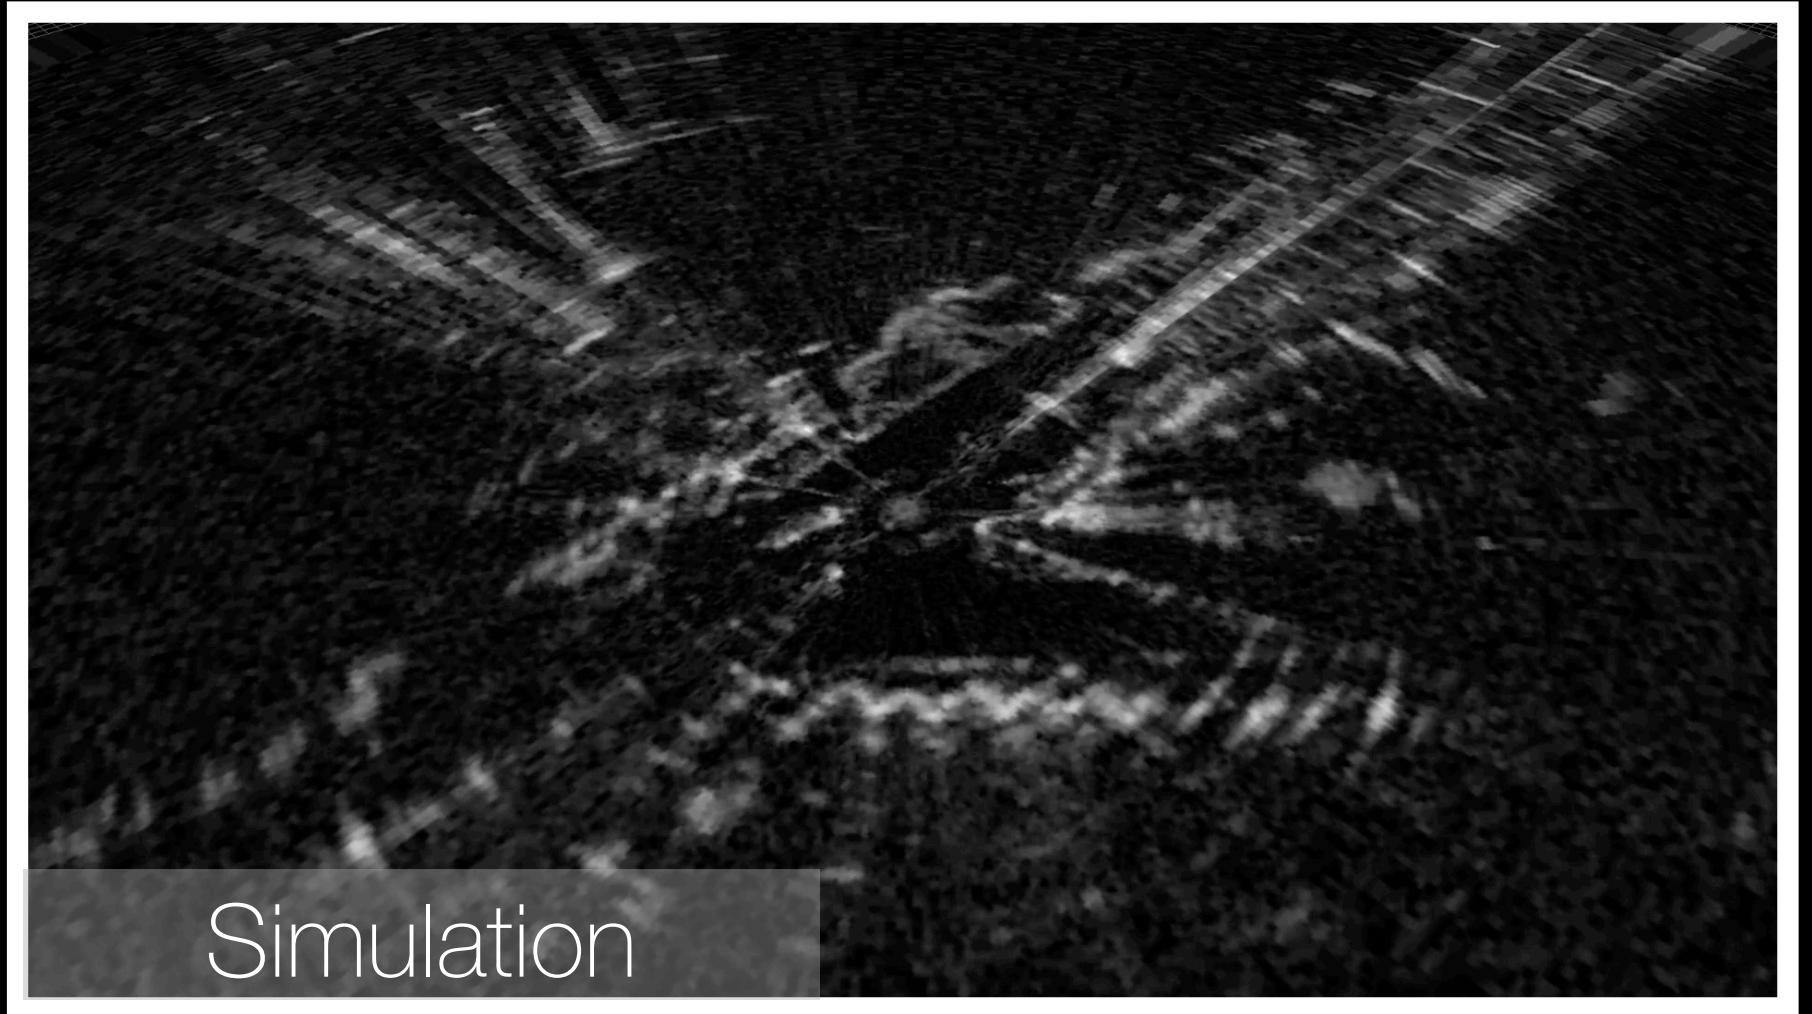
- Ours:**
  - (a) trained with **just** cycle consistency
  - (b) trained with cycle consistency **and** alignment

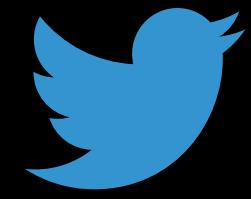
	data	$\mathcal{A}_z$	$\mathcal{G}_w$	$\mathcal{C}_w$	Mean Absolute Error (cm)		
					free	occ	mean
<b>Benchmark</b>							
direct regression	R	✓	-	-	18.7	7.5	13.1
<b>Ours</b>							
(a)	R, S	-	✓	✓	3.9 (0.4)	74.2 (1.8)	39.0 (0.9)
(b)	R, S	✓	✓	✓	4.1 (0.6)	41.1 (5.7)	22.6 (2.6)

Only a few centimetres off  
benchmark whilst **generalising** to  
regions of space where elevation  
labels do not exist

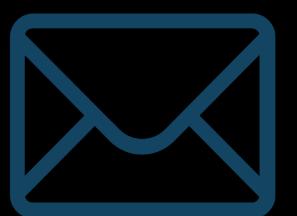
# Conclusion

- A **data-driven** approach to **stochastic** radar simulation **interfacing with commonly available** simulation environments
- Model the forward and backward processes **side-by-side**
  - Train on **simulated** inputs using adversarial losses, cyclical consistency and partial elevation labels
  - Train **new** radar models in simulation (for the first time)
    - Only 4 percentage points off a segmentation model trained in the real world
    - Backward model also infers the height state of the world to a **reasonable accuracy** (4-40cm) whilst remaining **regularised** where no elevation labels exist





robw4



robw@robots.ox.ac.uk

1524

Thanks for listening!

